

Homework 2: Deep Learning

Due June 2, beginning of exercise session*

Kristian Kersting, Alejandro Molina

{kersting, molina}@cs.tu-darmstadt.de upload link:

<https://www.dropbox.com/request/OEyNaA11EGyiexfyadSa>

0. Assuming that you have implemented a feed-forward neural network architecture with the following constraints.




Figure 1: A feed-forward neural network architecture consisting of two hidden layers.

- The sigmoid activation function is used for computing hidden activations.
- No activation function in the output layer.
- In this homework, please use the following code snippet to initialize all parameters, inputs and their ground truth in the FIRST cell of your notebook.

```
# Python CODE  
import numpy as np  
# DO NOT FORGET TO SPECIFY THE SAME SEED  
np.random.seed(12345)
```

*We will discuss the solutions in the exercise session. It is my suggestion that you try to address at least 50% of the exercise questions. Simply try hard to solve them. This way, you will get familiar with the technical terms and with the underlying ideas of the lecture.

```

def initialize(input_dim, hidden_dim, output_dim, batch_size):
    W1 = np.random.randn(hidden_dim, input_dim) * 0.01
    b1 = np.zeros((hidden_dim,))
    W2 = np.random.randn(hidden_dim, hidden_dim) * 0.01
    b2 = np.zeros((hidden_dim,))
    W3 = np.random.randn(output_dim, hidden_dim) * 0.01
    b3 = np.zeros((output_dim,))

    # list of all network parameters
    parameters = [W1, b1, W2, b2, W3, b3]

    # minibatch of input instances
    x = np.random.rand(input_dim, batch_size)

    # ground truths
    y = np.random.randn(output_dim, batch_size)

    return parameters, x, y

# initialize parameters, inputs and targets
parameters, x, y = initialize(3, 4, 2, 5)

```

And the squared loss function:

$$\mathcal{L}(\theta; \mathbf{x}, \mathbf{y}) = \frac{1}{M} \sum_{n=1}^M \frac{1}{2} \|f_{\theta}(\mathbf{x}_n) - \mathbf{y}_n\|^2 \quad (1)$$

where f_{θ} denotes the outputs of the network given inputs \mathbf{x} , \mathbf{y} is the targets, and θ denotes a set of the parameters.

1. Write a program that computes the *gradients* for every parameter given the loss using back propagation.
2. Compare to the gradients you obtained in the previous homework, and submit your notebook file that shows the following at the end:

- the numerical and backprop gradients of the loss function with respect to each parameter separately

$$\frac{\partial \mathcal{L}}{\partial \mathbf{W}_1}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_1}, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_2}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_2}, \frac{\partial \mathcal{L}}{\partial \mathbf{W}_3}, \frac{\partial \mathcal{L}}{\partial \mathbf{b}_3}$$