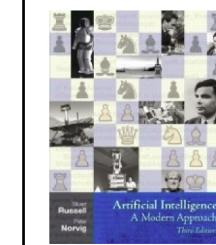


Outline

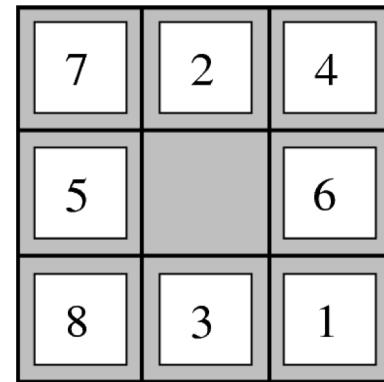
- Best-first search
 - Greedy best-first search
 - A* search
 - Heuristics
 - Admissible Heuristics
 - Graph Search
 - Consistent Heuristics
- Local search algorithms
 - Hill-climbing search
 - Beam search
 - Simulated annealing search
 - Genetic algorithms
- Constraint Satisfaction Problems



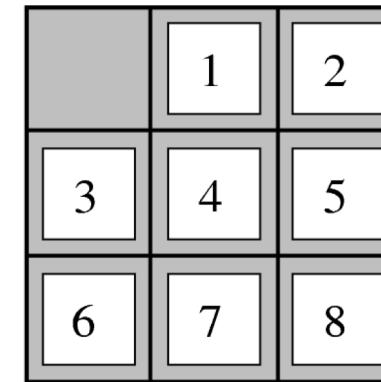
Many slides based on
Russell & Norvig's slides
[Artificial Intelligence:
A Modern Approach](#)

Motivation

- Uninformed search algorithms are too inefficient
 - they expand far too many unpromising paths
- Example:
 - 8-puzzle



Start State



Goal State

- Average solution depth = 22
- Breadth-first search to depth 22 has to expand about 3.1×10^{10} nodes

→ try to be more clever with what nodes to expand

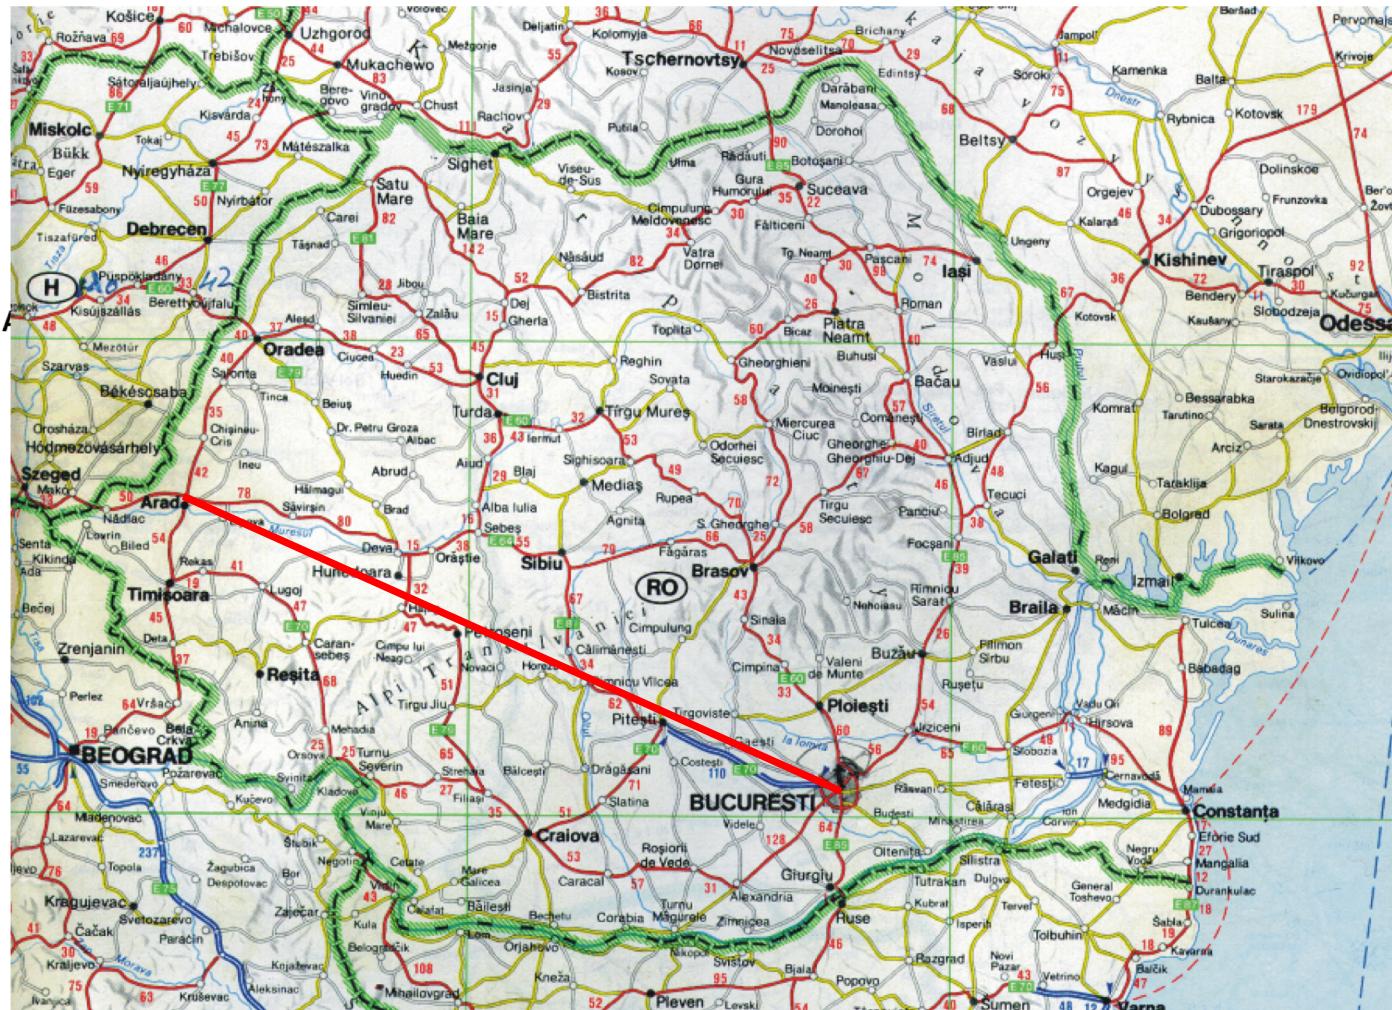
Best-First Search

- Recall
 - Search strategies are characterized by the order in which they expand the nodes of the search tree
 - Uninformed tree-search algorithms sort the nodes by problem-independent methods (e.g., recency)
- Basic Idea of Best-First Search
 - use a **heuristic evaluation function** $f(n)$ for each node
 - estimate of the "desirability" of the node's state
 - expand most desirable unexpanded node
- Implementation
 - use Tree Search algorithm
 - order the nodes in fringe in decreasing order of desirability
- Algorithms
 - Greedy best-first search
 - A* search

Heuristic

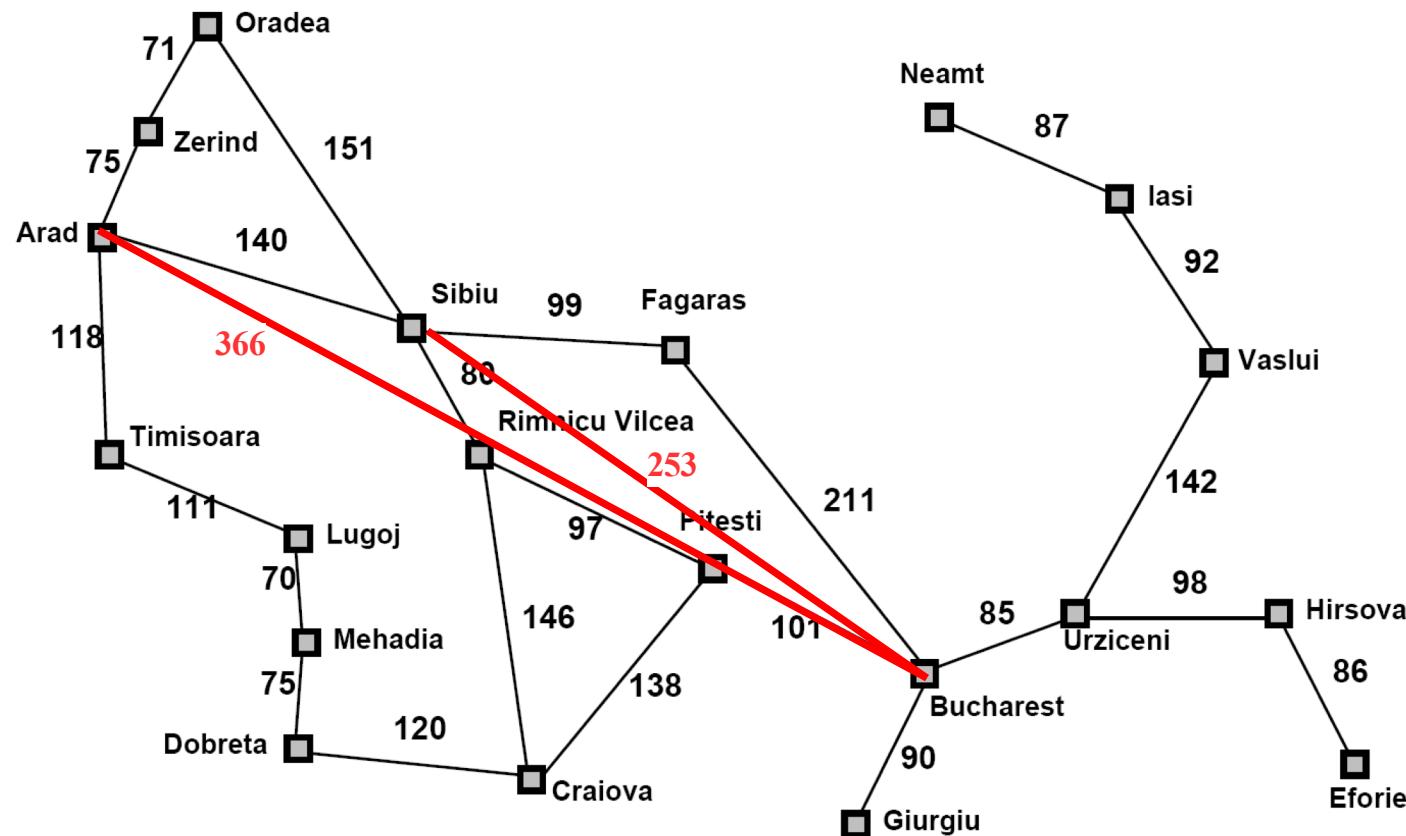
- Greek "heurisko" (εύρισκω) → "I find"
 - cf. also „Eureka!“
- informally denotes a „rule of thumb“
 - i.e., knowledge that may be helpful in solving a problem
 - note that heuristics may also go wrong!
- In tree-search algorithms, a heuristic denotes a function that estimates the remaining costs until the goal is reached
- Example:
 - straight-line distances may be a good approximation for the true distances on a map of Romania
 - and are easy to obtain (ruler on the map)
 - but cannot be obtained directly from the distances on the map

Romania Example: Straight-line Distances



Straight-line distance
to Bucharest

Romania Example: Straight-line Distances



Straight-line distance
to Bucharest

| | |
|----------------|-----|
| Arad | 366 |
| Bucharest | 0 |
| Craiova | 160 |
| Dobreta | 242 |
| Eforie | 161 |
| Fagaras | 176 |
| Giurgiu | 77 |
| Hirsova | 151 |
| Iasi | 226 |
| Lugoj | 244 |
| Mehadia | 241 |
| Neamt | 234 |
| Oradea | 380 |
| Pitesti | 98 |
| Rimnicu Vilcea | 193 |
| Sibiu | 253 |
| Timisoara | 329 |
| Urziceni | 80 |
| Vaslui | 199 |
| Zerind | 374 |

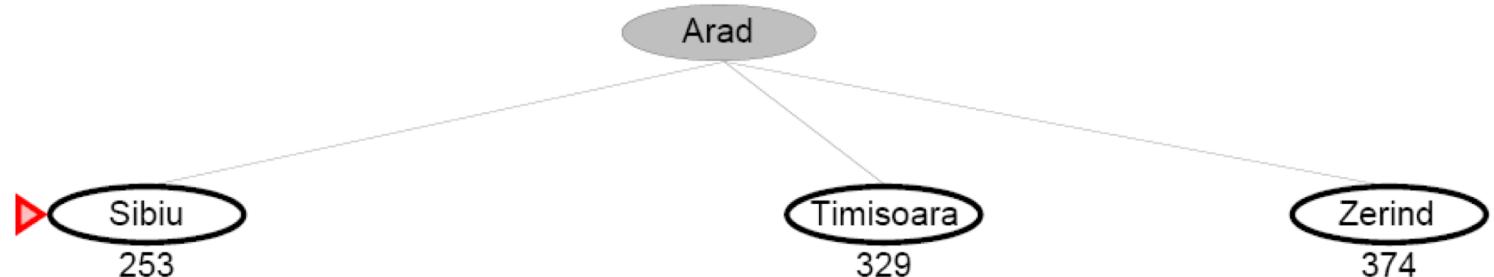
Greedy Best-First Search

- Evaluation function $f(n) = h(n)$ (*heuristic*)
 - estimates the cost from node n to *goal*
 - e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal
 - according to evaluation function
- Example:



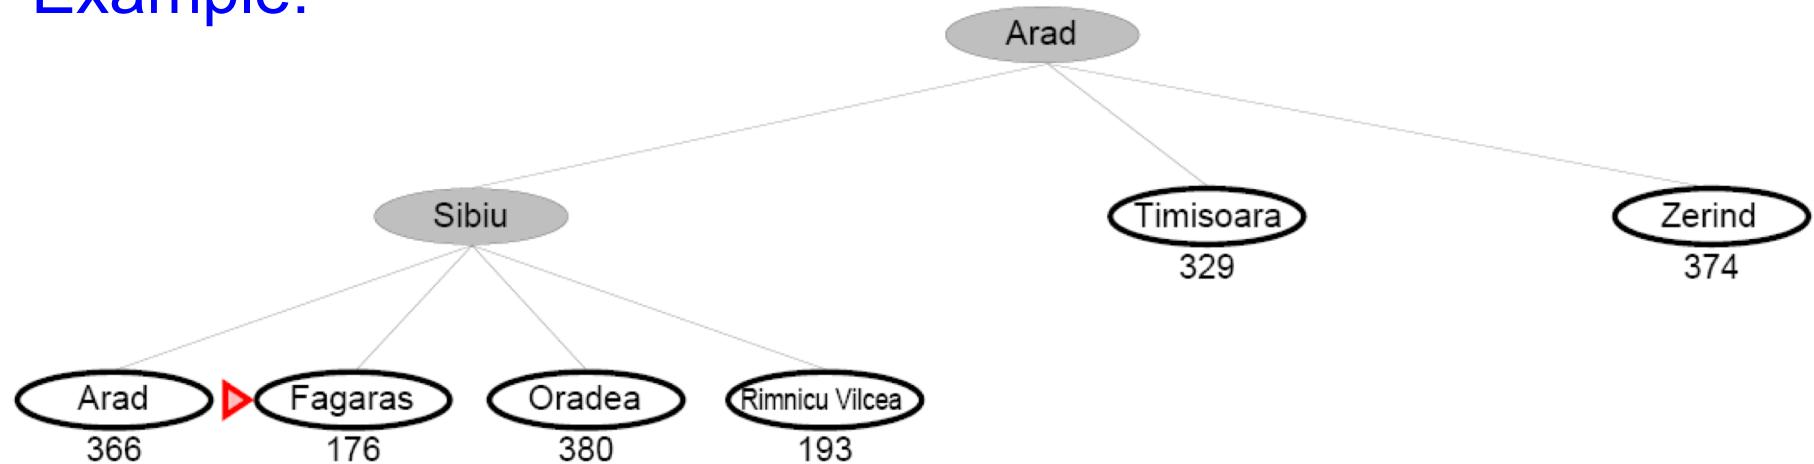
Greedy Best-First Search

- Evaluation function $f(n) = h(n)$ (*heuristic*)
 - estimates the cost from node n to *goal*
 - e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal
 - according to evaluation function
- Example:



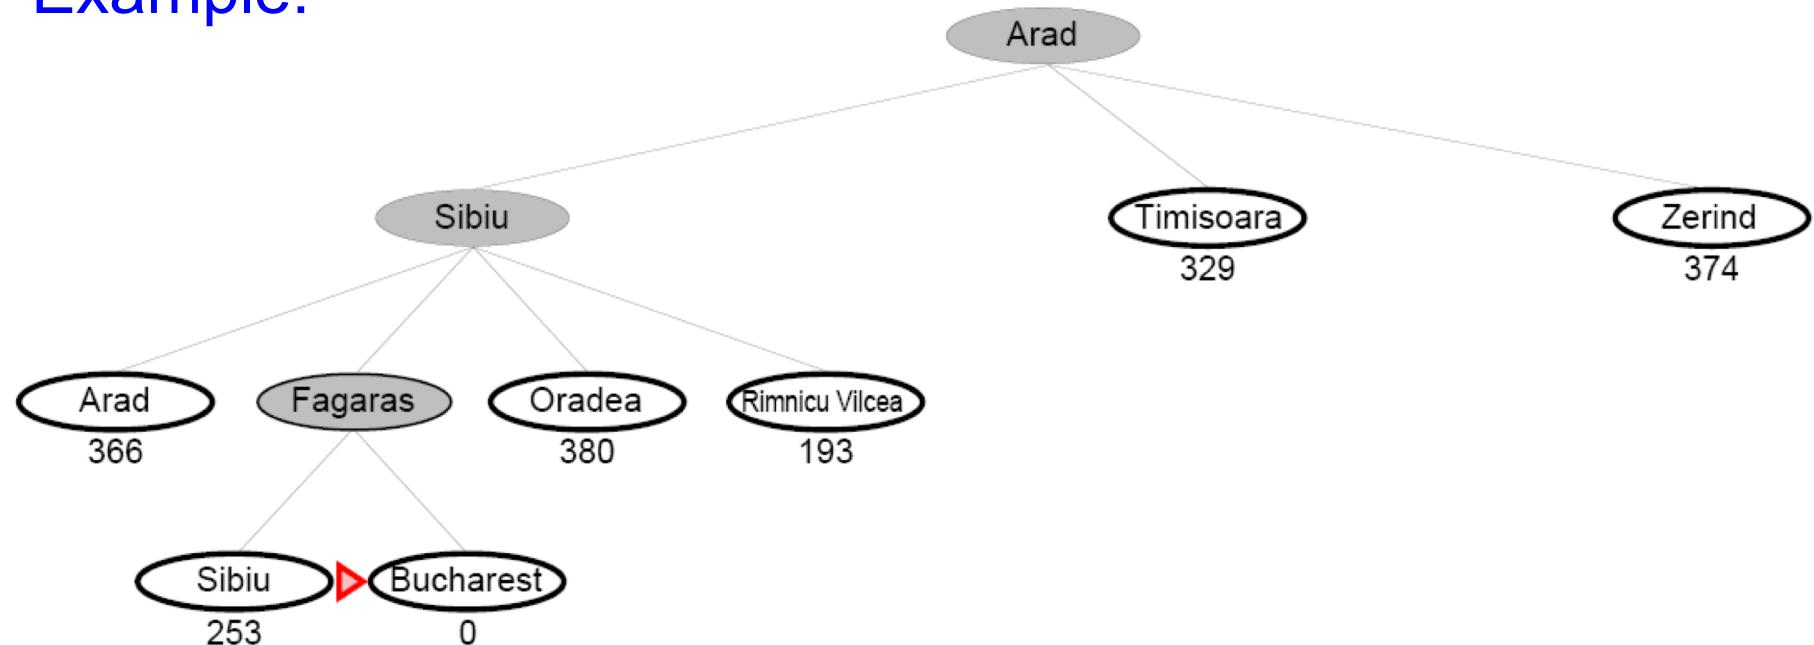
Greedy Best-First Search

- Evaluation function $f(n) = h(n)$ (*heuristic*)
 - estimates the cost from node n to *goal*
 - e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal
 - according to evaluation function
- **Example:**



Greedy Best-First Search

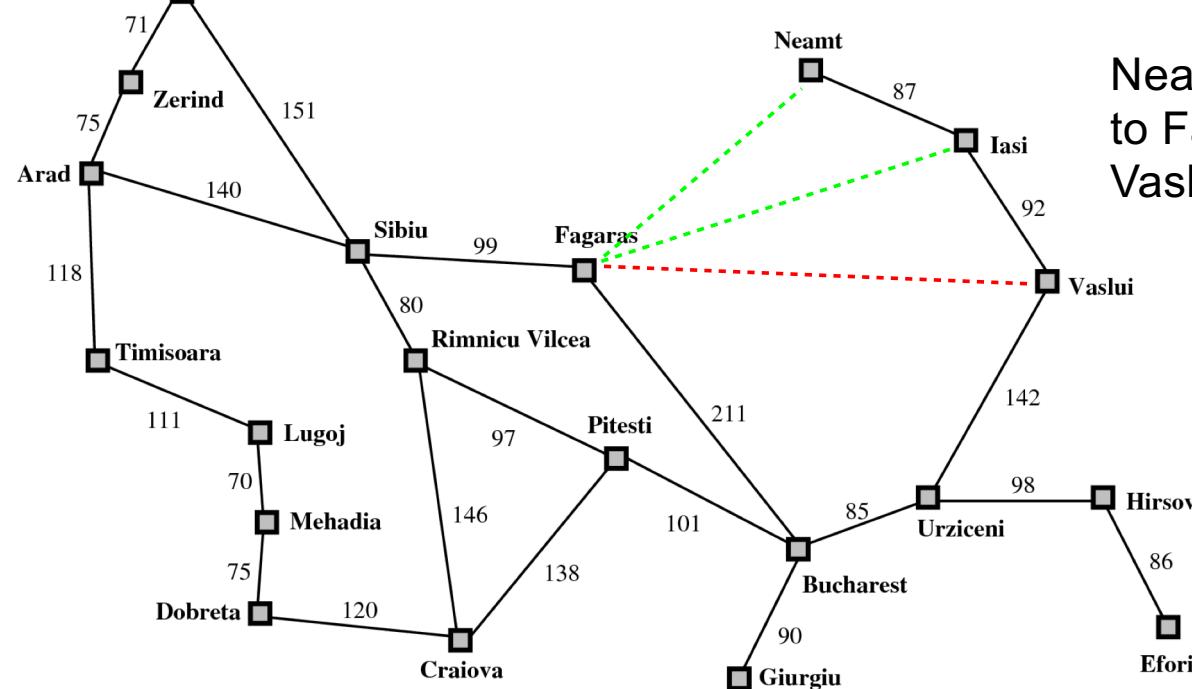
- Evaluation function $f(n) = h(n)$ (heuristic)
 - estimates the cost from node n to *goal*
 - e.g., $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal
 - according to evaluation function
- Example:



Properties of Greedy Best-First Search

- Completeness
 - No – can get stuck in loops
 - Example: We want to get from lasi to Fagaras
 - lasi → Neamt → lasi → Neamt → ...

Note:
These two are
different search
nodes referring
to the same state!



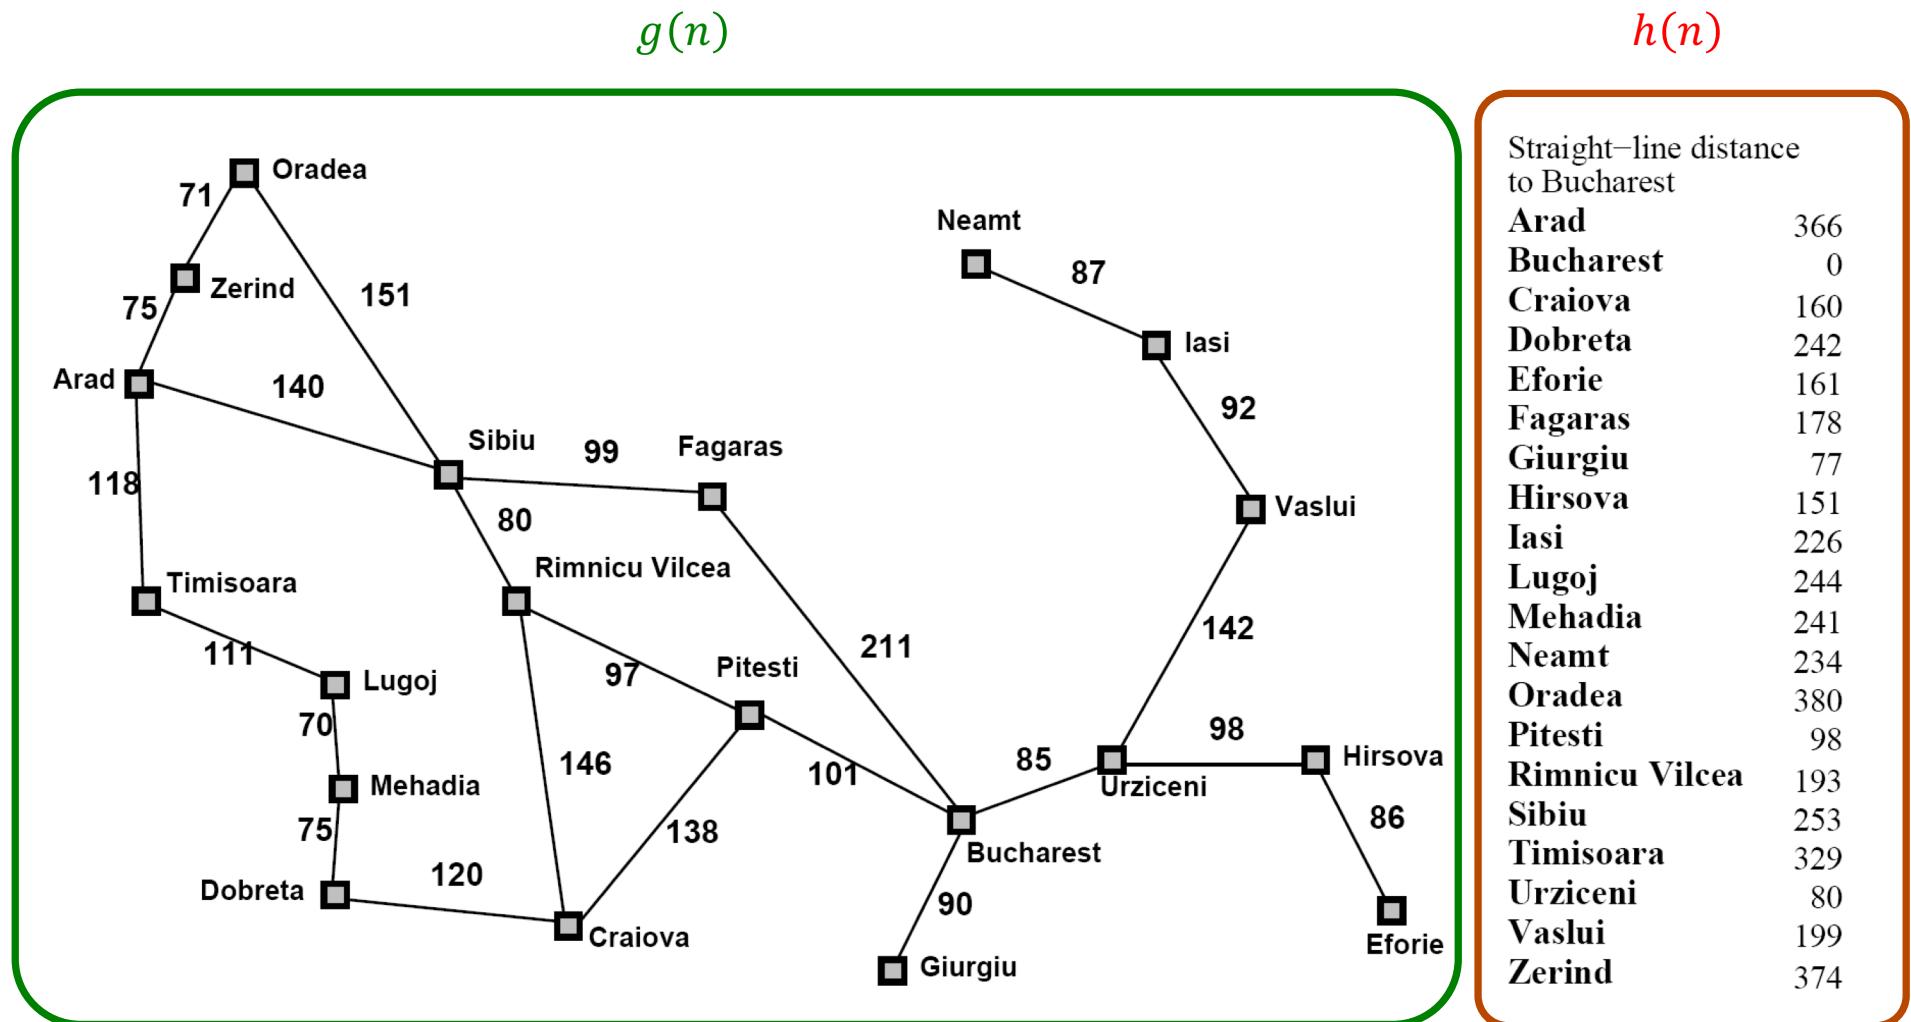
Properties of Greedy Best-First Search

- **Completeness**
 - No – can get stuck in loops
 - can be fixed with careful checking for duplicate states
→ **complete** in finite state space with repeated-state checking
- **Time Complexity**
 - $O(b^m)$, like depth-first search
 - but a good heuristic can give dramatic improvement
 - optimal case: best choice in each step → only d steps
 - a good heuristic improves chances for encountering optimal case
- **Space Complexity**
 - has to keep all nodes in memory → same as time complexity
- **Optimality**
 - **No**
 - **Example:**
 - solution Arad → Sibiu → Fagaras → Bucharest is not optimal

A* Search

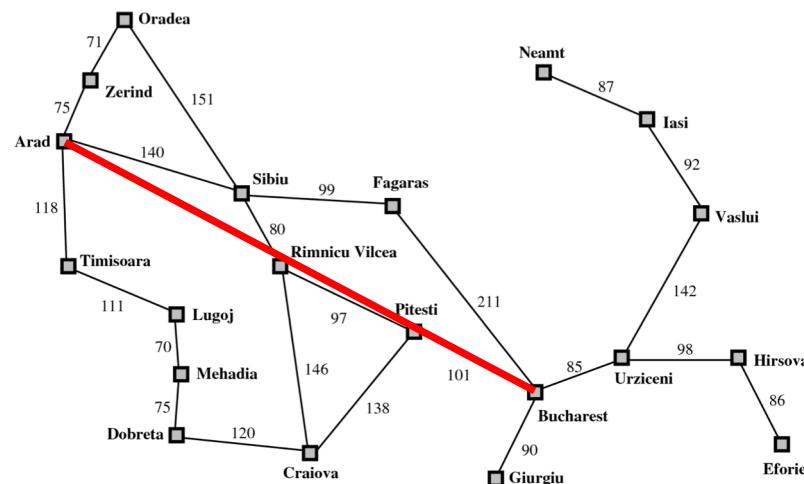
- Best-known form of best-first search
- Basic idea:
 - avoid expanding paths that are already expensive
→ evaluate complete path cost not only remaining costs
- Evaluation function: $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach node n
 - $h(n)$ = estimated cost to get from n to goal
 - $f(n)$ = estimated cost of path to goal via n

Beispiel

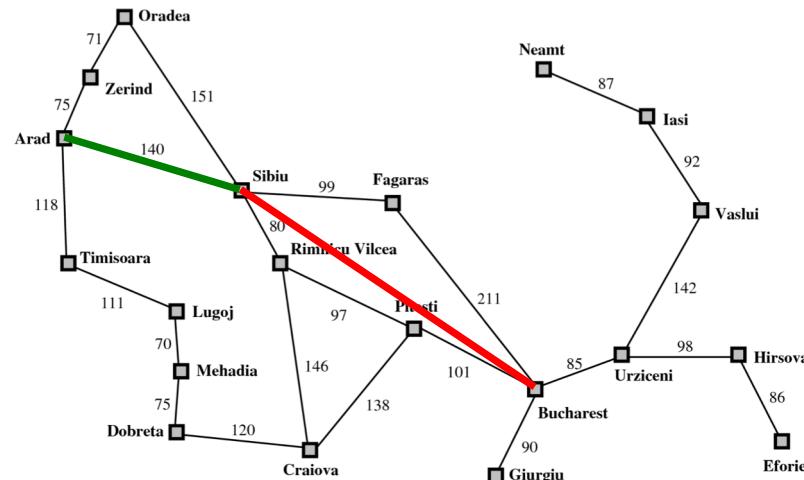
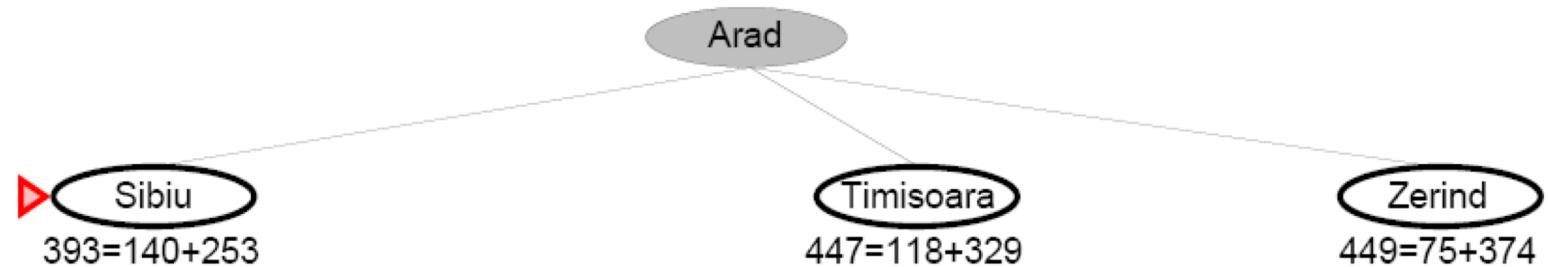


A* Search Example

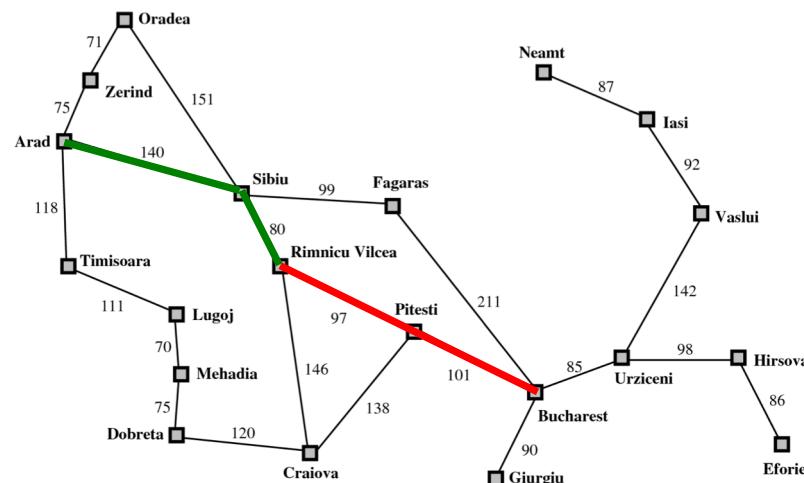
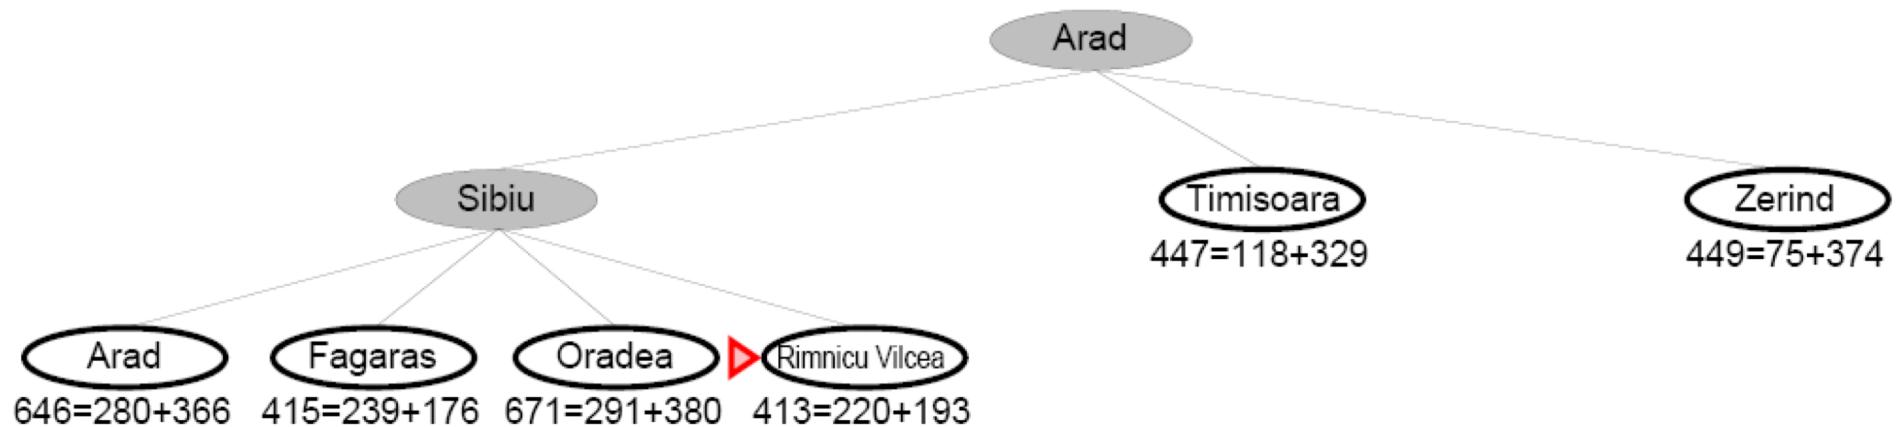
► Arad
 $366=0+366$



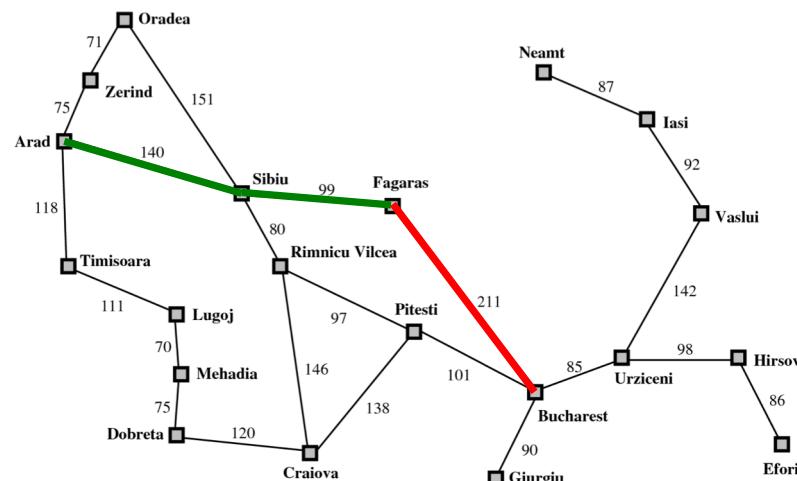
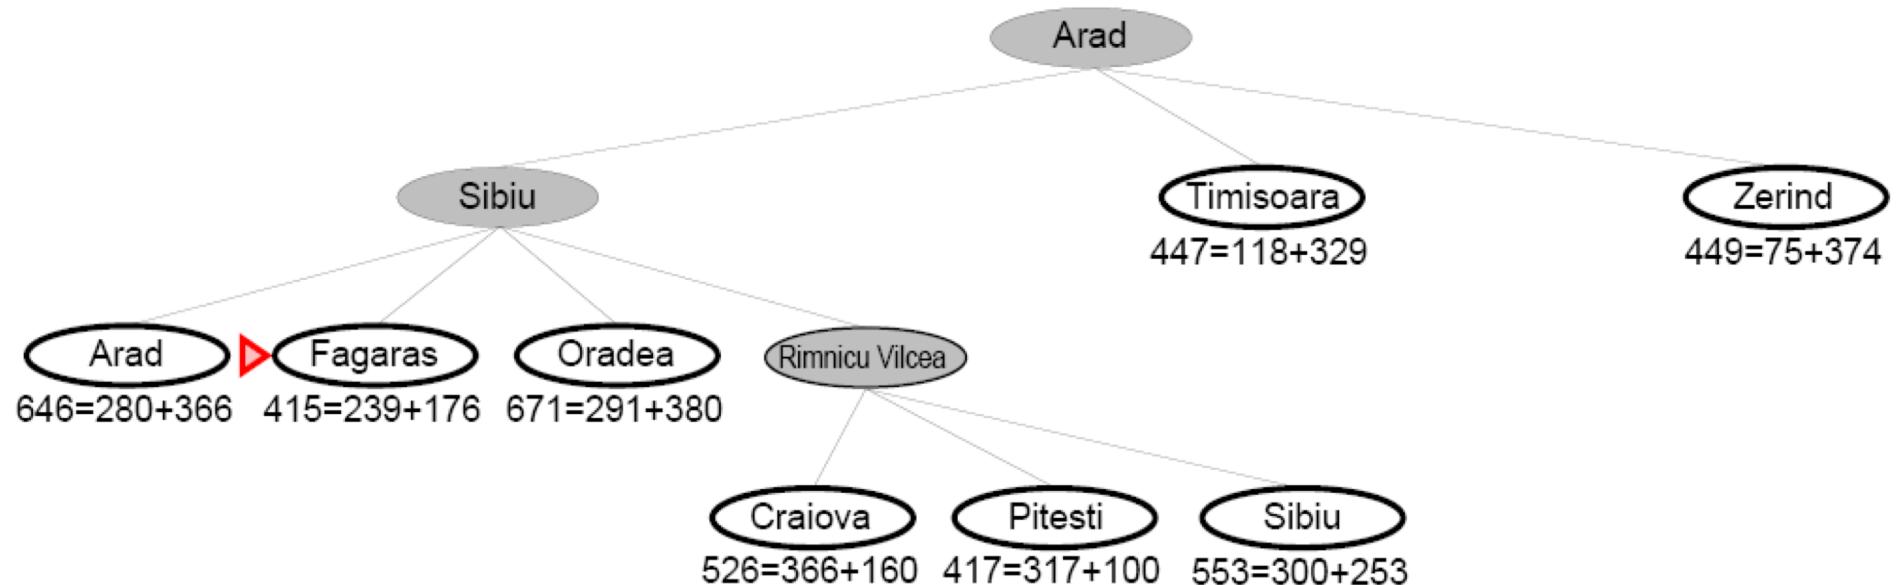
A* Search Example



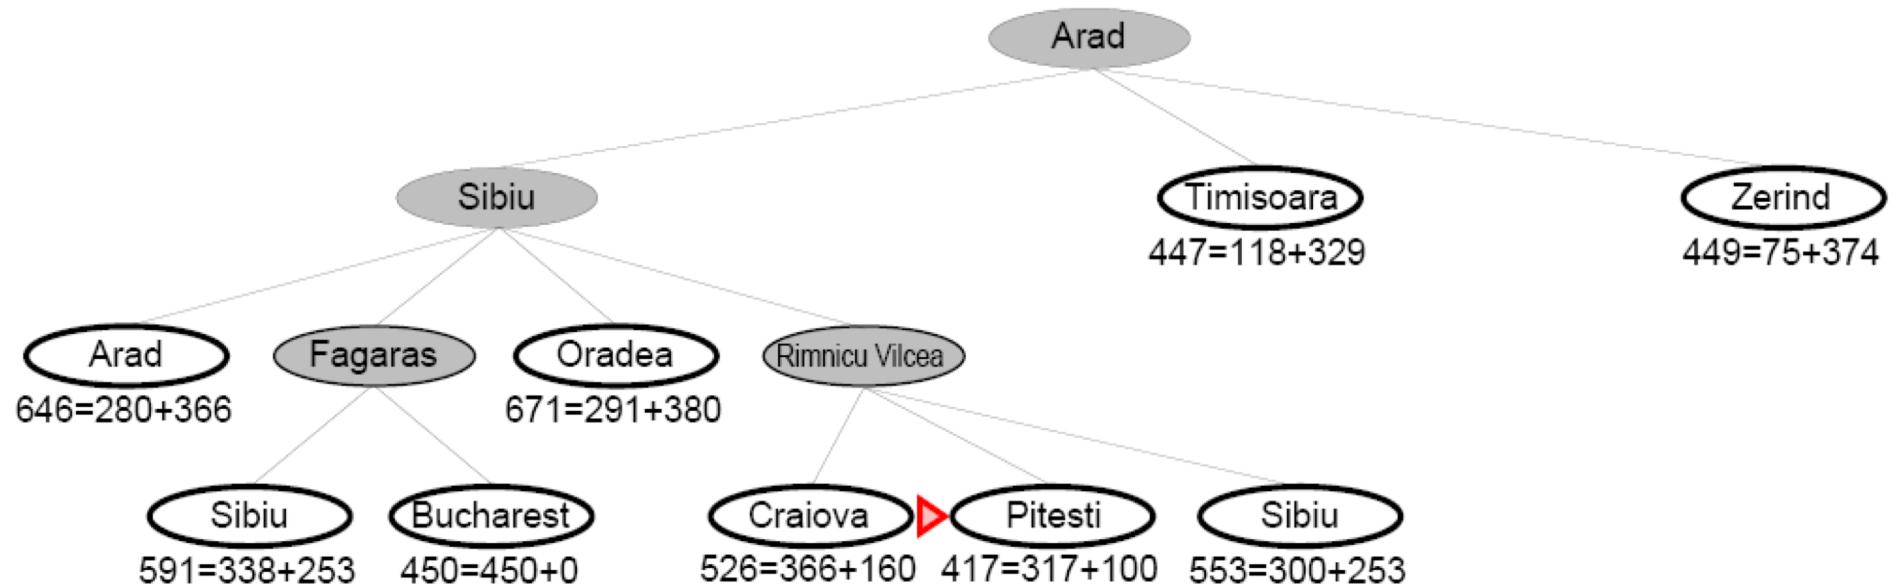
A* Search Example



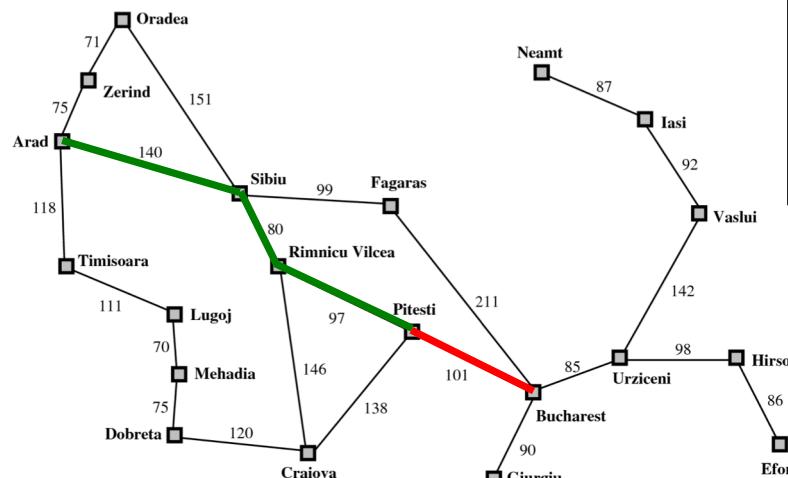
A* Search Example



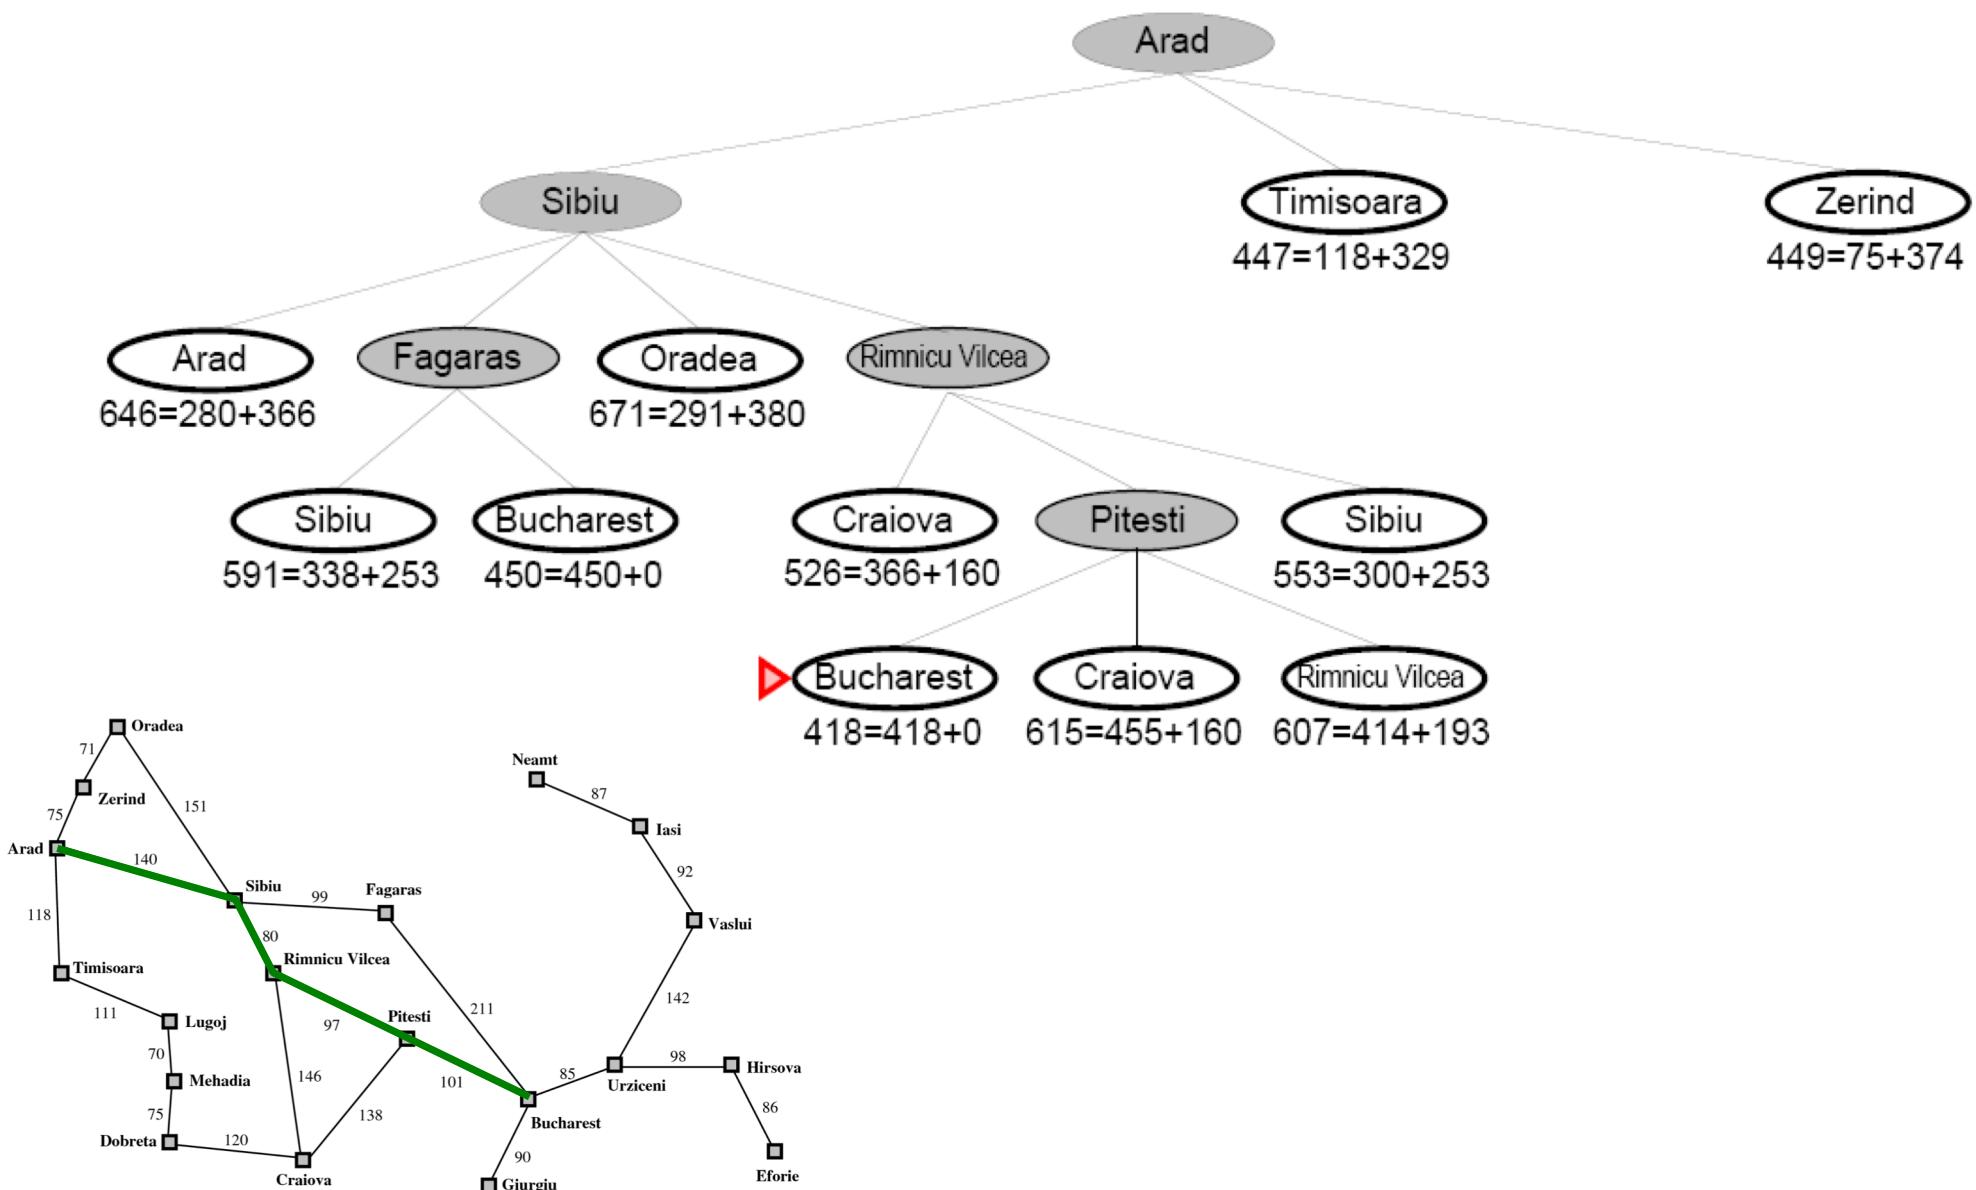
A* Search Example



Note that Pitesti will be expanded even though Bucharest is already in the fringe!
This is good, because we may still find a shorter way to Budapest.
Greedy Search would not do that.



A* Search Example



Properties of A*

- Completeness
 - Yes
 - unless there are infinitely many nodes with $f(n) \leq f(G)$
 - Time Complexity
 - it can be shown that the number of nodes grows exponentially unless the error of the heuristic $h(n)$ is bounded by the logarithm of the value of the actual path cost $h^*(n)$, i.e.
$$|h(n) - h^*(n)| \leq o(\log h^*(n))$$
 - Space Complexity
 - keeps all nodes in memory
 - typically the main problem with A*
 - Optimality
 - ???
- following pages

Admissible Heuristics

A heuristic is **admissible** if it *never overestimates* the cost to reach the goal

- Formally:
 - $h(n) \leq h^*(n)$ if $h^*(n)$ are the true cost from n to goal
- Example:
 - Straight-Line Distances h_{SLD} are an admissible heuristics for actual road distances h^*
- Note:
 - $h(n) \geq 0$ must also hold, so that $h(goal) = 0$

Theorem

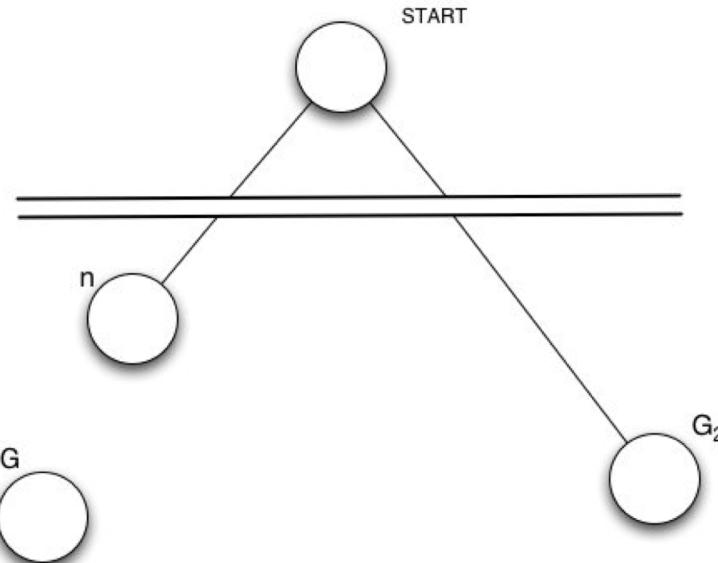
If $h(n)$ is admissible, A* using TREE-SEARCH is optimal.

Proof:

Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G with path cost $C^* = g(G)$

$$\begin{aligned} C^* &= g(n) + h^*(n) \\ f(n) &= g(n) + h(n) \\ h(n) &\leq h^*(n) \end{aligned}$$

because h admissible



Suppose some suboptimal goal G_2 has been generated and is in the fringe.

$$\begin{aligned} g(G_2) &> C^* \\ &\text{because } G_2 \text{ suboptimal} \\ f(G_2) &= g(G_2) \\ &\text{because } h(G_2) = 0 \\ &\text{(holds for all goal nodes)} \end{aligned}$$

$f(n) \leq C^* < f(G_2)$

G_2 will never be expanded, and G will be returned

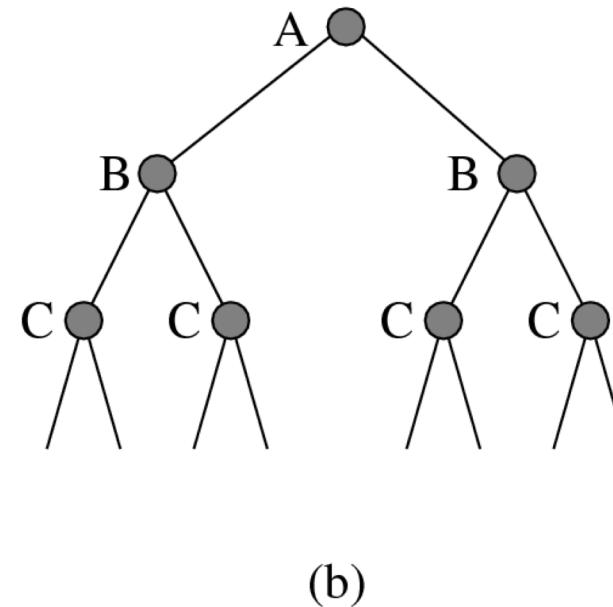
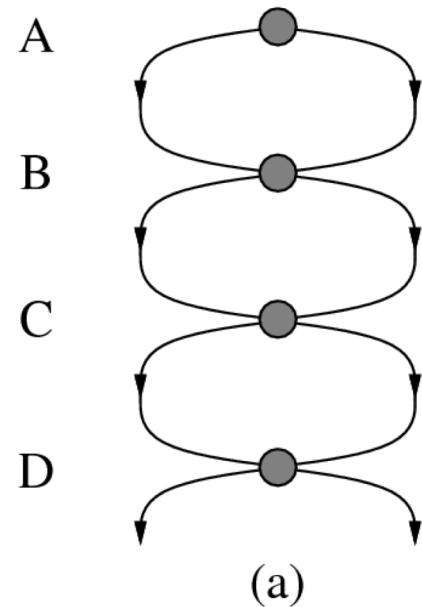
A* and Graph Search

- So far we only had tree search where each node in the search tree represents one possible path to the associated domain state
 - e.g., one can go directly from Arad to Sibiu or via Oradea
- Problem:
 - In some cases the path that is detected later may be the better path
 - so that a previously found solution starting from Sibiu has to re-investigated with the new, cheaper path
- Two solutions
 - Add ability to detect repeated states
 - → graph search
 - general solution that also works for BFS, DFS, ...
 - Or ensure that the cheaper path is always taken first
 - → consistent heuristics
 - specific solution for A*

Repeated States

- Failure to detect repeated states can turn a linear problem into an exponential one!

Ribbon Example

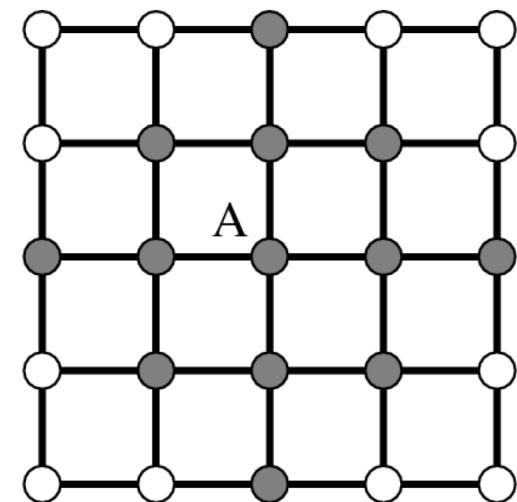


Repeated States

- Failure to detect repeated states can turn a linear problem into an exponential one!

(more realistic) Grid Example

- each square on grid has 4 neighboring states in
- thus, game tree w/o repetitions has 4^d nodes
- but only about $2d^2$ different states are reachable in d steps



Graph Search

- remembers the states that have been visited in a list *closed*
 - Note: the fringe list is often also called the **open list**

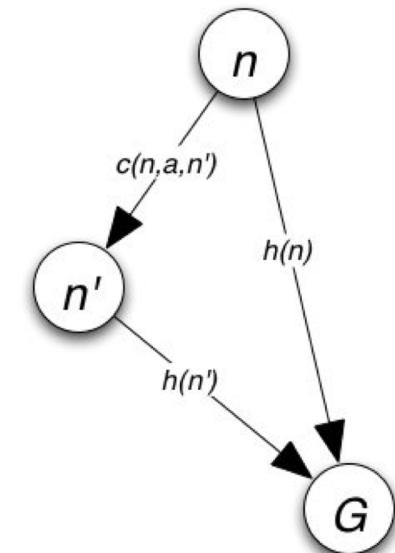
```
function GRAPH-SEARCH( problem, fringe) returns a solution, or failure
    closed  $\leftarrow$  an empty set
    fringe  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
    loop do
        if fringe is empty then return failure
        node  $\leftarrow$  REMOVE-FRONT(fringe)
        if GOAL-TEST(problem, STATE[node]) then return node
        if STATE[node] is not in closed then
            add STATE[node] to closed
            fringe  $\leftarrow$  INSERTALL(EXPAND(node, problem), fringe)
    end
```

- Example:
 - Dijkstra's algorithm is the graph-search variant of uniform cost search

Consistent Heuristics

- Graph-Search discards new paths to repeated state even though the new path may be cheaper
→ Previous proof breaks down
- 2 Solutions
 1. Add extra bookkeeping to remove the more expensive path
 2. Ensure that optimal path to any repeated state is always followed first
- Requirement for Solution 2:

A heuristic is **consistent** if for every node n and every successor n' generated by any action a it holds that $h(n) \leq c(n, a, n') + h(n')$



Lemma 1

Every consistent heuristic is admissible.

Proof Sketch by induction

Base Case: for all nodes n , in which an action a leads to goal G

$$h(n) \leq c(n, a, G) + h(G) = h^*(n)$$

By *induction on the path length from goal*, this argument can be extended to all nodes, so that eventually

$$\forall n: h(n) \leq h^*(n)$$

-
- Note:
 - not every admissible heuristic is consistent
 - but most of them are
 - it is hard to find non-consistent admissible heuristics

Lemma 2

If $h(n)$ is **consistent**, then the values of $f(n)$ along any path are **non-decreasing**.

Proof:

$$\begin{aligned} f(n) &= g(n) + h(n) \leq g(n) + c(n, a, n') + h(n') = \\ &g(n) + c(n, a, n') + h(n') = g(n') + h(n') = f(n') \end{aligned}$$

Theorem

If $h(n)$ is **consistent**, A* is optimal.

Proof:

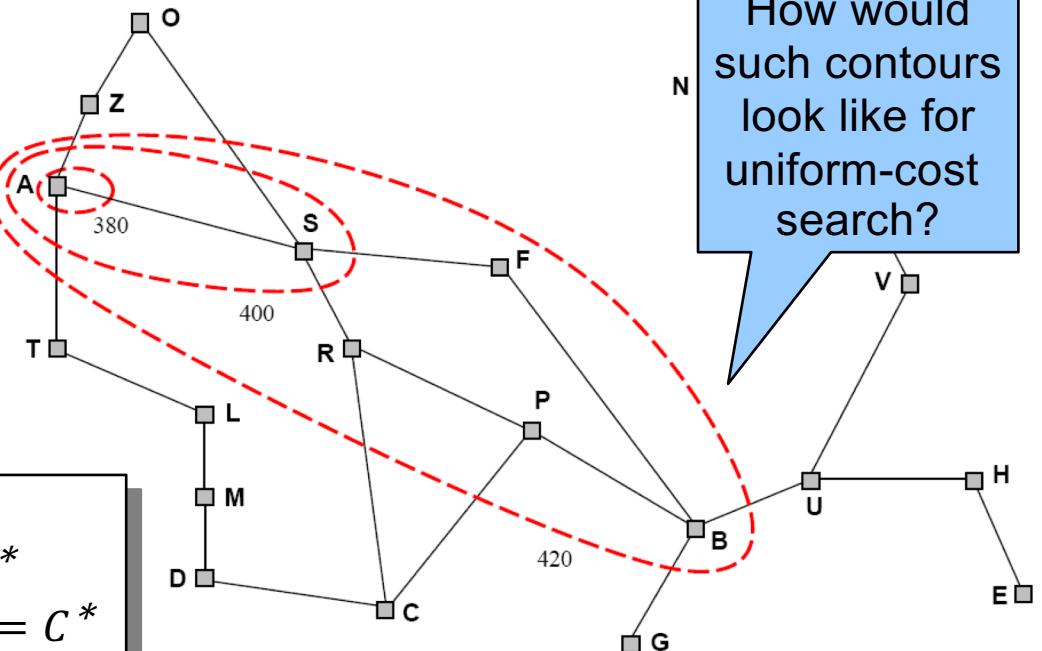
A* expands nodes in order of increasing f value

Contour labelled f_i
contains all nodes
with $f(n) < f_i$

Contours expand gradually
Cannot expand f_{i+1} until f_i is finished.

Eventually

- A* expands **all nodes with** $f(n) < C^*$
- A* expands **some nodes with** $f(n) = C^*$
- A* expands **no nodes with** $f(n) > C^*$



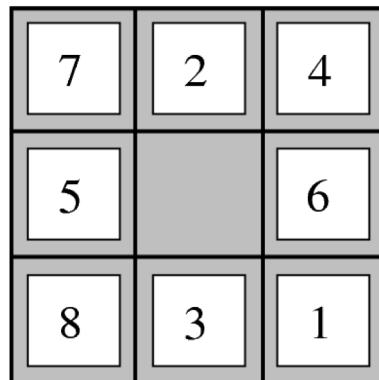
Memory-Bounded Heuristic Search

- Some **solutions** to A* space problems
(maintaining completeness and optimality)

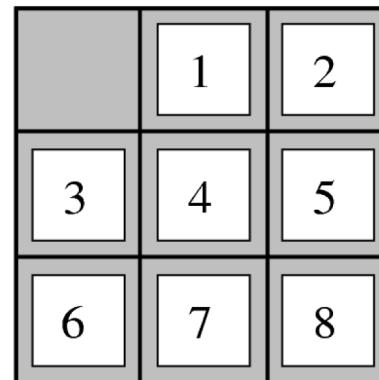
- Iterative-deepening A* (IDA*)
 - like iterative deepening
 - cutoff information is the f -cost ($g + h$) instead of depth
- Recursive best-first search (RBFS)
 - recursive algorithm that attempts to mimic
 - standard best-first search with linear space.
 - keeps track of the f -value of the best alternative
 - path available from any ancestor of current node
 - heuristic evaluations are updated with results of successors
- (Simple) Memory-bounded A* ((S)MA*)
 - drop the worst leaf node when memory is full
 - its value will be updated to its parent
 - May need to be re-searched later

Admissible Heuristics: 8-Puzzle

- $h_{MIS}(n)$ = number of misplaced tiles
 - admissible because each misplaced tile must be moved at least once
- $h_{MAN}(n)$ = total Manhattan distance
 - i.e., no. of squares from desired location of each tile
 - admissible because this is the minimum distance of each tile to its target square
- Example:



Start State



Goal State

$$h_{MIS}(\text{start}) = 8$$

$$h_{MAN}(\text{start}) = 18$$

$$h^*(\text{start}) = 26$$

Effective Branching Factor

- Evaluation Measure for a search algorithm:
 - assume we searched N nodes and found solution in depth d
 - the effective branching factor b^* is the branching factor of a uniform tree of depth d with $N+1$ nodes, i.e.

$$1 + N = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

- Measure is fairly constant for different instances of sufficiently hard problems
 - Can thus provide a good guide to the heuristic's overall usefulness.
 - A good value of b^* is 1

Efficiency of A* Search

- Comparison of number of nodes searched by A* and Iterative Deepening Search (IDS)
 - average of 100 different 8-puzzles with different solutions
 - Note:** heuristic $h_2 = h_{MAN}$ is always better than $h_1 = h_{MIS}$

| d | Suchkosten | | | Effektiver Verzweigungsfaktor | | |
|-----|------------|------------|------------|-------------------------------|------------|------------|
| | IDS | $A^*(h_1)$ | $A^*(h_2)$ | IDS | $A^*(h_1)$ | $A^*(h_2)$ |
| 2 | 10 | 6 | 6 | 2,45 | 1,79 | 1,79 |
| 4 | 112 | 13 | 12 | 2,87 | 1,48 | 1,45 |
| 6 | 680 | 20 | 18 | 2,73 | 1,34 | 1,30 |
| 8 | 6384 | 39 | 25 | 2,80 | 1,33 | 1,24 |
| 10 | 47127 | 93 | 39 | 2,79 | 1,38 | 1,22 |
| 12 | 3644035 | 227 | 73 | 2,78 | 1,42 | 1,24 |
| 14 | — | 539 | 113 | — | 1,44 | 1,23 |
| 16 | — | 1301 | 211 | — | 1,45 | 1,25 |
| 18 | — | 3056 | 363 | — | 1,46 | 1,26 |
| 20 | — | 7276 | 676 | — | 1,47 | 1,27 |
| 22 | — | 18094 | 1219 | — | 1,48 | 1,28 |
| 24 | — | 39135 | 1641 | — | 1,48 | 1,26 |

Dominance

If h_1 and h_2 are admissible, h_2 **dominates** h_1 if $\forall n: h_2(n) \geq h_1(n)$

- if h_2 dominates h_1 it will perform better because it will *always* be closer to the optimal heuristic h^*
- **Example:**
 - h_{MAN} dominates h_{MIS} because if a tile is misplaced, its Manhattan distance is ≥ 1

Theorem: (Combining admissible heuristics)

If h_1 and h_2 are two admissible heuristics than

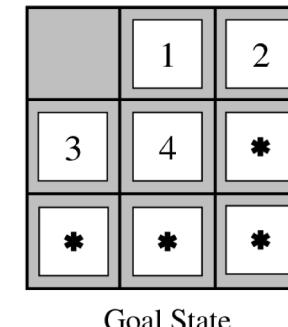
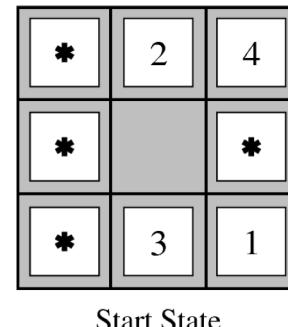
$h(n) = \max(h_1(n), h_2(n))$
is also admissible and dominates h_1 and h_2

Relaxed Problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- Examples:
 - If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then h_{MIS} gives the shortest solution
 - If the rules are relaxed so that a tile can move to **any adjacent square**, then h_{MAN} gives the shortest solution
- Thus, looking for relaxed problems is a good strategy for **inventing admissible heuristics**.

Pattern Databases

- Admissible heuristics can also be derived from the solution cost of a subproblem of a given problem.
 - This cost is a lower bound on the cost of the real problem.
- Pattern databases store the **exact solution** (length) for every possible **subproblem** instance
 - constructed once for all by searching backwards from the goal and recording every possible pattern
- **Example:**
 - store exact solution costs for solving 4 tiles of the 8-puzzle
 - sample pattern:



Learning of Heuristics

- Another way to find a heuristic is through learning from experience
- Experience:
 - states encountered when solving lots of 8-puzzles
 - states are encoded using features, so that similarities between states can be recognized
- Features:
 - for the 8-puzzle, features could, e.g. be
 - the number of misplaced tiles
 - number of pairs of adjacent tiles that are also adjacent in goal
 - ...
- An inductive learning algorithm can then be used to predict costs for other states that arise during search.
- No guarantee that the learned function is admissible!

Learning of Heuristics

- Another way to find a heuristic is through learning from experience
- Experience:
 - states encountered when solving lots of 8-puzzles
 - states are encoded using features, so that similarities between states can be recognized
- Features:
 - for the 8-puzzle, features could, e.g. be
 - the number of misplaced tiles
 - number of pairs of adjacent tiles that are also adjacent in goal
 - ...
- An inductive learning algorithm can then be used to predict costs for other states that arise during search.
- No guarantee that the learned function is admissible!

Summary

- Heuristic functions estimate the costs of shortest paths
- Good heuristics can dramatically reduce search costs
- Greedy best-first search expands node with lowest estimated remaining cost
 - incomplete and not always optimal
- A* search minimizes the path costs so far plus the estimated remaining cost
 - complete and optimal, also **optimally efficient**:
 - no other search algorithm can be more efficient, because they all have search the nodes with $f(n) < c^*$
 - otherwise it could miss a solution
- Admissible search heuristics can be derived from exact solutions of reduced problems
 - problems with less constraints
 - subproblems of the original problem