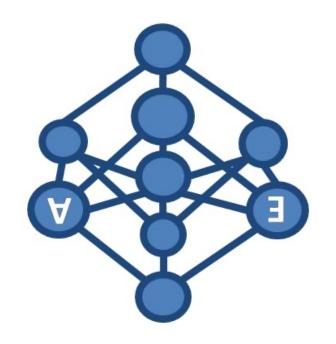# Probabilistic Graphical Models*
Dynamic Bayesian Networks

*Thanks to Carlos Guestrin, Pedro Domingos and many others for making their slides publically available

So far

- Bayesian networks

- Inference within Bayesian networks

- Learning Bayesian networks

- **All rather static situations …**

What's next

- The world is not static but evolves over time
    - Monitoring freeway traffic (from an autonomous driver or for traffic analysis)
    - Monitoring patient's symptoms
- **How can we model dynamic environments?**

- Hidden Markov Models
- (Bayes Filter and Kalman Filters)
- Dynamic Bayesian Networks

Hidden Markov Models

Let's consider an example:

Part-of-speech tagging

Kristian Kersting - Probabilistic Graphical Models

Example 1: Parts of Speech

- **Part of speech (POS):** Noun, verb, adjective, preposition, adverb, article, interjection, pronoun, conjunction, etc.

- **POS tagging**: The process of assigning (labeling) a part-of-speech or other lexical class marker to each word in a sentence (or a corpus)

  - Decide whether each word is a noun, verb, adjective, …

- **Also called word-category disambiguation**

Kristian Kersting - Probabilistic Graphical Models

POS Tagging Statistical Approaches: Idea 1

Simply assign each word its most likely POS.

Success rate: 91%!

Word	POS listings in Brown		
heat	noun/89	verb/5	
oil	noun/87		
in	prep/20731	noun/1	adv/462
a	det/22943	noun/50	noun-proper/30
large	adj/354	noun/2	adv/5
pot	noun/27		

Statistical Approaches: Idea 2

For a string of words

$$W = w_1 w_2 w_3 ... w_n$$

find the string of POS tags

$$T = t_1\, t_2\, t_3\, ...t_n$$

which maximizes P(T|W)

- i.e., the probability of tag string T given that the word string was W, i.e., that W was tagged T

The Sparse Data Problem …

A simple but unrealistic approach to compute P(T|W):

Count instances of the string "heat oil in a large pot" in the training corpus, and **pick the most common tag assignment** to the string.

So, what can we estimate ?

- What parameters can we estimate with a million words of hand tagged training data?

  - Assume a uniform distribution of 5000 words and 40 part of speech tags..

Event	Count	Estimate Quality?
tags	40	Excellent
tag bigrams	1600	Excellent
tag trigrams	64,000	OK
tag 4-grams	2.5M	Poor
words	5000	Very Good
word bigrams	25M	Poor
word x tag pairs	200,000	OK

- **Rich Models often require vast amounts of data**

- Good estimates of models with bad assumptions often outperform better models which are badly estimated (recall Naïve Bayes!)

Kristian Kersting - Probabilistic Graphical Models

A Practical Statistical Tagger

By Bayes' Rule:

$$\mathbf{P(T|W)} = \frac{P(W|T) * P(T)}{P(W)}$$

so to maximize $P(T|W)$, need to maximize $P(W|T) * P(T)$.

To compute $\mathbf{P(T)}$: By the chain rule,

$$P(T) = P(t_1) * P(t_2|t_1) * P(t_3|t_1 t_2) * \ldots * P(t_n|t_1 \ldots t_{n-1})$$

A Practical Statistical Tagger II

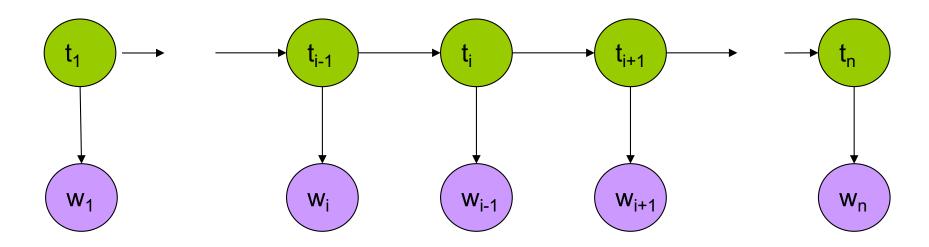But we cannot accurately estimate more than tag bigrams! Thus, we simply change to a model that we CAN estimate:

A Markov Assumption: $P(t_i | t_1 \ldots t_{i-1}) = P(t_i | t_{i-1})$

By which

$$P(T) = P(t_1) * P(t_2 | t_1) * P(t_3 | t_2) * \ldots * P(t_n | t_{n-1})$$

A Graphical Representation

$$P(T) * P(W|T) \; = \;$$

$$P(t_1) * P(t_2|t_1) * P(t_3|t_2) * \ldots * P(t_n|t_{n-1}) *$$
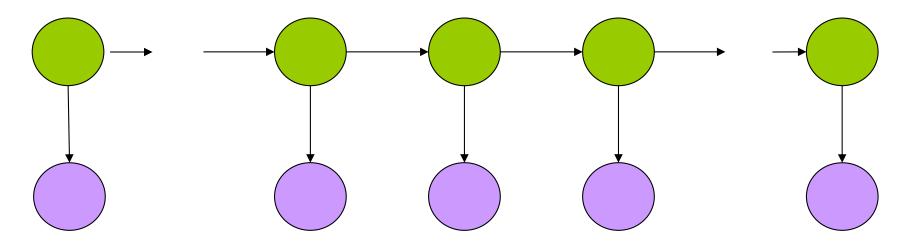$$P(w_1|t_1) * P(w_2|t_2) * \ldots * P(w_n|t_n)$$

Kristian Kersting - Probabilistic Graphical Models

- This is an example of a Hidden Markov Model (HMM)

- So, what are HMMs?

Kristian Kersting - Probabilistic Graphical Models

What is an HMM?
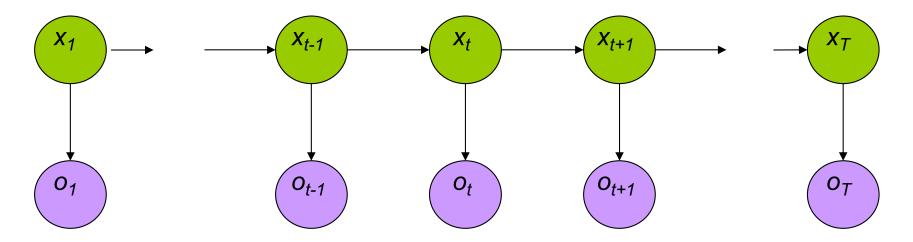
- Graphical Model
- Circles indicate states
- Arrows indicate probabilistic dependencies between states

HMM = Hidden Markov Model
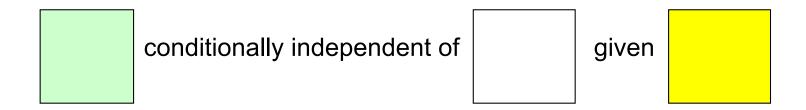
HMMs and Bayesian Nets
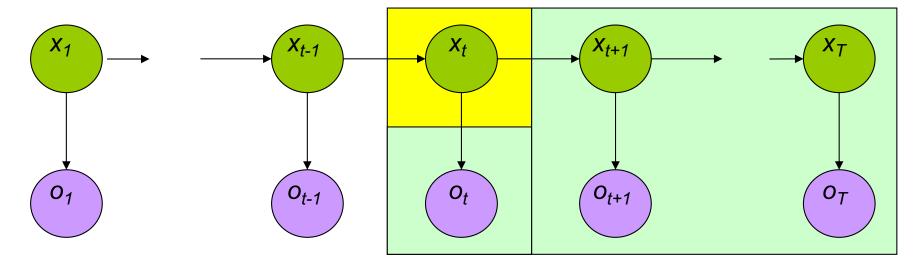
$$P(x_1...x_T, o_1...o_T) = P(x_1)P(o_1 \mid x_1)\prod_{i=1}^{T-1} P(x_{i+1} \mid x_i).P(o_{i+1} \mid x_{i+1})$$

$$= \pi_{x_1} b_{x_1 o_1} \prod_{t=1}^{T-1} a_{x_t x_{t+1}} b_{x_{t+1} o_{t+1}}$$
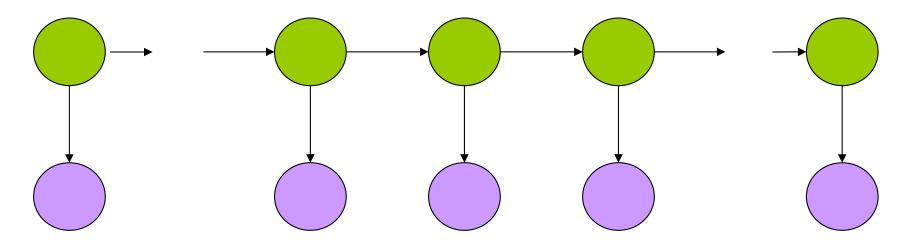
Independence Assumption

conditionally independent of □ given ▨

Because of d-separation "The past is independent of the future given the present."
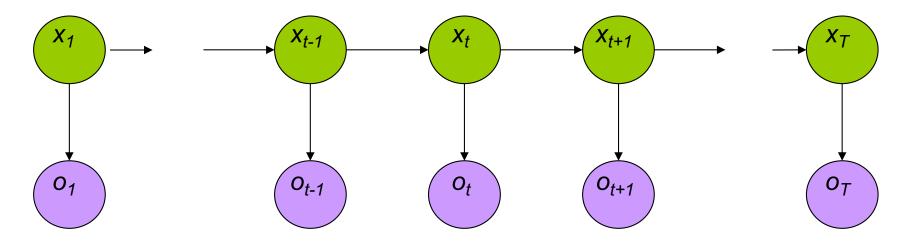
Kristian Kersting - Probabilistic Graphical Models

Inference in an HMM

1. Compute the probability of a given observation sequence

2. Given an observation sequence, compute the most likely hidden state sequence

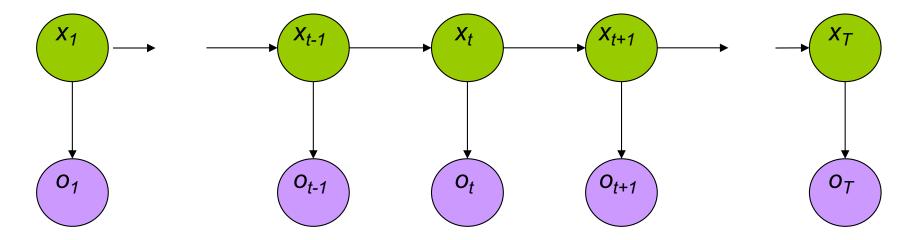3. Given an observation sequence and set of possible models, which model most closely fits the data?

(1) Decoding

$$P(O \mid X, \mu) = b_{x_1 o_1} b_{x_2 o_2} ... b_{x_T o_T}$$

μ are the parameters

(1) Decoding

$$P(O \mid X, \mu) = b_{x_1 o_1} b_{x_2 o_2} ... b_{x_T o_T}$$

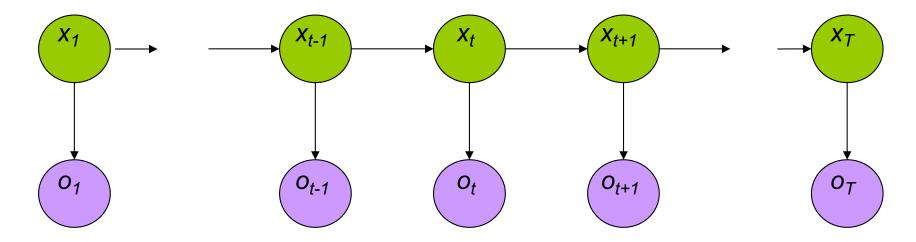$$P(X \mid \mu) = \pi_{x_1} a_{x_1 x_2} a_{x_2 x_3} ... a_{x_{T-1} x_T}$$
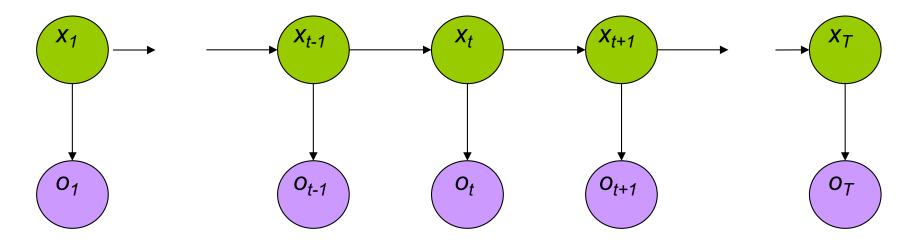
(1) Decoding

$$P(O \mid X, \mu) = b_{x_1 o_1} b_{x_2 o_2} ... b_{x_T o_T}$$

$$P(X \mid \mu) = \pi_{x_1} a_{x_1 x_2} a_{x_2 x_3} ... a_{x_{T-1} x_T}$$

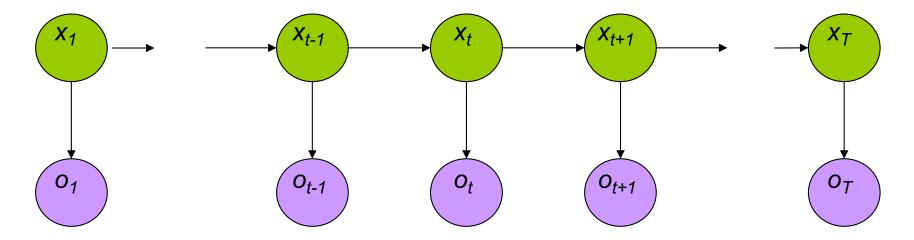$$P(O, X \mid \mu) = P(O \mid X, \mu) P(X \mid \mu)$$

(1) Decoding

$$P(O \mid X, \mu) = b_{x_1 o_1} b_{x_2 o_2} ... b_{x_T o_T}$$

$$P(X \mid \mu) = \pi_{x_1} a_{x_1 x_2} a_{x_2 x_3} ... a_{x_{T-1} x_T}$$

$$P(O, X \mid \mu) = P(O \mid X, \mu) P(X \mid \mu)$$
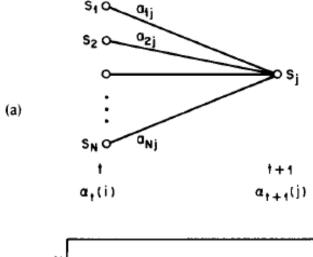
$$P(O \mid \mu) = \sum_X P(O \mid X, \mu) P(X \mid \mu)$$

Kristian Kersting — Probabilistic Graphical Models

(1) Decoding

$$P(O \mid \mu) = \sum_{X} \pi_{x_1} b_{x_1 o_1} \prod_{t=1}^{T-1} a_{x_t x_{t+1}} b_{x_{t+1} o_{t+1}}$$

Complexity $O(N^T . 2T)$

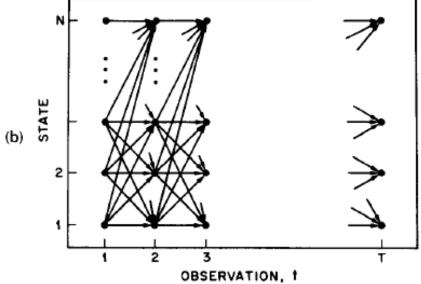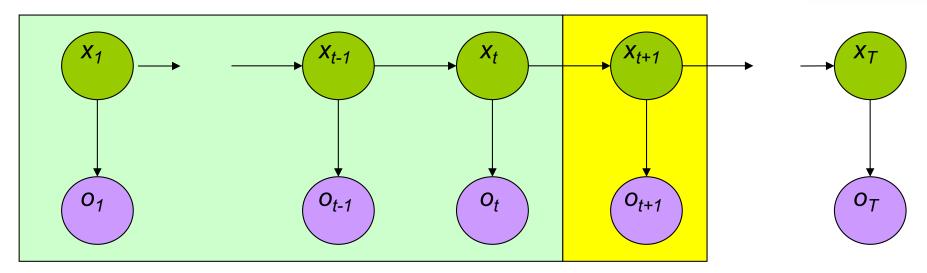$E.g. \, N = 5, T = 100$ gives $2.100.5^{100} \approx 10^{72}$

Fig. 4. (a) Illustration of the sequence of operations required for the computation of the forward variable $\alpha_{t+1}(j)$. (b) Implementation of the computation of $\alpha_t(i)$ in terms of a lattice of observations t, and states i.

Dynamic Programming to the rescue!

Forward Procedure
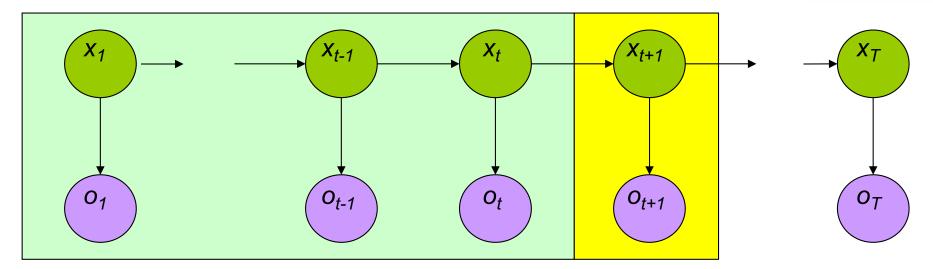
- Special structure gives us an efficient solution using *dynamic programming*

- **Intuition**: Probability of the first *t* observations is the same for all possible *t*+1 length state sequences

- **Define:**

$$\alpha_i(t) = P(o_1...o_t, x_t = i \mid \mu)$$

$$\alpha_i(1) = P(o_1, x_1 = i \mid \mu)$$

$$= \pi_i \cdot b_{io_1}$$

Forward Procedure

$$\alpha_j(t+1)$$

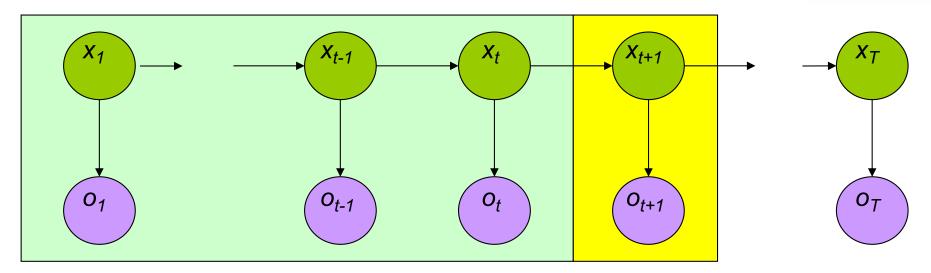$$= P(o_1...o_{t+1}, x_{t+1} = j) \quad \text{definition (μ omitted)}$$

$$= P(o_1...o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t \mid x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t, x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)$$

Forward Procedure

$$\alpha_j(t+1)$$

$$= P(o_1...o_{t+1}, x_{t+1} = j)$$

$$= P(o_1...o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j) \qquad \textbf{\color{red}chain rule}$$

$$= P(o_1...o_t \mid x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t, x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)$$

Forward Procedure

$$\alpha_j(t+1)$$

$$= P(o_1...o_{t+1}, x_{t+1} = j)$$
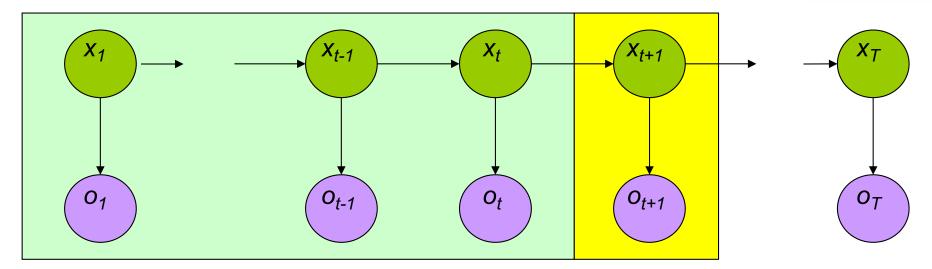
$$= P(o_1...o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t \mid x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t, x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j) \qquad \textcolor{red}{\textbf{chain rule}}$$

Forward Procedure

$$\alpha_j(t+1)$$

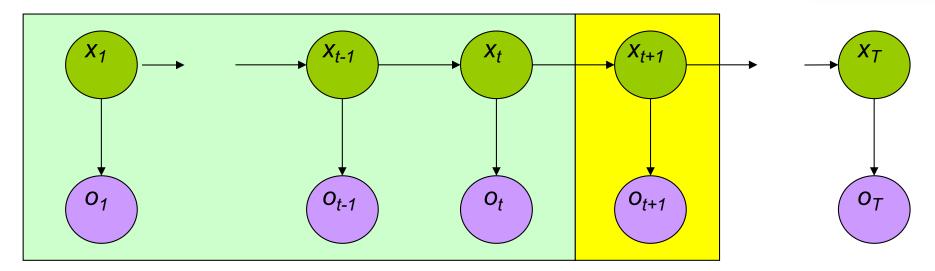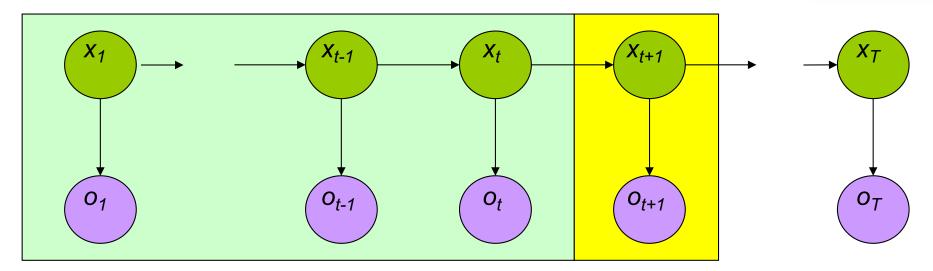$$= P(o_1...o_{t+1}, x_{t+1} = j)$$

$$= P(o_1...o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t \mid x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)P(x_{t+1} = j)$$

$$= P(o_1...o_t, x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)$$

chain rule

Forward Procedure
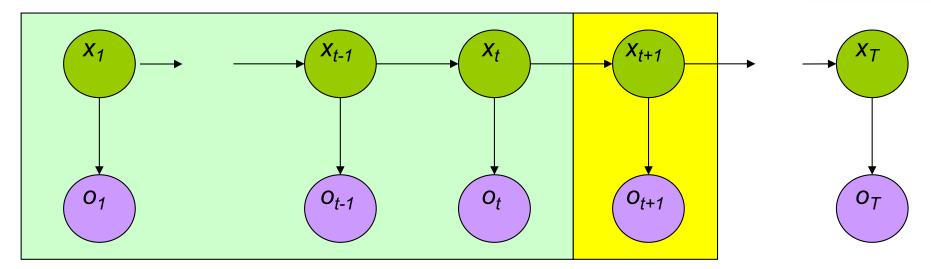
$$= \sum_{i=1\ldots N} P(o_1\ldots o_t, x_t = i, x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j) \quad \textbf{Total prob}$$

$$= \sum_{i=1\ldots N} P(o_1\ldots o_t, x_{t+1} = j \mid x_t = i)P(x_t = i)P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1\ldots N} P(o_1\ldots o_t, x_t = i)P(x_{t+1} = j \mid x_t = i)P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1\ldots N} \alpha_i(t)a_{ij}b_{jo_{t+1}}$$

29

Forward Procedure

$$= \sum_{i=1...N} P(o_1...o_t, x_t = i, x_{t+1} = j)P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} P(o_1...o_t, x_{t+1} = j \mid x_t = i)P(x_t = i)P(o_{t+1} \mid x_{t+1} = j)$$

Chain rule

$$= \sum_{i=1...N} P(o_1...o_t, x_t = i)P(x_{t+1} = j \mid x_t = i)P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} \alpha_i(t) a_{ij} b_{jo_{t+1}}$$
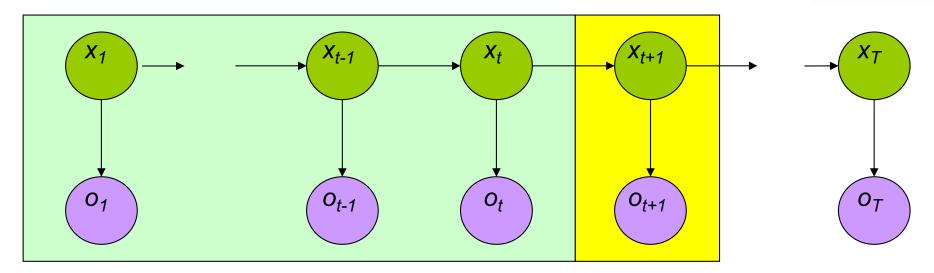
30

Forward Procedure

$$= \sum_{i=1...N} P(o_1...o_t, x_t = i, x_{t+1} = j) P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} P(o_1...o_t, x_{t+1} = j \mid x_t = i) P(x_t = i) P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} P(o_1...o_t, x_t = i) P(x_{t+1} = j \mid x_t = i) P(o_{t+1} \mid x_{t+1} = j)$$

2x Chain rule

$$= \sum_{i=1...N} \alpha_i(t) a_{ij} b_{jo_{t+1}}$$
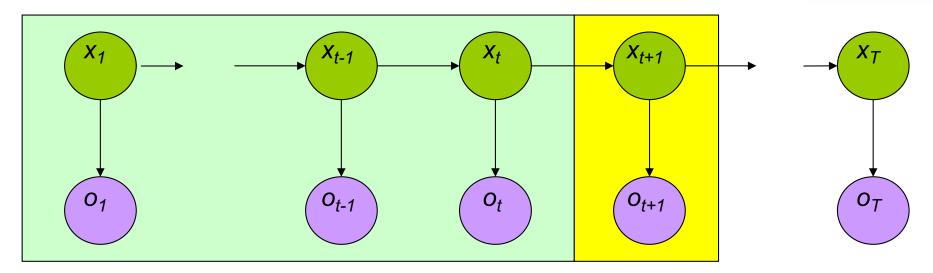
Forward Procedure
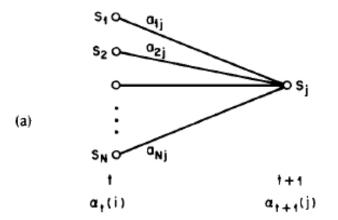
$$= \sum_{i=1...N} P(o_1...o_t, x_t = i, x_{t+1} = j) P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} P(o_1...o_t, x_{t+1} = j \mid x_t = i) P(x_t = i) P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} P(o_1...o_t, x_t = i) P(x_{t+1} = j \mid x_t = i) P(o_{t+1} \mid x_{t+1} = j)$$

$$= \sum_{i=1...N} \alpha_i(t) a_{ij} b_{jo_{t+1}} \qquad \textbf{Definition + given parameters}$$

32

(a)

(b)

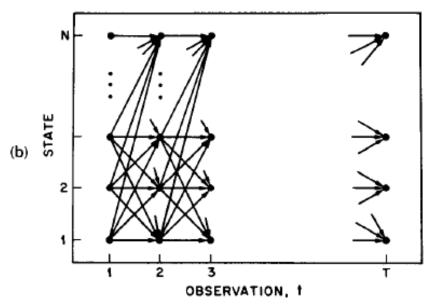Fig. 4. (a) Illustration of the sequence of operations required for the computation of the forward variable $\alpha_{t+1}(j)$. (b) Implementation of the computation of $\alpha_t(i)$ in terms of a lattice of observations t, and states i.

Dynamic Programming
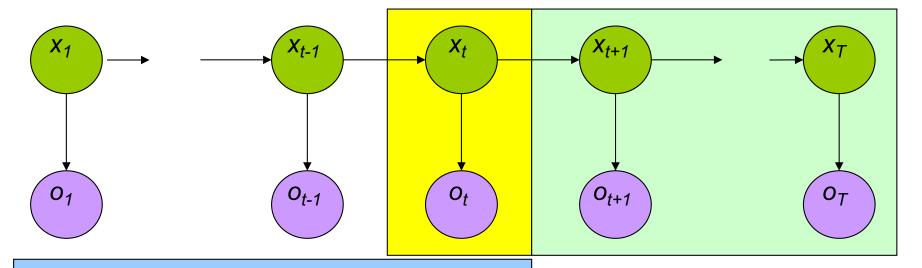
$$\alpha_j(t+1) = \sum_{i=1...N} \alpha_i(t) a_{ij} b_{jo_{t+1}}$$

Complexity $O(N^2 T)$

$E.g. \, N = 5, T = 100$ gives ≈ 3000

Similar, Backward Procedure

$$\beta_i(T) = 1$$

$$\beta_i(t) = P(o_{t+1}...o_T \mid x_t = i)$$

$$\beta_i(t) = \sum_{j=1...N} a_{ij} b_{io_{t+1}} \beta_j(t+1)$$

Probability of the rest of the states given the first state

Kristian Kersting - Probabilistic Graphical Models
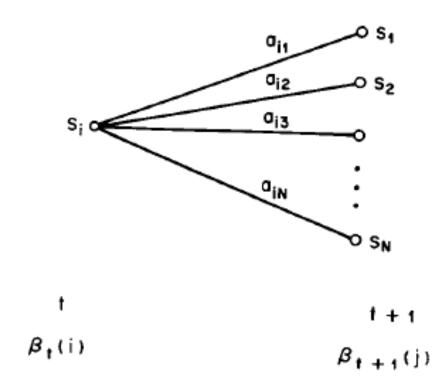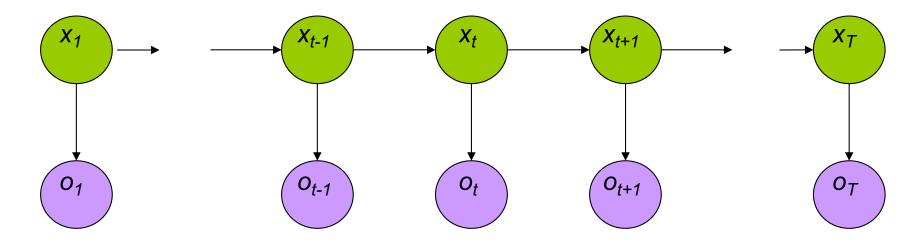
Fig. 5. Illustration of the sequence of operations required for the computation of the backward variable $\beta_t(i)$.

Kristian Kersting - Probabilistic Graphical Models

(1) Decoding: Solution

$$P(O \mid \mu) = \sum_{i=1}^{N} \alpha_i(T)$$ **Forward Procedure**

$$P(O \mid \mu) = \sum_{i=1}^{N} \pi_i \beta_i(1)$$ **Backward Procedure**
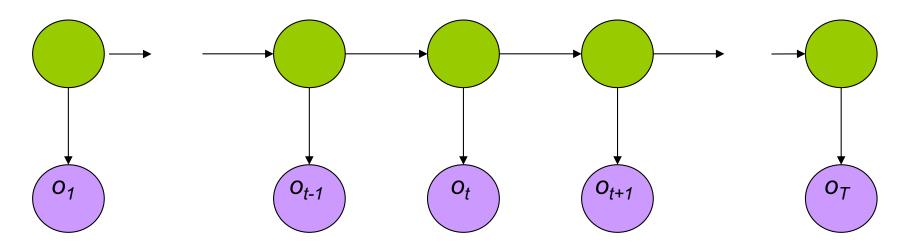
$$P(O \mid \mu) = \sum_{i=1}^{N} \alpha_i(t)\beta_i(t)$$ **Combination**
Important for (2) and (3)

(2) Best State Sequence

- **Find the state sequence that best explains the observations**

- **Two approaches**
  - **Individually most likely states**
  - **Most likely sequence (Viterbi)**

$$\arg\max_X P(X \mid O)$$

Best State Sequence (1)

$$\gamma_i(t) = P(X_t = i \mid O, \mu)$$

target

$$= \frac{P(X_t = i, O \mid \mu)}{P(O \mid \mu)}$$

definition of conditional prob

$$= \frac{\alpha_i(t).\beta_i(t)}{\sum_{j=1}^{n} \alpha_j(t).\beta_j(t)}$$
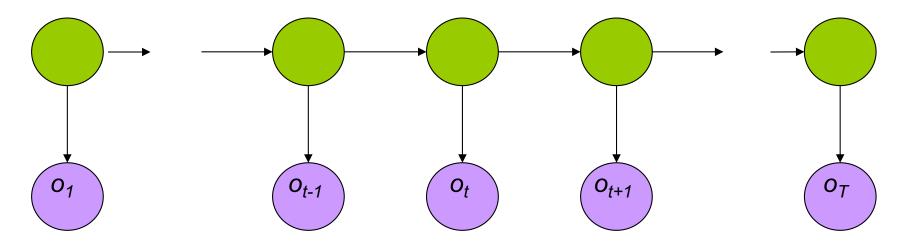
use what we know already

Most likely state at each point in time

$$\hat{X}_t = \arg\max \gamma_i(t)$$

Best State Sequence (2)
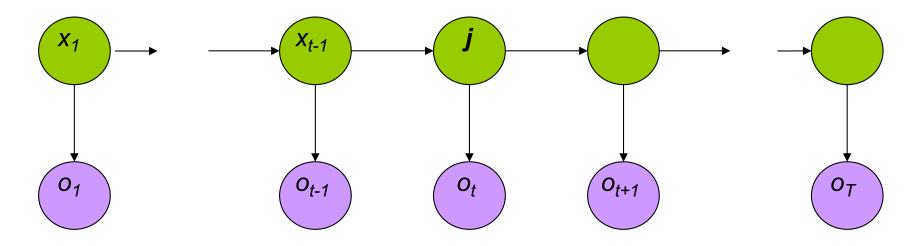
- Find the state sequence that best explains the observations, i.e., MPE

- **Viterbi algorithm**

$$\arg\max_X P(X \mid O)$$

Viterbi Algorithm
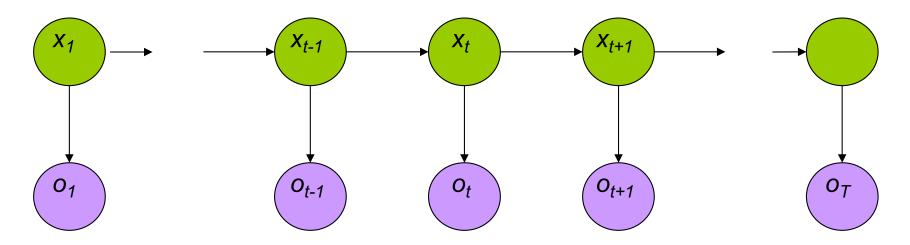
$$\delta_j(t) = \max_{x_1...x_{t-1}} P(x_1...x_{t-1}, o_1...o_{t-1}, x_t = j, o_t)$$

The state sequence which maximizes the probability of seeing the observations to time t-1, landing in state j, and seeing the observation at time t

Viterbi Algorithm

$$\delta_j(t) = \max_{x_1 \ldots x_{t-1}} P(x_1 \ldots x_{t-1}, o_1 \ldots o_{t-1}, x_t = j, o_t)$$

$$\delta_j(t+1) = \max_i \delta_i(t) a_{ij} b_{jo_{t+1}}$$

$$\psi_j(t+1) = \arg\max_i \delta_i(t) a_{ij} b_{jo_{t+1}}$$
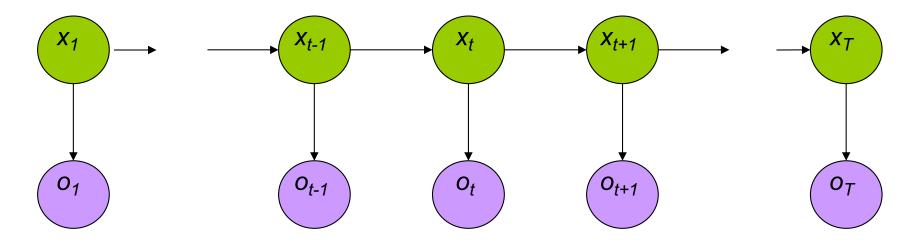
$$Initialization$$

$$\delta_1(i) = \pi_i b_{io_1}$$

$$\psi_1(i) = 0$$

Viterbi Algorithm

$$\hat{X}_T = \arg\max_i \delta_i(T)$$

$$\hat{X}_t = \psi_{\hat{X}_{t+1}}(t+1)$$
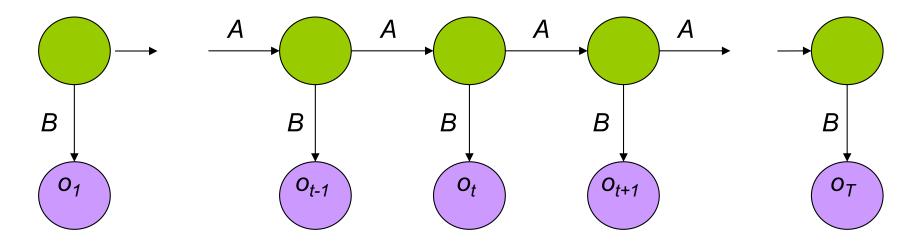
$$P(\hat{X}) = \arg\max_i \delta_i(T)$$

Compute the most likely state sequence by working backwards
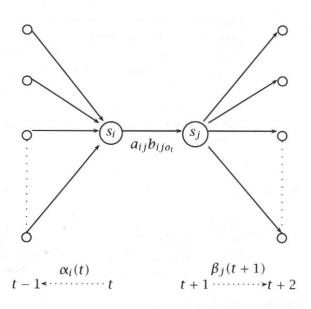
Parameter Estimation using EM

- Given an observation sequence, find the model that is most likely to produce that sequence.

- No analytic method $\arg\max_{\mu} P(O_{training} \mid \mu)$

- Given a model and observation sequence, update the model parameters to better fit the observations.

Figure 9.7 The probability of traversing an arc. Given an observation sequence and a model, we can work out the probability that the Markov process went from state s_i to s_j at time t.

$$\alpha_i(t) = P(o_1 ... o_t, x_t = i \mid \mu)$$

$$\beta_i(t) = P(o_{t+1} ... o_T \mid x_t = i)$$

$$P(O \mid \mu) = \sum_{i=1}^{N} \alpha_i(t)\beta_i(t)$$

Probability of making a transition from i to j at time t

$$p_t(i,j) = P(X_t = i, X_{t+1} = j \mid O, \mu) \quad \text{target}$$

$$= \frac{P(X_t = i, X_{t+1} = j, O \mid \mu)}{P(O \mid \mu)} \quad \text{def of cond prob}$$

Parameter Estimation

$$p_t(i,j) = \frac{\alpha_i(t)a_{ij}b_{jo_{t+1}}\beta_j(t+1)}{\sum\limits_{m=1\ldots N}\alpha_m(t)\beta_m(t)}$$
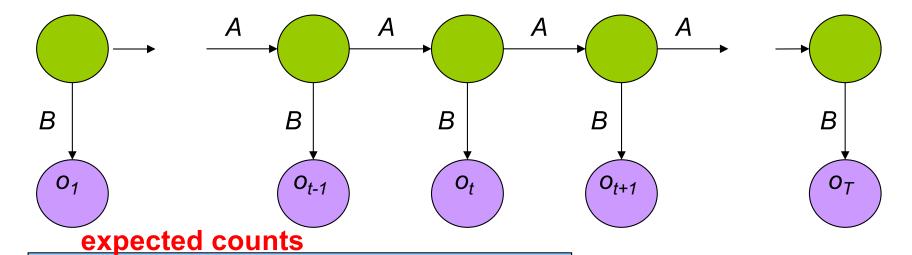
Probability of traversing an arc

$$\gamma_i(t) = \sum\limits_{j=1\ldots N} p_t(i,j)$$

Probability of being in state i

Parameter Estimation

expected counts

$$\hat{\pi}_i = \gamma_i(1)$$

$$\hat{a}_{ij} = \frac{\sum_{t=1}^{T-1} p_t(i,j)}{\sum_{t=1}^{T-1} \gamma_i(t)}$$

$$\hat{b}_{ik} = \frac{\sum_{\{t:o_t=k\}} \gamma_i(t)}{\sum_{t=1}^{T} \gamma_i(t)}$$

Now we can compute the new estimates of the model parameters.

Kristia

46

Instance of EM
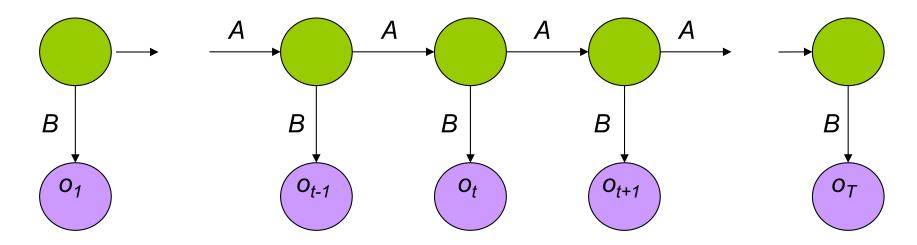
- We have that

$$P(O \mid \hat{\mu}) \geq P(O \mid \mu)$$

- We may get stuck in local maximum (or even saddle point)
- Nevertheless, Baum-Welch usually effective

The Most Important Thing

We can use the special structure of this model to do a lot of neat math and solve problems that are otherwise not solvable.

Kalman Filers
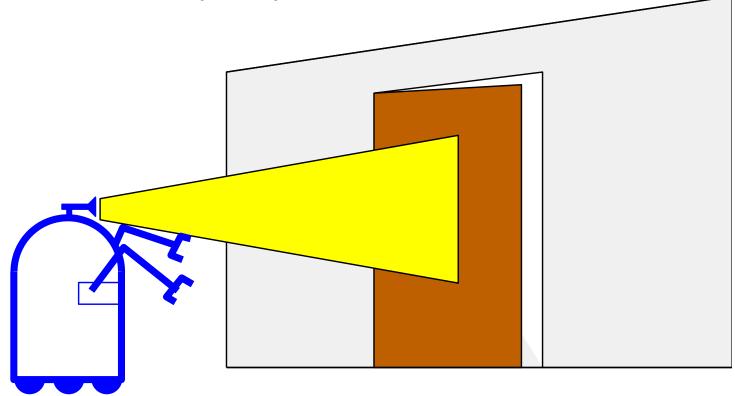
- What about continuous states as e.g. encountered in robotics?

- The Kalman filter is a recursive filter that estimates the state of a linear dynamic system from a series of noisy measurements

Simple Example of State Estimation

- Suppose a robot obtains measurement z
- What is *P(open|z)?*

Example

- *P(z|open) = 0.6*   *P(z|¬open) = 0.3*
- *P(open) = P(¬open) = 0.5*

$$P(open \mid z) = \frac{P(z \mid open)P(open)}{P(z \mid open)p(open) + P(z \mid \neg open)p(\neg open)}$$

$$P(open \mid z) = \frac{0.6 \cdot 0.5}{0.6 \cdot 0.5 + 0.3 \cdot 0.5} = \frac{2}{3} = 0.67$$

Observing z raises the probability that the door is open.

Combining Evidence

- Suppose our robot obtains another observation z_2.

- How can we integrate this new information?

- More generally, how can we estimate $P(x|\ z_1...z_n)$?

Kristian Kersting - Probabilistic Graphical Models

Recursive Bayesian Updating

$$P(x \mid z_1, \ldots, z_n) = \frac{P(z_n \mid x, z_1, \ldots, z_{n-1}) \, P(x \mid z_1, \ldots, z_{n-1})}{P(z_n \mid z_1, \ldots, z_{n-1})}$$

Markov assumption: z_n is independent of $z_1, \ldots, z_{n-1}$ if we know x

$$P(x \mid z_1, \ldots, z_n) = \frac{P(z_n \mid x) \, P(x \mid z_1, \ldots, z_{n-1})}{P(z_n \mid z_1, \ldots, z_{n-1})}$$

$$= \eta \, P(z_n \mid x) \, P(x \mid z_1, \ldots, z_{n-1})$$

$$= \eta_{1 \ldots n} \prod_{i=1 \ldots n} P(z_i \mid x) \, P(x)$$

Kristian Kersting - Probabilistic Graphical Models

Example: Second Measurement

- $P(z_2 | open) = 0.5$ $P(z_2 | \neg open) = 0.6$
- $P(open | z_1) = 2/3$

$$P(open \mid z_2, z_1) = \frac{P(z_2 \mid open)\, P(open \mid z_1)}{P(z_2 \mid open)\, P(open \mid z_1) + P(z_2 \mid \neg open)\, P(\neg open \mid z_1)}$$

$$= \frac{\frac{1}{2} \cdot \frac{2}{3}}{\frac{1}{2} \cdot \frac{2}{3} + \frac{3}{5} \cdot \frac{1}{3}} = \frac{5}{8} = 0.625$$

z_2 lowers the probability that the door is open.

Actions

- As you know already, often the world is **dynamic** since

  - **actions carried out by the robot**,

  - **actions carried out by other agents**,

  - or just the **time** is passing by

  change the world.

- How can we **incorporate** such **actions**?

Typical Actions for a Robot

- The robot **turns its wheels** to move

- The robot **uses its manipulator** to grasp an object

- Plants grow over **time**…

- Actions are **never carried out with absolute certainty**.

- In contrast to measurements, **actions generally increase the uncertainty**.
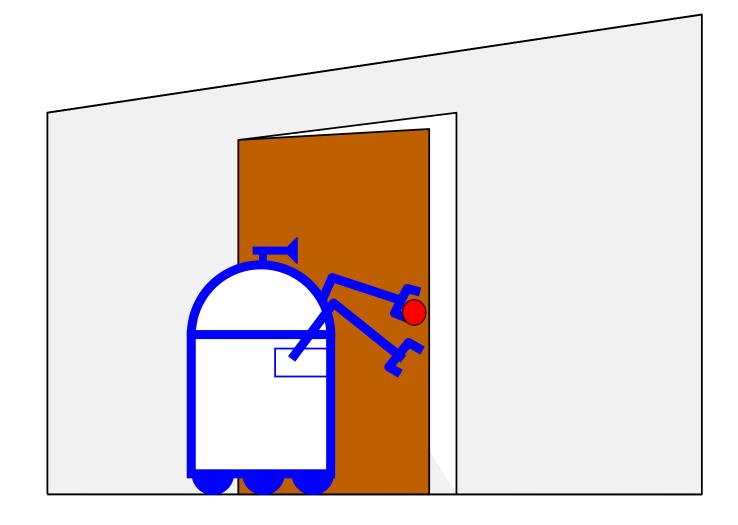
Modeling Action

- To incorporate the outcome of an action *u* into the current "belief", we use the conditional pdf

$$P(x|u,x')$$

- This term specifies the pdf that **executing *u* changes the state from *x' to x*.**
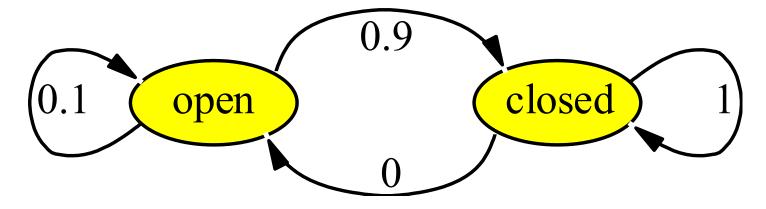
Example: Closing the door

P(x|u,x') for *u* = "close door":

If the door is open, the action "close door" succeeds in 90% of all cases.

Integrating the Outcome of Actions

Continuous case:

$$P(x \mid u) = \int P(x \mid u, x')P(x')dx'$$

Discrete case:

$$P(x \mid u) = \sum_{x'} P(x \mid u, x')P(x')$$

Example: The Resulting Belief

$$P(closed \mid u) = \sum P(closed \mid u, x')P(x')$$

$$= P(closed \mid u, open)P(open)$$

$$+ P(closed \mid u, closed)P(closed)$$

$$= \frac{9}{10} * \frac{5}{8} + \frac{1}{1} * \frac{3}{8} = \frac{15}{16}$$

$$P(open \mid u) = \sum P(open \mid u, x')P(x')$$

$$= P(open \mid u, open)P(open)$$

$$+ P(open \mid u, closed)P(closed)$$

$$= \frac{1}{10} * \frac{5}{8} + \frac{0}{1} * \frac{3}{8} = \frac{1}{16}$$

$$= 1 - P(closed \mid u)$$

Bayes Filters: Framework

- **Given:**

  - Stream of observations *z* and action data *u:*

  $$d_t = \{u_1, z_2 \ldots, u_{t-1}, z_t\}$$

  - Sensor model *P(z|x).*

  - Action model *P(x|u,x').*

  - Prior probability of the system state *P(x).*

- **Wanted:**
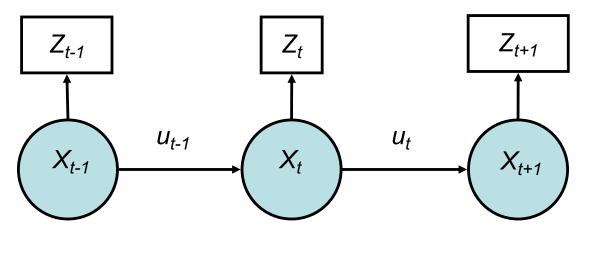
  - Estimate of the state *X* of a dynamical system.

  - The posterior of the state is also called **Belief**:

  $$Bel(x_t) = P(x_t \mid u_1, z_2 \ldots, u_{t-1}, z_t)$$

Kristian Kersting - Probabilistic Graphical Models

Markov Assumption

$$p(d_t, d_{t-1}, ..., d_0 \mid x_t, d_{t+1}, d_{t+2}, ...) = p(d_t, d_{t-1}, ..., d_0 \mid x_t)$$

$$p(d_t, d_{t+1}, ... \mid x_t, d_1, d_2, ..., d_{t-1}) = p(d_t, d_{t+1}, ... \mid x_t)$$

$$p(x_t \mid u_{t-1}, x_{t-1}, d_{t-2}, ..., d_0) = p(x_t \mid u_{t-1}, x_{t-1})$$

Underlying Assumptions

- Static world; Independent noise; Perfect model, no approximation errors

Bayes Filters

$$\boxed{Bel(x_t) =} P(x_t \mid u_1, z_2 \ldots, u_{t-1}, z_t)$$

Bayes
$$= \eta \, P(z_t \mid x_t, u_1, z_2, \ldots, u_{t-1}) \, P(x_t \mid u_1, z_2, \ldots, u_{t-1})$$

Markov
$$= \eta \, P(z_t \mid x_t) \, P(x_t \mid u_1, z_2, \ldots, u_{t-1})$$

Total prob.
$$= \eta \, P(z_t \mid x_t) \int P(x_t \mid u_1, z_2, \ldots, u_{t-1}, x_{t-1})$$
$$P(x_{t-1} \mid u_1, z_2, \ldots, u_{t-1}) \, dx_{t-1}$$

Markov
$$= \eta \, P(z_t \mid x_t) \int P(x_t \mid u_{t-1}, x_{t-1}) \, P(x_{t-1} \mid u_1, z_2, \ldots, u_{t-1}) \, dx_{t-1}$$

$$= \boxed{\eta \, P(z_t \mid x_t) \int P(x_t \mid u_{t-1}, x_{t-1}) \, Bel(x_{t-1}) \, dx_{t-1}}$$

$$Bel(x_t) = \eta \, P(z_t \mid x_t) \int P(x_t \mid u_{t-1}, x_{t-1}) \, Bel(x_{t-1}) \, dx_{t-1}$$

1. Algorithm **Bayes_filter**(*Bel(x),d*):

2. $\eta = 0$

3. if *d* is a perceptual data item *z* then

4.   For all *x* do

5. $$Bel'(x) = P(z \mid x) Bel(x)$$

6. $$\eta = \eta + Bel'(x)$$

7.   For all *x* do

8. $$Bel'(x) = \eta^{-1} Bel'(x)$$

9. else if *d* is an action data item *u* then

10.   For all *x* do

11. $$Bel'(x) = \int P(x \mid u, x') \, Bel(x') \, dx'$$

12. return *Bel'(x)*

Bayes Filters are Familiar!

$$Bel(x_t) = \eta \, P(z_t \mid x_t) \int P(x_t \mid u_{t-1}, x_{t-1}) \, Bel(x_{t-1}) \, dx_{t-1}$$

- Kalman filters

- Particle filters

- Hidden Markov models

- Dynamic Bayes networks (later this class)

- Partially Observable Markov Decision Processes (POMDPs; maybe we touch them later this class)
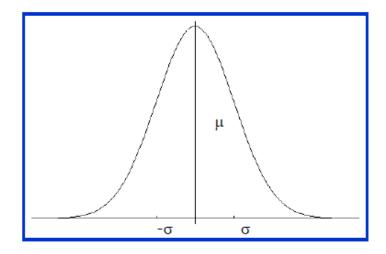
Gaussian

$p(x) \sim N(\mu, \sigma^2):$

$$p(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}}$$
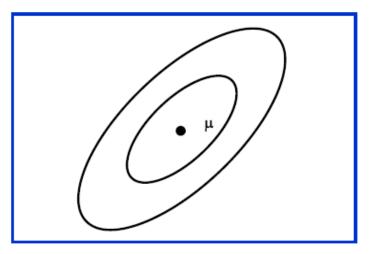
Univariate

$p(\mathbf{x}) \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma}):$

$$p(\mathbf{x}) = \frac{1}{(2\pi)^{d/2}|\boldsymbol{\Sigma}|^{1/2}} e^{-\frac{1}{2}(\mathbf{x}-\boldsymbol{\mu})^t \boldsymbol{\Sigma}^{-1}(\mathbf{x}-\boldsymbol{\mu})}$$

Multivariate

Propertis of Gaussians

$$\left.\begin{array}{l} X \sim N(\mu, \Sigma) \\ Y = AX + B \end{array}\right\} \Rightarrow \quad Y \sim N(A\mu + B, A\Sigma A^T)$$

$$\left.\begin{array}{l} X_1 \sim N(\mu_1, \Sigma_1) \\ X_2 \sim N(\mu_2, \Sigma_2) \end{array}\right\} \Rightarrow p(X_1) \cdot p(X_2) \sim N\left(\frac{\Sigma_2}{\Sigma_1 + \Sigma_2} \mu_1 + \frac{\Sigma_1}{\Sigma_1 + \Sigma_2} \mu_2, \quad \frac{1}{\Sigma_1^{-1} + \Sigma_2^{-1}} \right)$$

- We stay in the "Gaussian world" as long as we start with Gaussians and perform only linear transformations.

Discrete Time Kalman Filter

Estimates the state x of a discrete-time controlled process that is governed by the linear stochastic difference equation

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$$

with a measurement

$$z_t = C_t x_t + \delta_t$$

Components of a Kalman Filter

$\boxed{A_t}$ Matrix (nxn) that describes how the state evolves from t to $t\text{-}1$ without controls or noise.

$\boxed{B_t}$ Matrix (nxl) that describes how the control u_t changes the state from t to $t\text{-}1$.

$\boxed{C_t}$ Matrix (kxn) that describes how to map the state x_t to an observation z_t.

$\boxed{\varepsilon_t}$

$\boxed{\delta_t}$ Random variables representing the process and measurement noise that are assumed to be independent and normally distributed with covariance R_t and Q_t respectively.

Initialization

- Initial belief is normally distributed:

$$bel(x_0) = N(x_0; \mu_0, \Sigma_0)$$

Dynamics

- Dynamics are linear function of state and control plus additive noise:

$$x_t = A_t x_{t-1} + B_t u_t + \varepsilon_t$$

$$p(x_t \mid u_t, x_{t-1}) = N(x_t; A_t x_{t-1} + B_t u_t, R_t)$$

$$\overline{bel}(x_t) = \int p(x_t \mid u_t, x_{t-1}) \qquad\qquad bel(x_{t-1}) \, dx_{t-1}$$

$$\Downarrow \qquad\qquad\qquad\qquad \Downarrow$$

$$\sim N(x_t; A_t x_{t-1} + B_t u_t, R_t) \quad \sim N(x_{t-1}; \mu_{t-1}, \Sigma_{t-1})$$

Belief Update: Prediction

$$\overline{bel}(x_t) = \int p(x_t \mid u_t, x_{t-1}) \qquad\qquad bel(x_{t-1}) \, dx_{t-1}$$

$$\Downarrow \qquad\qquad\qquad\qquad \Downarrow$$

$$\sim N(x_t; A_t x_{t-1} + B_t u_t, R_t) \quad \sim N(x_{t-1}; \mu_{t-1}, \Sigma_{t-1})$$

$$\Downarrow$$

$$\overline{bel}(x_t) = \eta \int \exp\left\{ -\frac{1}{2}(x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1}(x_t - A_t x_{t-1} - B_t u_t) \right\}$$

$$\exp\left\{ -\frac{1}{2}(x_{t-1} - \mu_{t-1})^T \Sigma_{t-1}^{-1}(x_{t-1} - \mu_{t-1}) \right\} dx_{t-1}$$

$$\overline{bel}(x_t) = \begin{cases} \overline{\mu}_t = A_t \mu_{t-1} + B_t u_t \\ \overline{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t \end{cases}$$

Observations

- Observations are linear function of state plus additive noise:

$$z_t = C_t x_t + \delta_t$$

$$p(z_t \mid x_t) = N(z_t; C_t x_t, Q_t)$$

$$bel(x_t) = \quad \eta \quad p(z_t \mid x_t) \qquad\qquad \overline{bel}(x_t)$$

$$\Downarrow \qquad\qquad\qquad \Downarrow$$

$$\sim N(z_t; C_t x_t, Q_t) \qquad \sim N\left(x_t; \overline{\mu}_t, \overline{\Sigma}_t\right)$$

Belief Update: Correction

$$bel(x_t) = \quad \eta \quad p(z_t \mid x_t) \qquad\qquad \overline{bel}(x_t)$$

$$\Downarrow \qquad\qquad\qquad\qquad \Downarrow$$

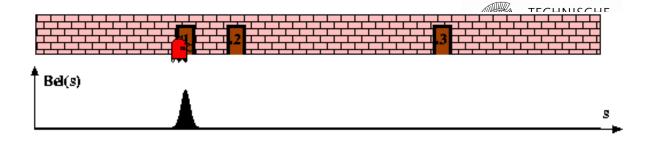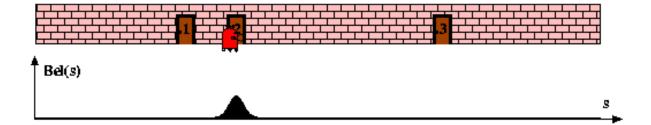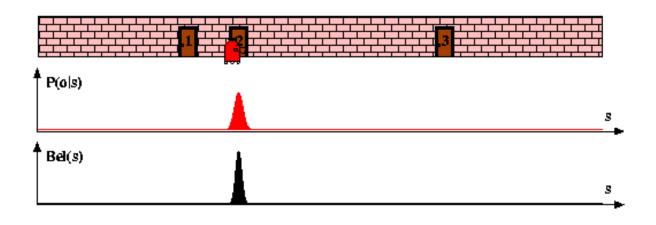$$\sim N(z_t; C_t x_t, Q_t) \qquad \sim N(x_t; \overline{\mu}_t, \overline{\Sigma}_t)$$

$$\Downarrow$$

$$bel(x_t) = \eta \exp\left\{-\frac{1}{2}(z_t - C_t x_t)^T Q_t^{-1}(z_t - C_t x_t)\right\} \exp\left\{-\frac{1}{2}(x_t - \overline{\mu}_t)^T \overline{\Sigma}_t^{-1}(x_t - \overline{\mu}_t)\right\}$$

$$bel(x_t) = \begin{cases} \mu_t = \overline{\mu}_t + K_t(z_t - C_t \overline{\mu}_t) \\ \Sigma_t = (I - K_t C_t)\overline{\Sigma}_t \end{cases} \quad \text{with} \quad K_t = \overline{\Sigma}_t C_t^T (C_t \overline{\Sigma}_t C_t^T + Q_t)^{-1}$$
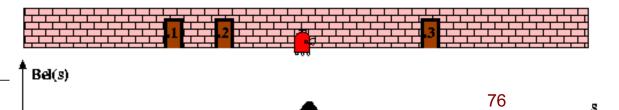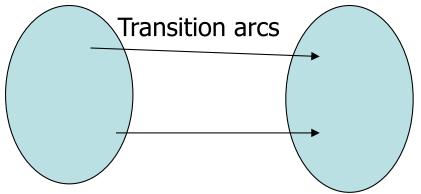
Kalman Filters

We stopped here

So far

- Sometimes, belief states admit compact representations & manipulations:
    - E.g. Kalman filters (assuming Gaussian processes)

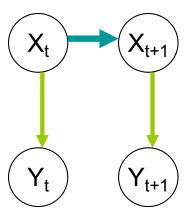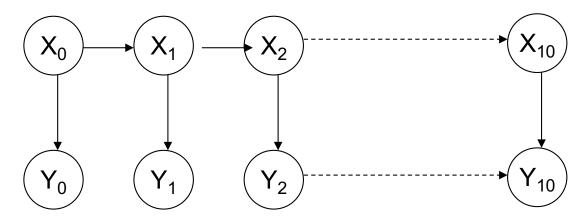- What about Dynamic Bayesian Networks, i.e., the general case ?

Dynamic Bayesian Networks (DBNs)

Transition arcs

Arbitrary Bayesian
Network at time t

The same Bayesian
Network at time t+1

$X_t \rightarrow X_{t+1}$

$X_t \rightarrow Y_t$

$X_{t+1} \rightarrow Y_{t+1}$

$X_0 \rightarrow X_1 \rightarrow X_2 \dashrightarrow X_{10}$

$X_0 \rightarrow Y_0$

$X_1 \rightarrow Y_1$

$X_2 \rightarrow Y_2 \dashrightarrow Y_{10}$

$X_{10} \rightarrow Y_{10}$
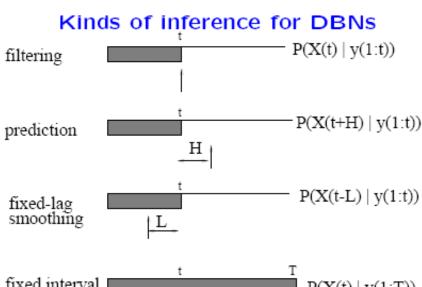
Unrolled DBN for t=0 to t=10

Types of Inference Tasks

Filtering: $\mathbf{P}(\mathbf{X}_t|\mathbf{e}_{1:t})$
  belief state—input to the decision process of a rational agent

Prediction: $\mathbf{P}(\mathbf{X}_{t+k}|\mathbf{e}_{1:t})$ for $k > 0$
  evaluation of possible action sequences;
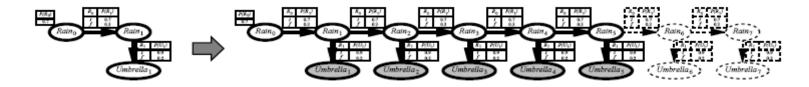 like filtering without the evidence

Smoothing: $\mathbf{P}(\mathbf{X}_k|\mathbf{e}_{1:t})$ for $0 \le k < t$
  better estimate of past states, essential for learning

Most likely explanation: $\arg\max_{\mathbf{x}_{1:t}} P(\mathbf{x}_{1:t}|\mathbf{e}_{1:t})$
  speech recognition, decoding with a noisy channel

Kinds of inference for DBNs

filtering — $P(X(t) \mid y(1{:}t))$

prediction — $P(X(t{+}H) \mid y(1{:}t))$
  H

fixed-lag smoothing — $P(X(t{-}L) \mid y(1{:}t))$
  L

fixed interval smoothing (offline) — $P(X(t) \mid y(1{:}T))$

Exact inference in DBNs

Naive method: unroll the network and run any exact algorithm

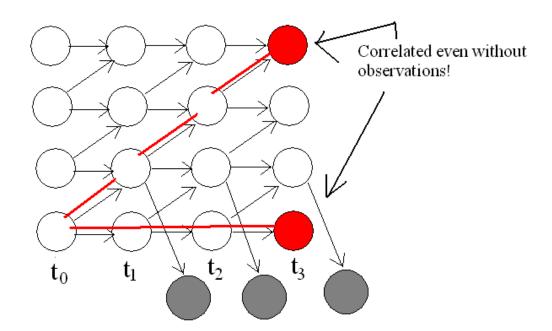Problem: inference cost for each update grows with t

Rollup filtering: add slice $t + 1$, "sum out" slice t using variable elimination

Largest factor is $O(d^{n+1})$, update cost $O(d^{n+2})$
(cf. HMM update cost $O(d^{2n})$)

Problem

- Bayesian Network: a decomposed structure to represent the full joint distribution and to do efficient inference

- Does it imply easy decomposition for the *belief state*?

Correlated even without observations!

- No!
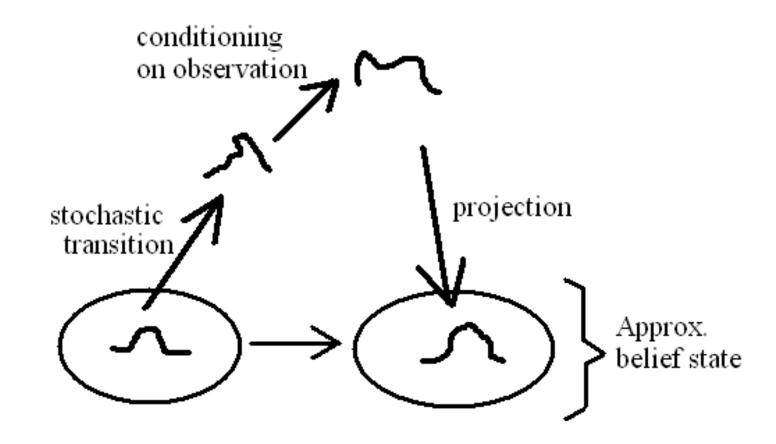
Tractable, approximate representation

- Exact inference in DBN is intractable

- Need approximation

  - Maintain an *approximate belief state*

- E.g. Factored belief state
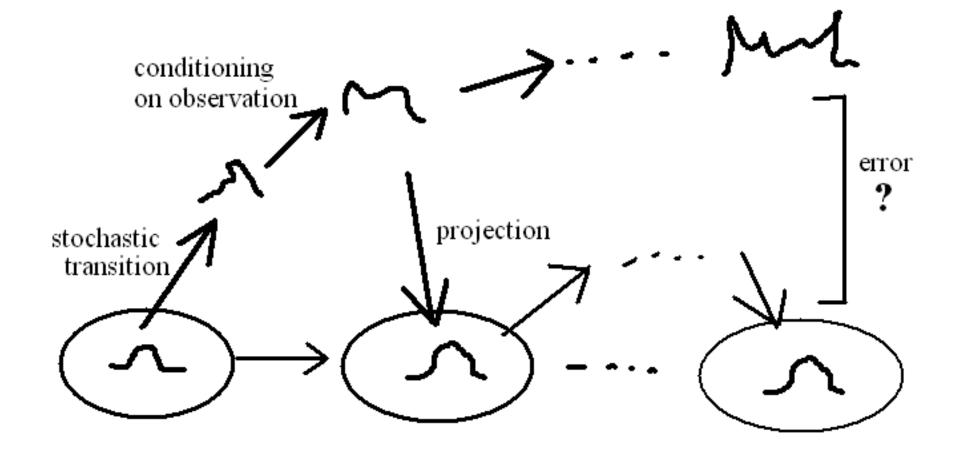
Assumed Density Filtering: Idea

Use a decomposable representation for the belief state (pre-assume some independency)

Is there a Problem?

- What about the approximation errors?
  - It might accumulate and grow unbounded...

Contraction property

- Boyen and Koller have proven that

    - Under reasonable assumptions about the stochasticity of the process, every state transition results in a *contraction* of the distance between the two distributions by a constant factor

    - Since approximation errors from previous steps decrease exponentially, the overall error remains bounded *indefinitely*