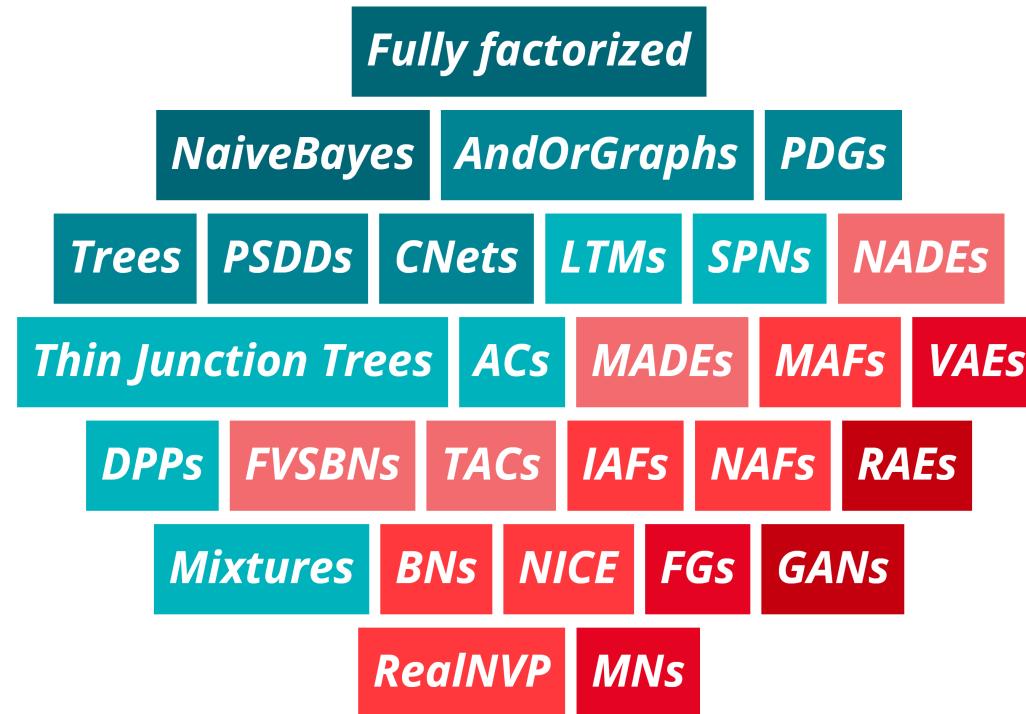


# Tractable Probabilistic Models



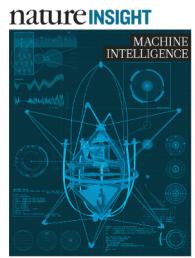
Thanks to Pedro Domingos for making  
slides publically available  
This is also based on joint work with  
Alejandro Molina, Antonio Vergari,  
Robert Peharz, Guy van den Broeck,  
Karl Stelzner and many others. Thanks!

**tractability is a spectrum**





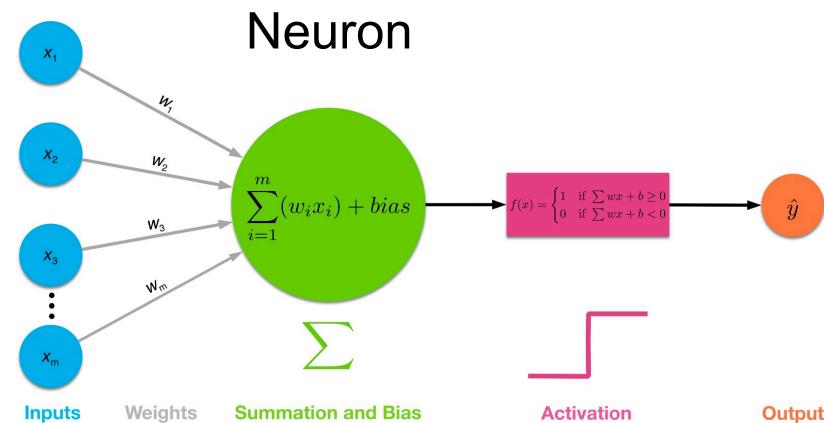
# The Current Second Wave of AI



## Deep Neural Networks

Potentially much more powerful than shallow architectures, represent computations

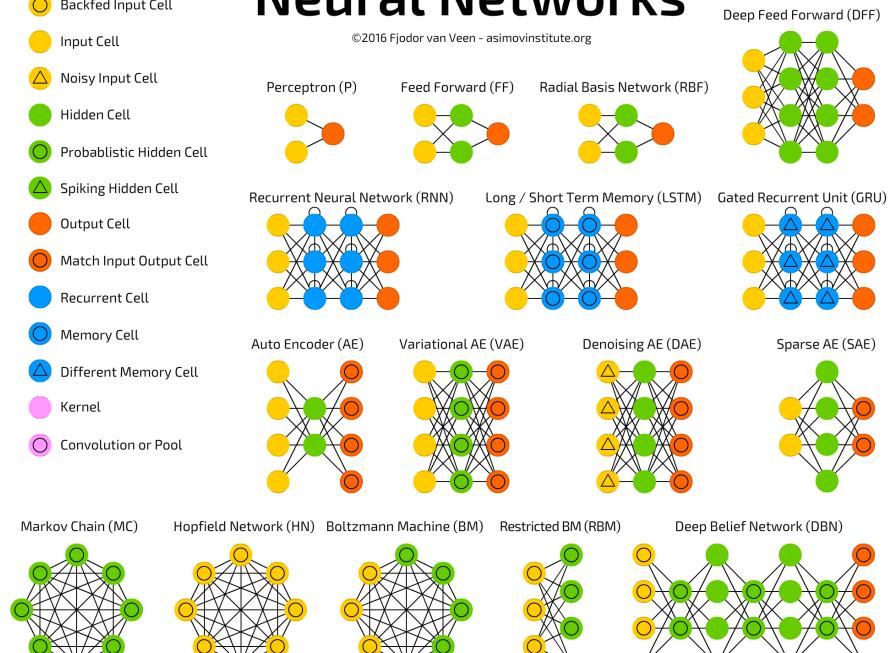
[LeCun, Bengio, Hinton Nature 521



- Backfed Input Cell
- Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Different Memory Cell
- Kernel
- Convolution or Pool

## A mostly complete chart of Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org



**Differentiable Programming**

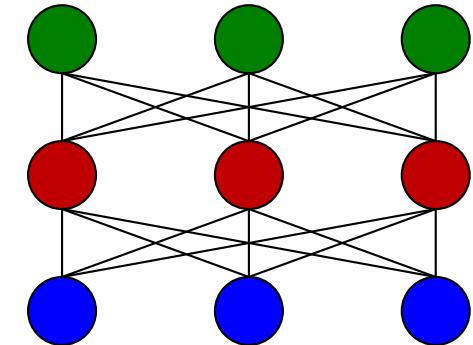
# Deep Learning

# Stack many layers

E.g.: DBN [Hinton & Salakhutdinov,2006]

CDBN [Lee et al.,2009]

## DBM [Salakhutdinov & Hinton,2010]



**Potentially much more powerful than shallow architectures, represent computations** [Bengio, 2009]

**But ...**

- Often no probabilistic semantics
  - Learning requires extensive efforts



DNNs often have no probabilistic semantics. They are not calibrated joint distributions.

$$P(Y|X) \neq P(Y,X)$$

MNIST



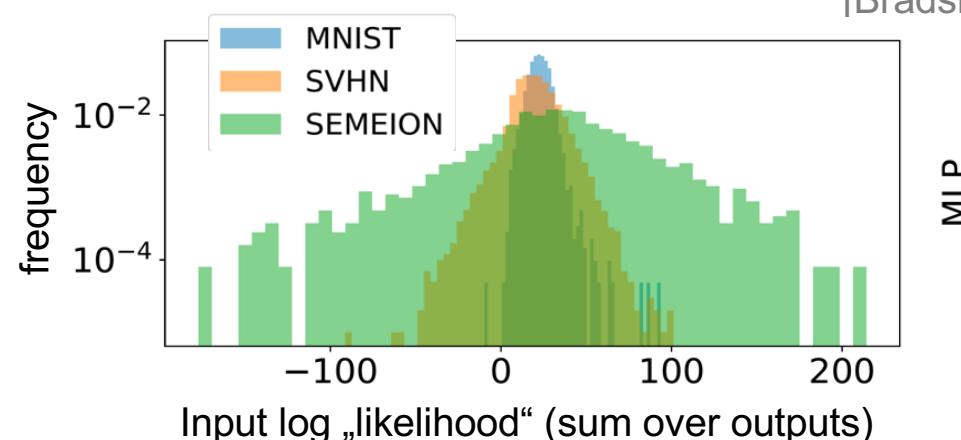
SVHN



SEMEION



Train & Evaluate



Transfer Testing

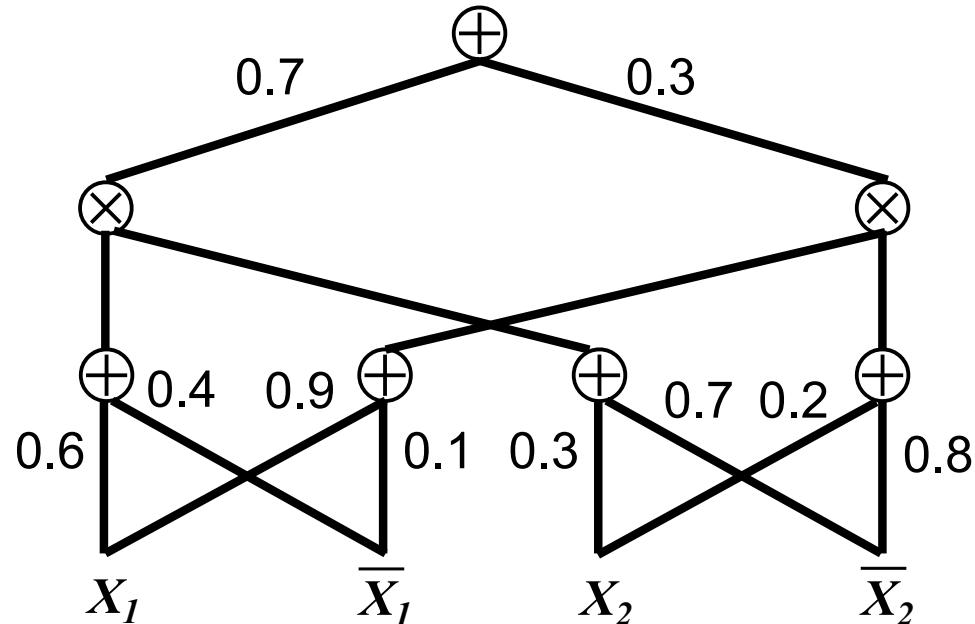
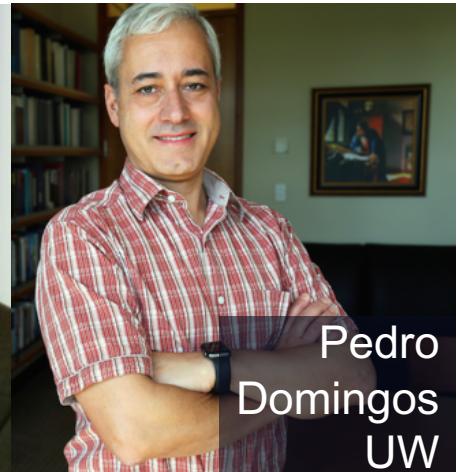
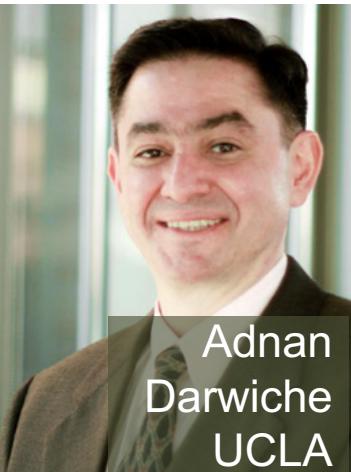
[Bradshaw et al. arXiv:1707.02476 2017]

Many DNNs cannot distinguish the datasets

[Peharz, Vergari, Molina, Stelzner, Trapp, Kersting, Ghahramani UAI 2019]

# Probabilistic Circuits

a deep probabilistic learning framework  
(here sum-product networks)



Computational graph  
(kind of TensorFlow graphs) that encodes how to compute probabilities

Inference is linear in size of network



# Alternative Representation: Joint Distributions as Deep Networks

$X_1$	$X_2$	$P(X)$
1	1	0.4
1	0	0.2
0	1	0.1
0	0	0.3

$$\begin{aligned}
 P(X) = & 0.4 \cdot I[X_1=1] \cdot I[X_2=1] \\
 & + 0.2 \cdot I[X_1=1] \cdot I[X_2=0] \\
 & + 0.1 \cdot I[X_1=0] \cdot I[X_2=1] \\
 & + 0.3 \cdot I[X_1=0] \cdot I[X_2=0]
 \end{aligned}$$

# Alternative Representation: Joint Distributions as (Deep) Networks

$X_1$	$X_2$	$P(X)$
1	1	0.4
1	0	0.2
0	1	0.1
0	0	0.3

$$\begin{aligned}
 P(X) = & \mathbf{0.4} \cdot I[X_1=1] \cdot I[X_2=1] \\
 & + 0.2 \cdot I[X_1=1] \cdot I[X_2=0] \\
 & + 0.1 \cdot I[X_1=0] \cdot I[X_2=1] \\
 & + 0.3 \cdot I[X_1=0] \cdot I[X_2=0]
 \end{aligned}$$

# Shorthand for Indicators

$X_1$	$X_2$	$P(X)$
1	1	0.4
1	0	0.2
0	1	0.1
0	0	0.3

$$\begin{aligned}P(X) = & 0.4 \cdot X_1 \cdot X_2 \\& + 0.2 \cdot X_1 \cdot \bar{X}_2 \\& + 0.1 \cdot \bar{X}_1 \cdot X_2 \\& + 0.3 \cdot \bar{X}_1 \cdot \bar{X}_2\end{aligned}$$

# Sum Out Variables

$X_1$	$X_2$	$P(X)$
1	1	0.4
1	0	0.2
0	1	0.1
0	0	0.3

$$e: X_1 = 1$$

$$\begin{aligned} P(e) &= \mathbf{0.4} \cdot X_1 \cdot X_2 \\ &\quad + \mathbf{0.2} \cdot X_1 \cdot \bar{X}_2 \\ &\quad + 0.1 \cdot \bar{X}_1 \cdot X_2 \\ &\quad + 0.3 \cdot \bar{X}_1 \cdot \bar{X}_2 \end{aligned}$$

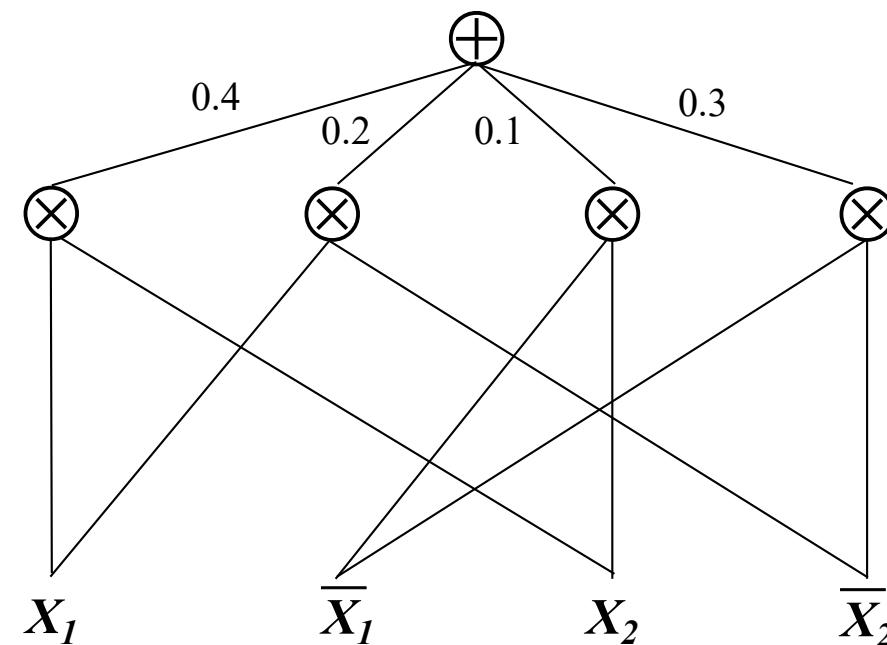
Set  $X_1 = 1, \bar{X}_1 = 0, X_2 = 1, \bar{X}_2 = 1$

Easy: Set both indicators of  $X_2$  to 1



# Idea: Deeper Network Representation of a Graphical Model that encodes how to compute probabilities

$X_1$	$X_2$	$P(X)$
1	1	0.4
1	0	0.2
0	1	0.1
0	0	0.3



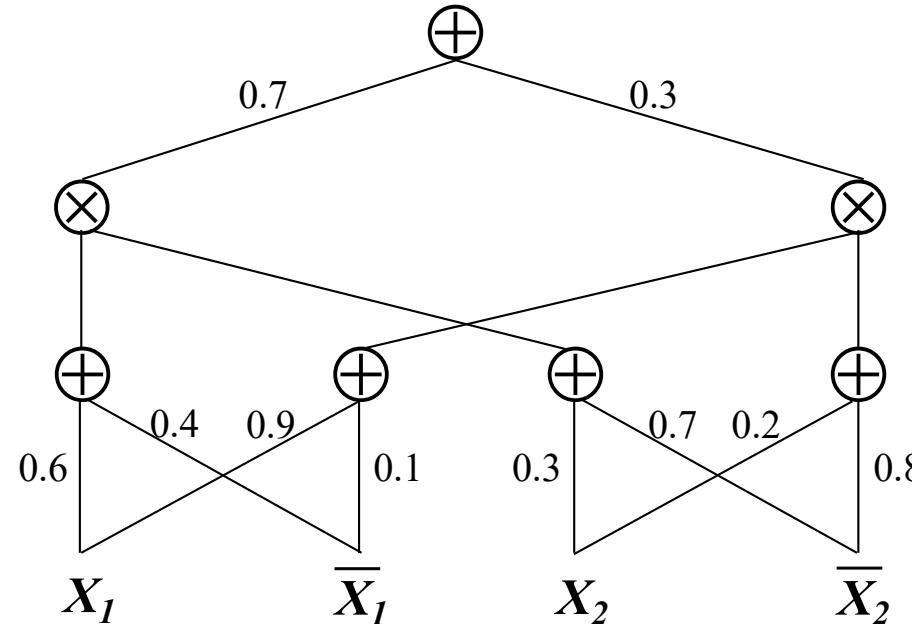
# Sum-Product Networks\* (SPNs)



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

[Poon, Domingos UAI 2011]

A SPN **S** is a rooted DAG where:  
Nodes: Sum, product, input indicator  
Weights on edges from sum to children



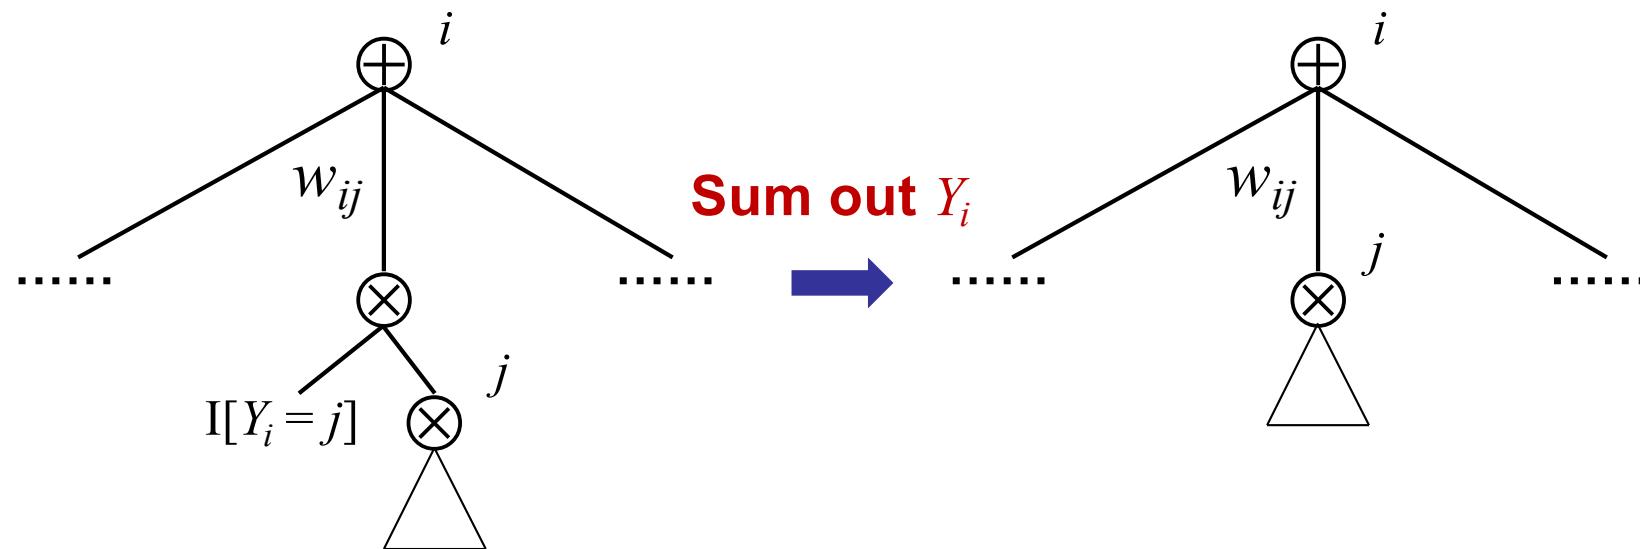
\*SPNs are an instance of Arithmetic Circuits (ACs). ACs have been introduced into the AI literature more than 15 years ago as a tractable representation of probability distributions  
[Darwiche CACM 48(4):608-647 2001]

# Semantics of Sums and Products



Product ~ **Feature** → Form feature hierarchy

Sum ~ **Mixture** (with hidden var. summed out)



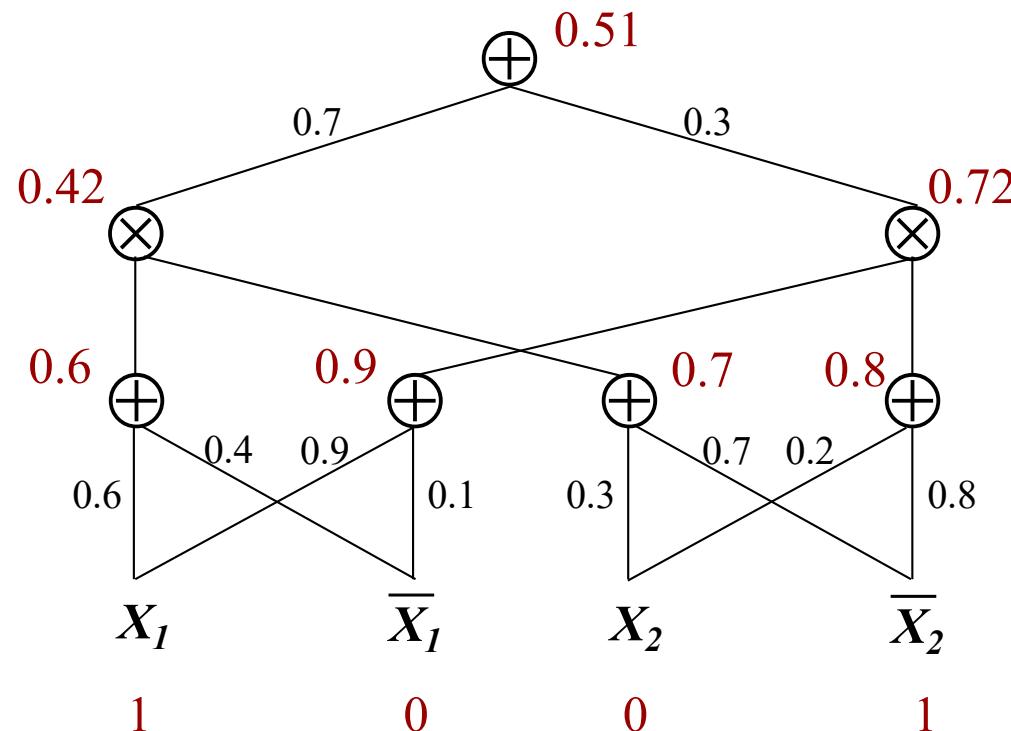
# Inference: Linear in Size of Network

As long as weights sum to 1  
at each sum node

$$P(X) = S(X)$$

$$X: X_1 = 1, X_2 = 0$$

$X_1$	1
$\bar{X}_1$	0
$X_2$	0
$\bar{X}_2$	1

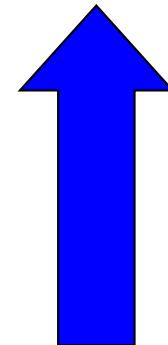
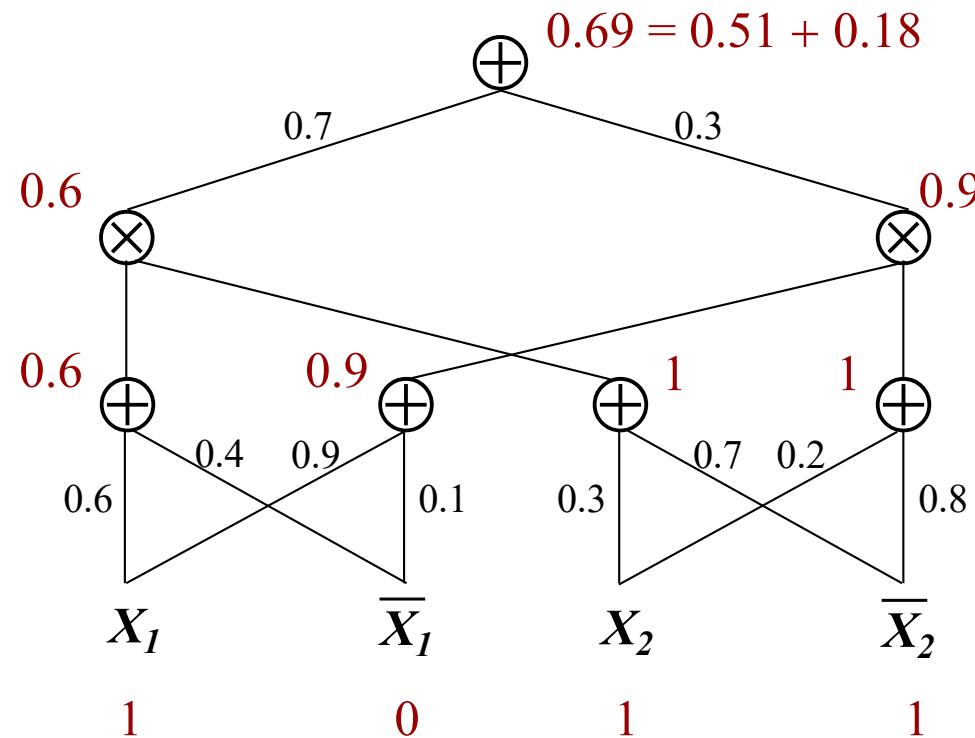


# Inference: Linear in Size of Network

Marginal:  $P(e) = S(e)$

$e: X_1 = 1$

$X_1$	1
$\bar{X}_1$	0
$X_2$	1
$\bar{X}_2$	1



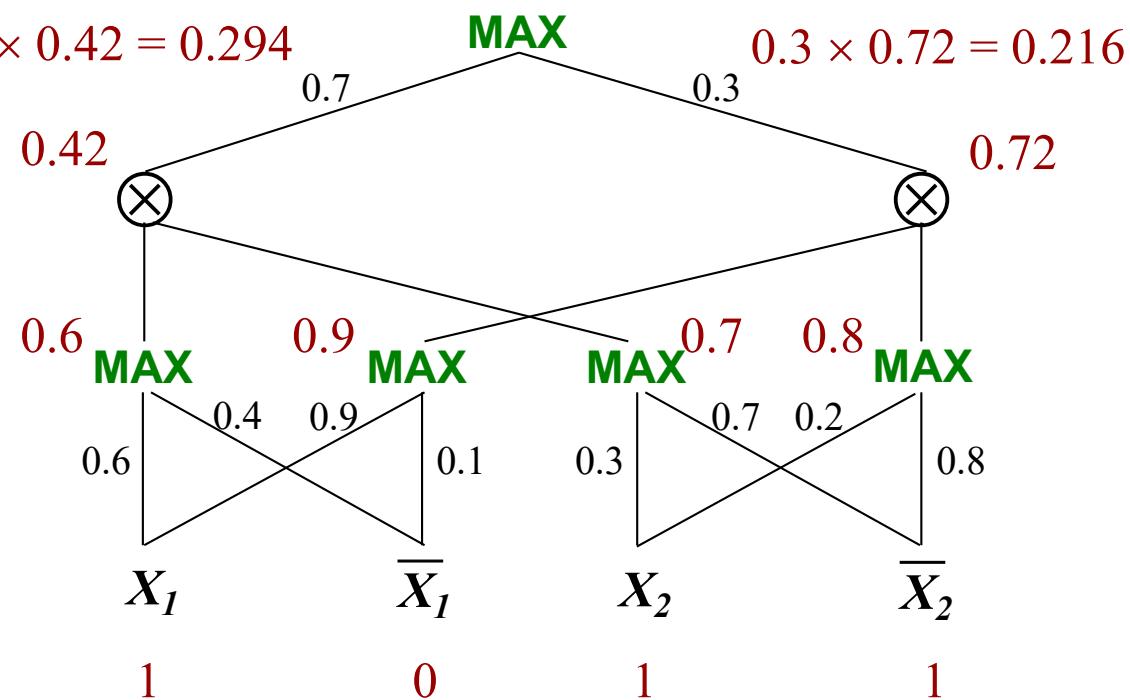
# Inference: Linear in Size of Network

MAP: Replace sums with maxs

$$e: X_1 = 1$$

$$0.7 \times 0.42 = 0.294$$

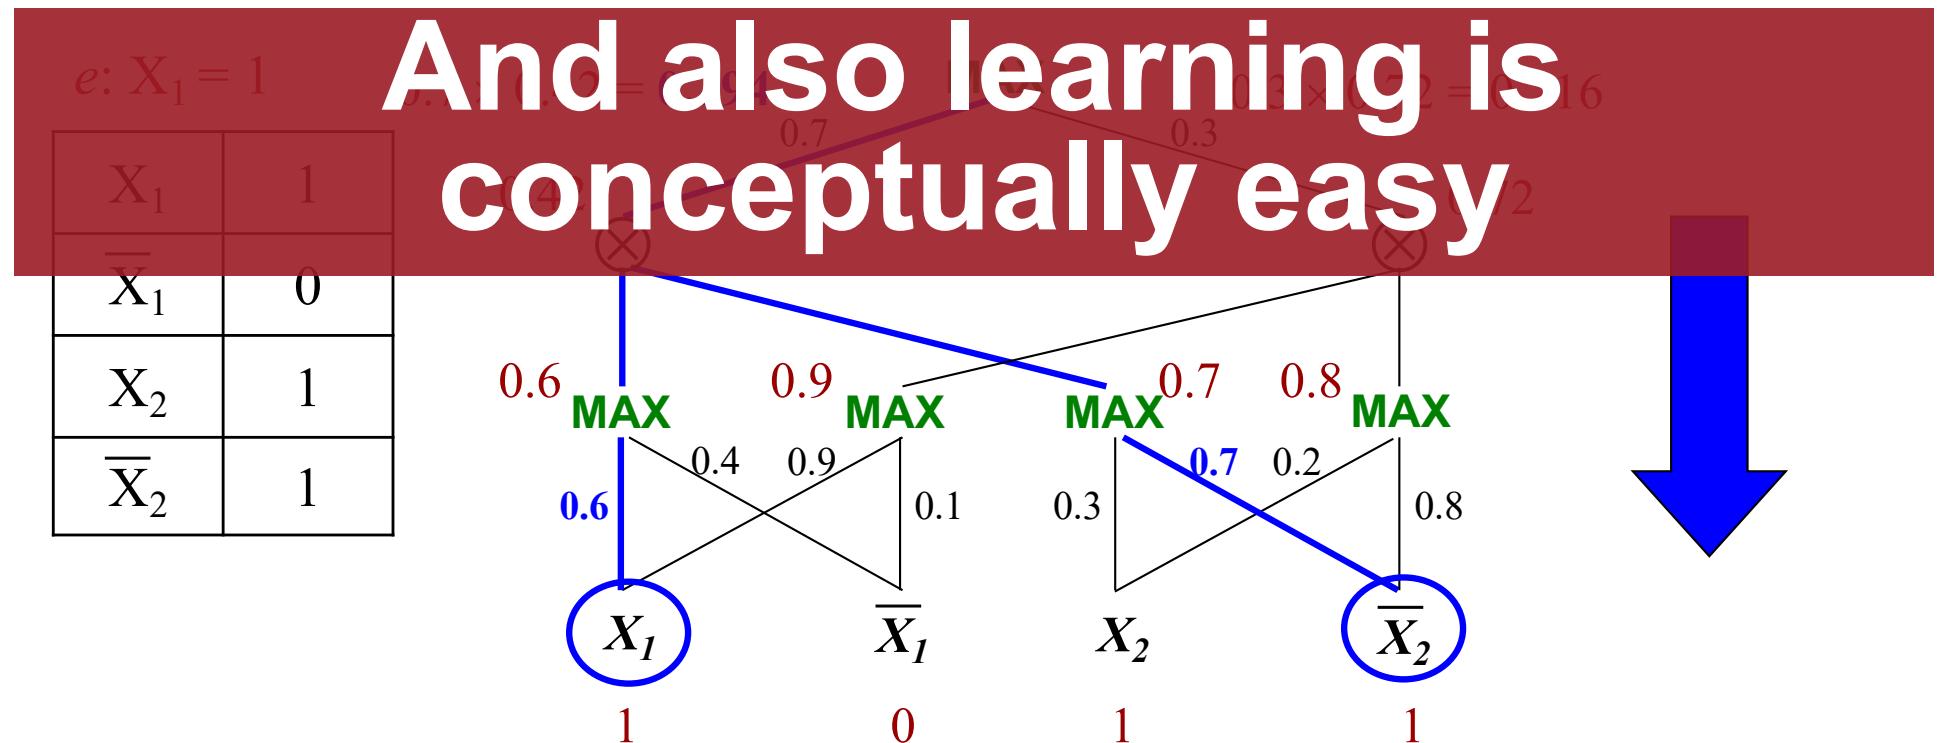
$X_1$	1
$\bar{X}_1$	0
$X_2$	1
$\bar{X}_2$	1



# Inference: Linear in Size of Network

MAX: Pick child with highest value

MAP State:  $X_1 = 1, X_2 = 0$



# General Approach via Parameter Estimation assuming a fixed network (like in Deep Learning)

- Start with a dense SPN
- Find the structure by (online) learning weights  
Zero weights signify absence of connections
- (Hard) EM beneficial to avoid gradient vanishing  
Each sum node is a mixture over children

In principle you can turn a given SPN into a TensorFlow computation graph and apply any known algorithm from there



# Image Completion

**Main evaluation: Caltech-101 [Fei-Fei et al., 2004]**

- 101 categories, e.g., faces, cars, elephants
- Each category: 30 – 800 images

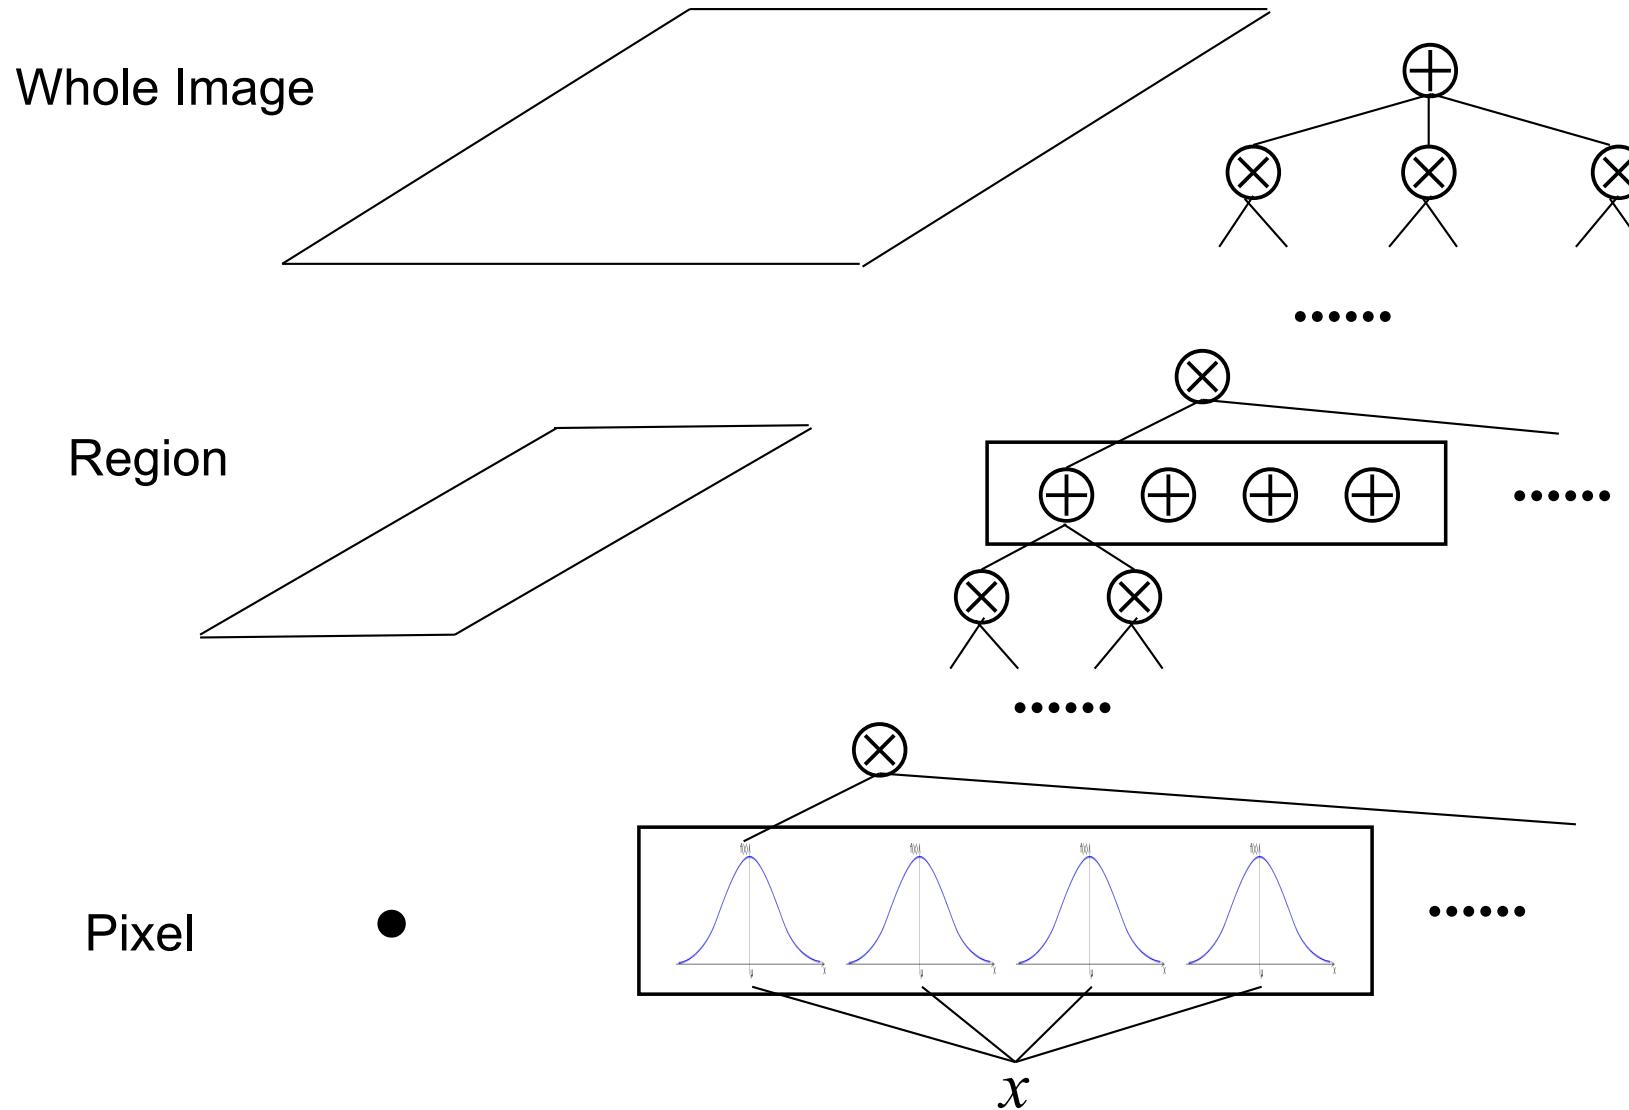
Also, Olivetti [Samaria & Harter, 1994] (400 faces)

Each category: Last third for test

**Test images: Unseen objects**



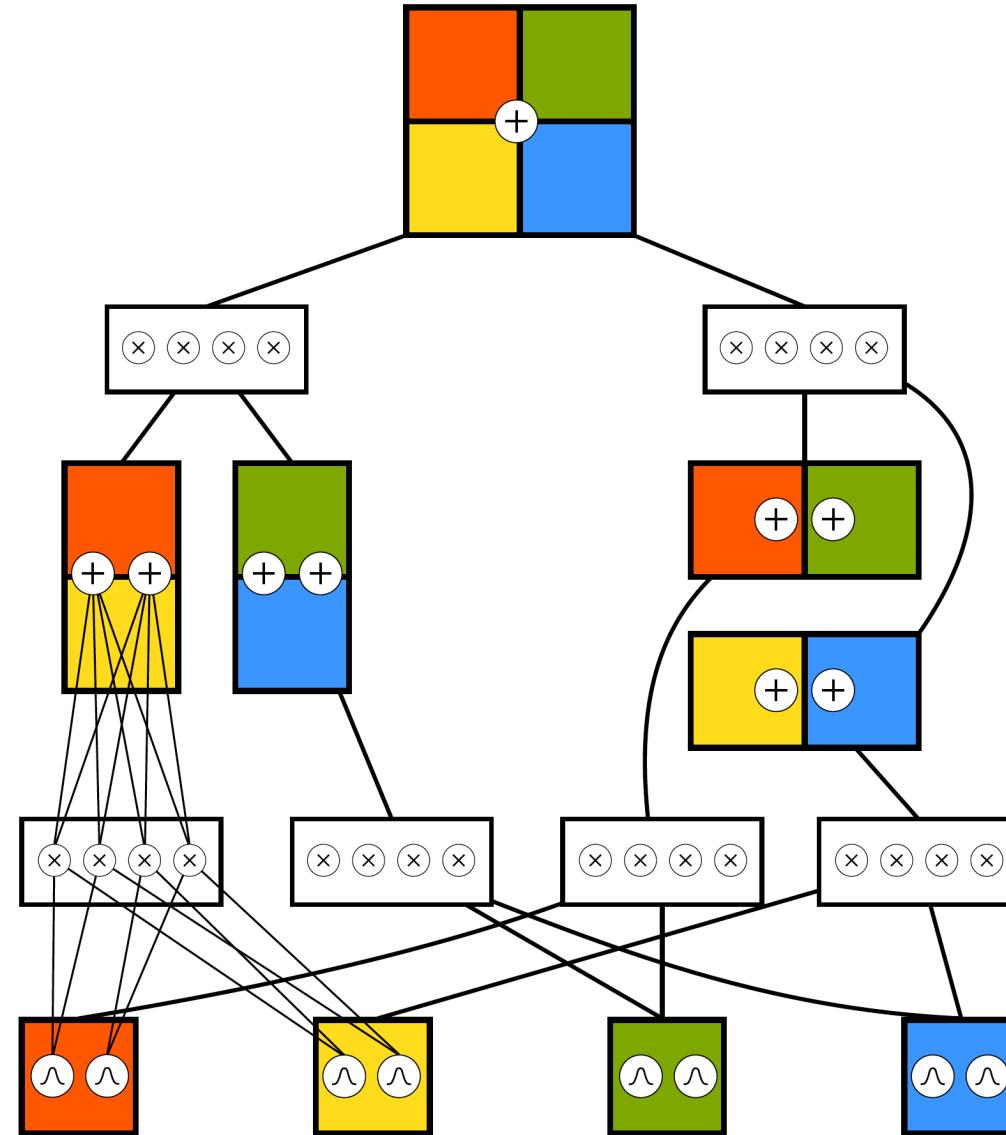
# SPN Architecture



# SPN Architecture

Whole Image

Regions



# Caltech: Mean-Square Errors

	LEFT	BOTTOM
<b>SPN</b>	<b>3551</b>	<b>3270</b>
<b>DBM</b>	9043	9792
<b>DBN</b>	4778	4492
<b>PCA</b>	4234	4465
<b>Nearest Neighbor</b>	4887	5505



# SPN vs. DBM / DBN

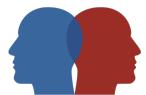
SPN is order of magnitude faster

	SPN	DBM / DBN
Learning	2-3 hours	Days
Inference	< 1 second	Minutes or hours

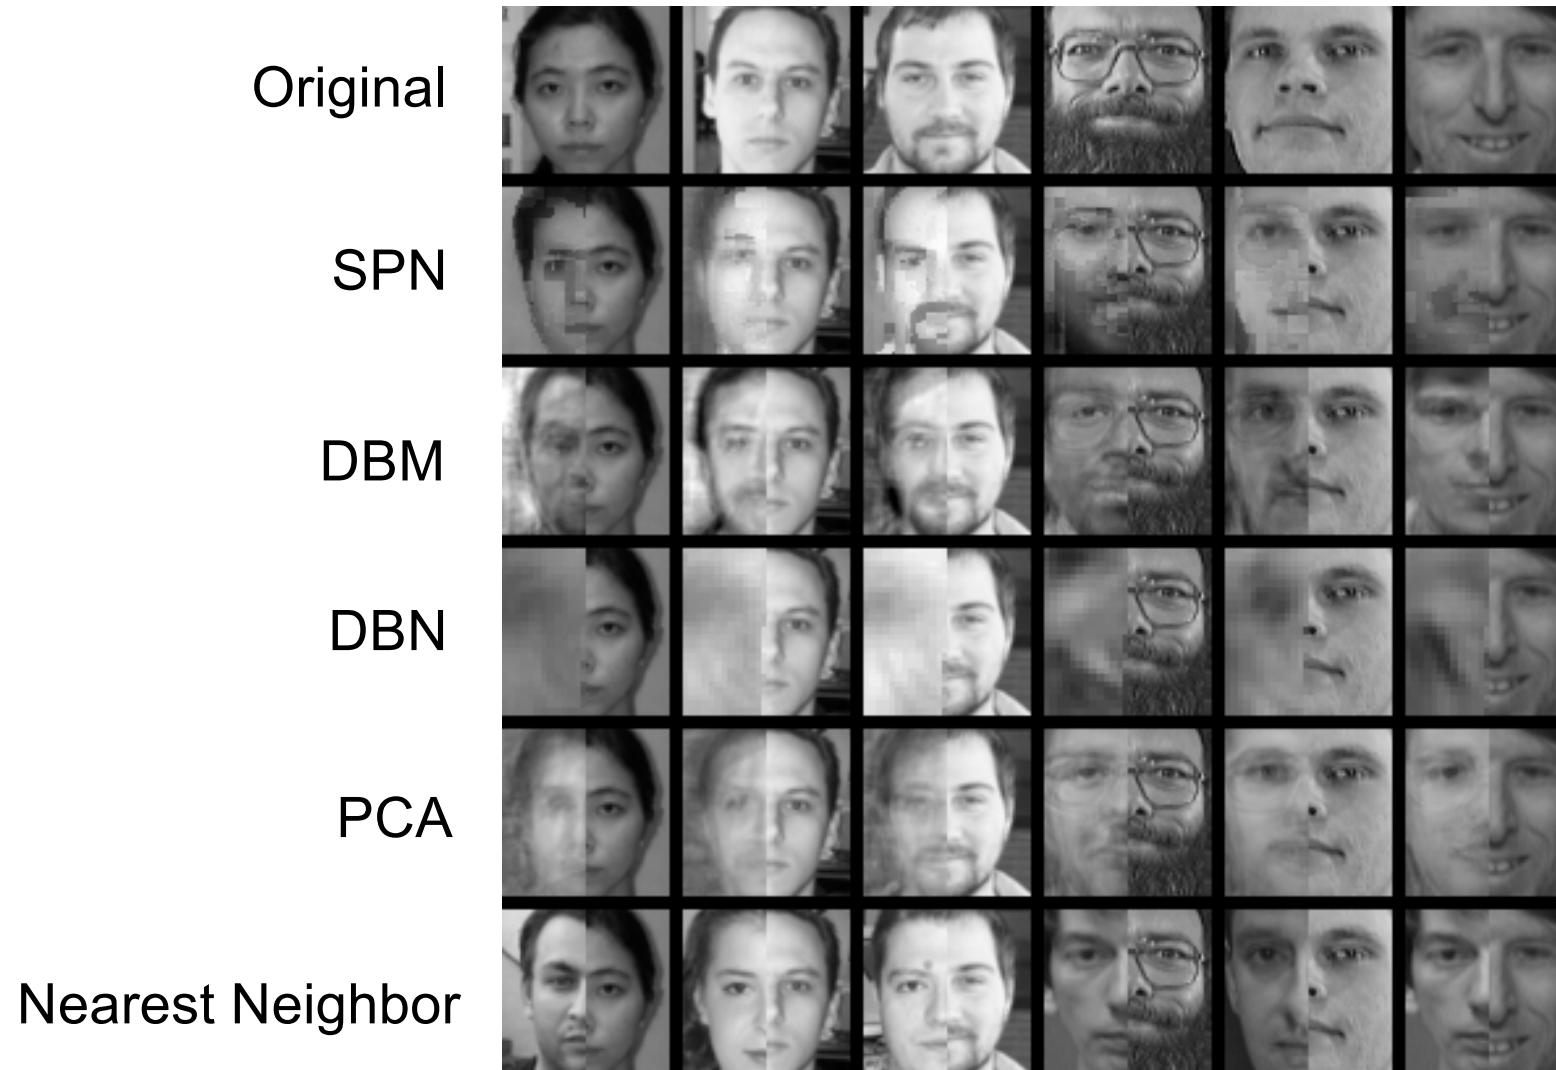
No elaborate preprocessing, tuning

Reduced errors by 30-60%

Learned up to 46 layers

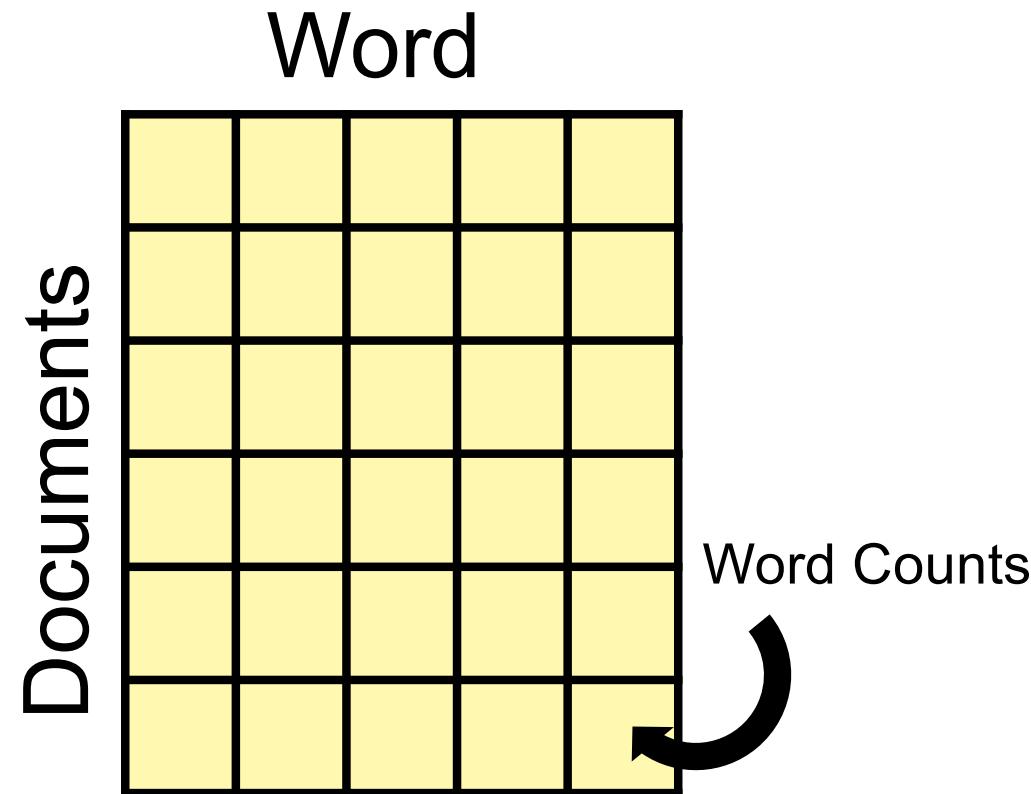


# Example Completions



# Or we learn directly (Tree-)SPNs

Testing independence of  
Poisson random variables

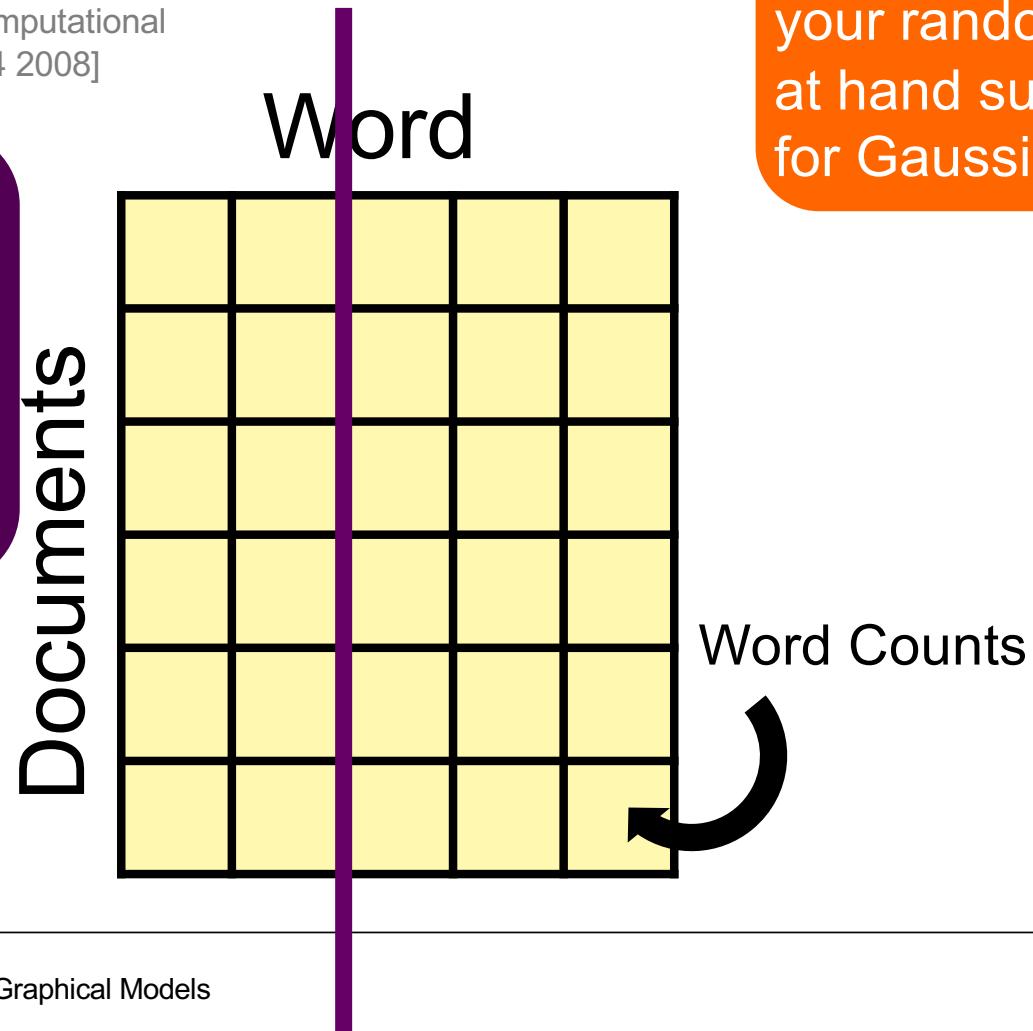


# Or we learn directly (Tree-)SPNs

## Testing independence of Poisson random variables

[Zeileis, Hothorn, Hornik Journal of Computational And Graphical Statistics 17(2):492–514 2008]

Learn Poisson model trees for  $P(x|V-x)$  and  $P(y|V-y)$ . Check whether X resp. Y is significant in  $P(y|V-x)$  resp.  $P(x|V-y)$



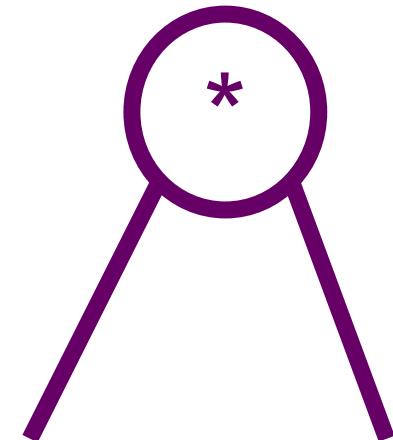
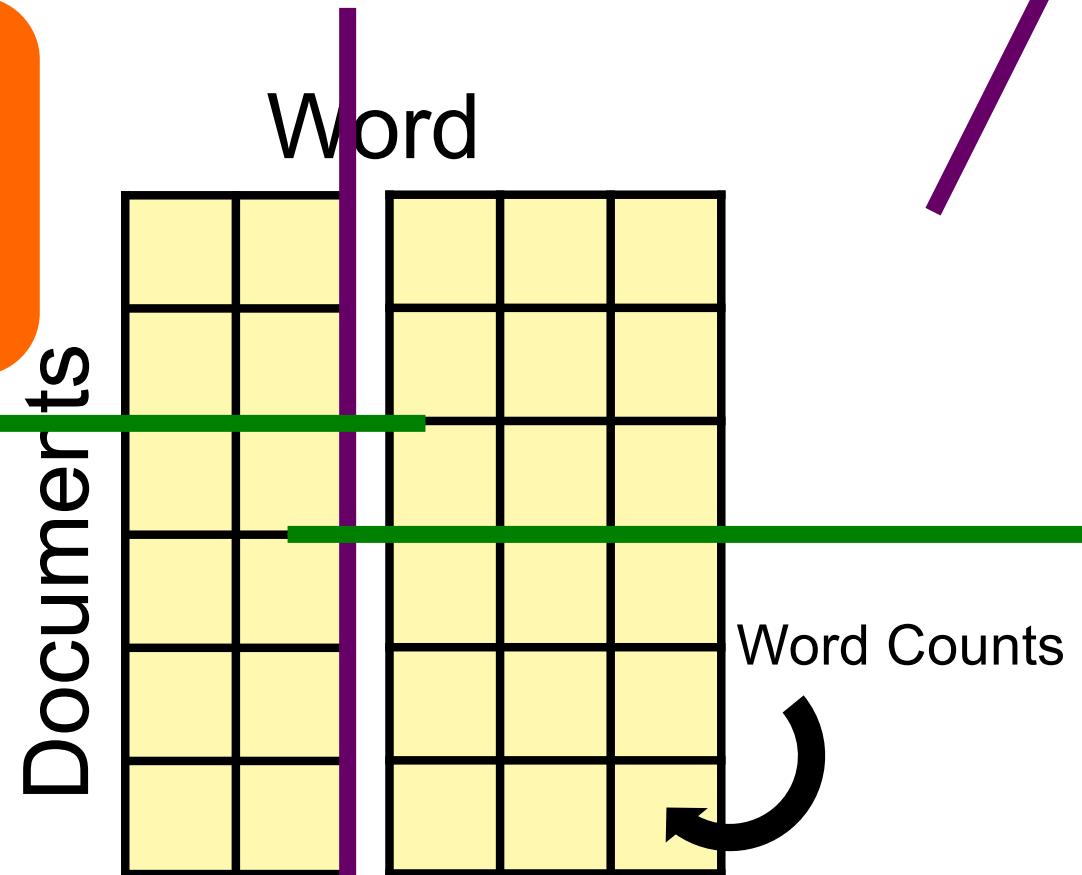
In general use the independency test for your random variables at hand such as g-test for Gaussians

# Or we learn directly (Tree-)SPNs

Testing independence of  
Poisson random variables

In general some clustering for your random variables at hand such as kMeans for Gaussians

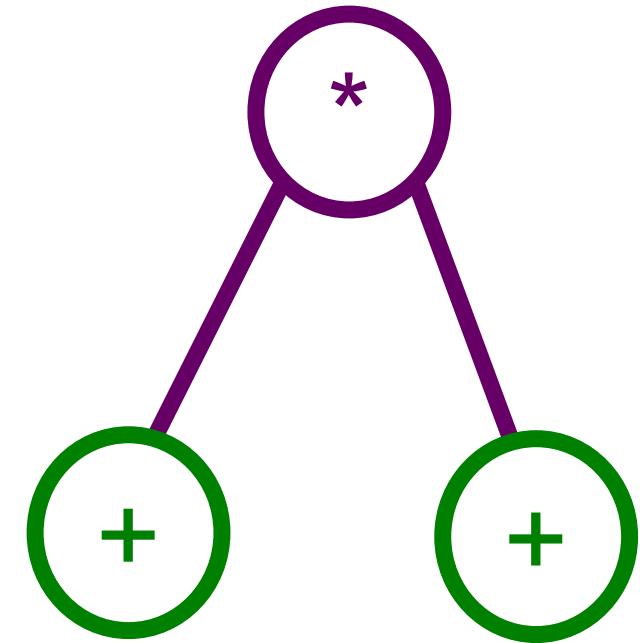
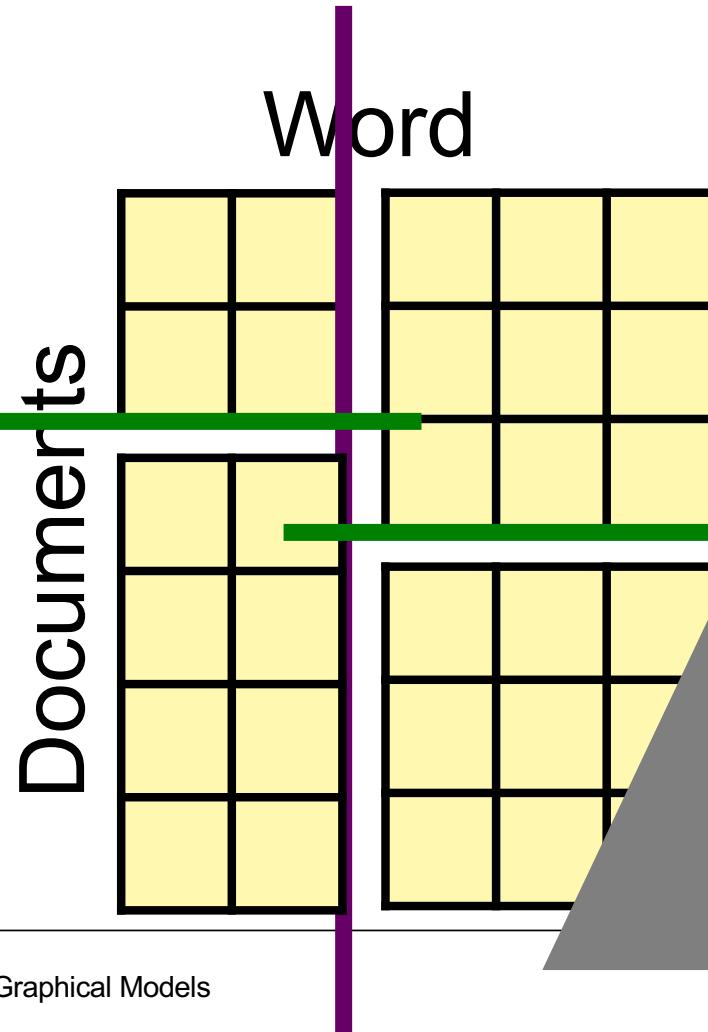
Mixture of Poisson Dependency Networks or random splits



# Or we learn directly (Tree-)SPNs

Testing independence of  
Poisson random variables

Mixture of  
Poisson  
Dependency  
Networks or  
random splits



keep growing  
alternatingly  
and + layers

[Poon, Domingos UAI'11; Molina, Natarajan, Kersting AAAI'17; Vergari, Peharz, Di Mauro, Molina, Kersting, Esposito AAAI '18; Molina, Vergari, Di Mauro, Esposito, Natarajan, Kersting AAAI '18, Peharz et al. UAI 2019, Stelzner, Peharz, Kersting iCML 2019]

# FL<sup>+</sup> SPFlow: An Easy and Extensible Library XW for Sum-Product Networks



UNIVERSITÀ  
DEGLI STUDI DI BARI  
ALDO MORO



[Molina, Vergari, Stelzner, Peharz, Subramani, Poupart, Di Mauro, Kersting arXiv:1901.03704, 2019]



Federal Ministry  
of Education  
and Research

195 commits

2 branches

0 releases

6 contrib.....

Branch: master ▾

New pull request

Create new file

Upload files

Find file

Clone or download ▾

<https://github.com/SPFlow/SPFlow>

```
from spn.structure.leaves.parametric.Parametric import Categorical
from spn.structure.Base import Sum, Product
from spn.structure.base import assign_ids, rebuild_scopes_bottom_up

p0 = Product(children=[Categorical(p=[0.3, 0.7], scope=1), Categorical(p=[0.4, 0.6], scope=2)])
p1 = Product(children=[Categorical(p=[0.5, 0.5], scope=1), Categorical(p=[0.6, 0.4], scope=2)])
s1 = Sum(weights=[0.3, 0.7], children=[p0, p1])
p2 = Product(children=[Categorical(p=[0.2, 0.8], scope=0), s1])
p3 = Product(children=[Categorical(p=[0.2, 0.8], scope=0), Categorical(p=[0.3, 0.7], scope=1)])
p4 = Product(children=[p3, Categorical(p=[0.4, 0.6], scope=2)])
spn = Sum(weights=[0.4, 0.6], children=[p2, p4])

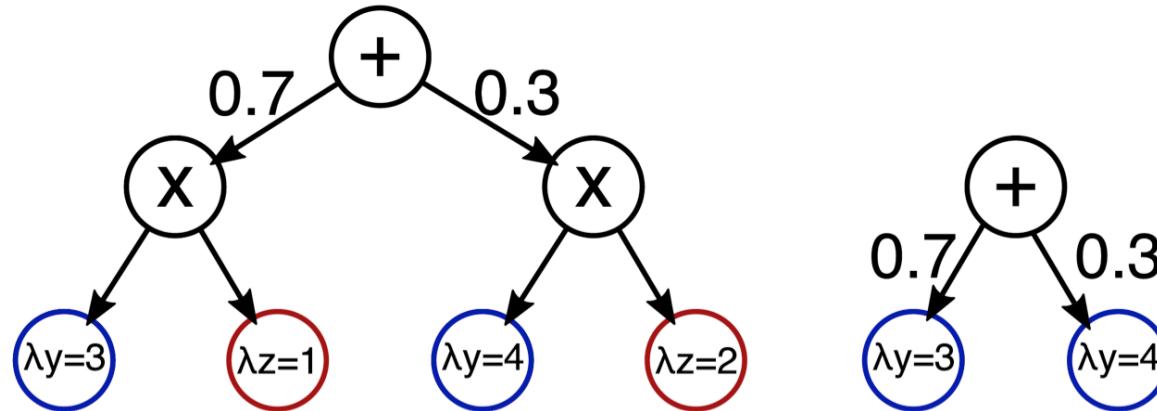
assign_ids(spn)
rebuild_scopes_bottom_up(spn)

return spn
```

**Domain Specific Language,  
Inference, EM, and Model  
Selection as well as  
Compilation of SPNs into TF  
and PyTorch and also into flat,  
library-free code even suitable  
for running on devices:  
C/C++, GPU, FPGA**

SPFlow, an open-source Python library providing a simple interface to inference, learning and manipulation routines for deep and tractable probabilistic models called Sum-Product Networks (SPNs). The library allows one to quickly create SPNs both from data and through a domain specific language (DSL). It efficiently implements several probabilistic inference routines like computing marginals, conditionals and (approximate) most probable explanations (MPEs) along with sampling

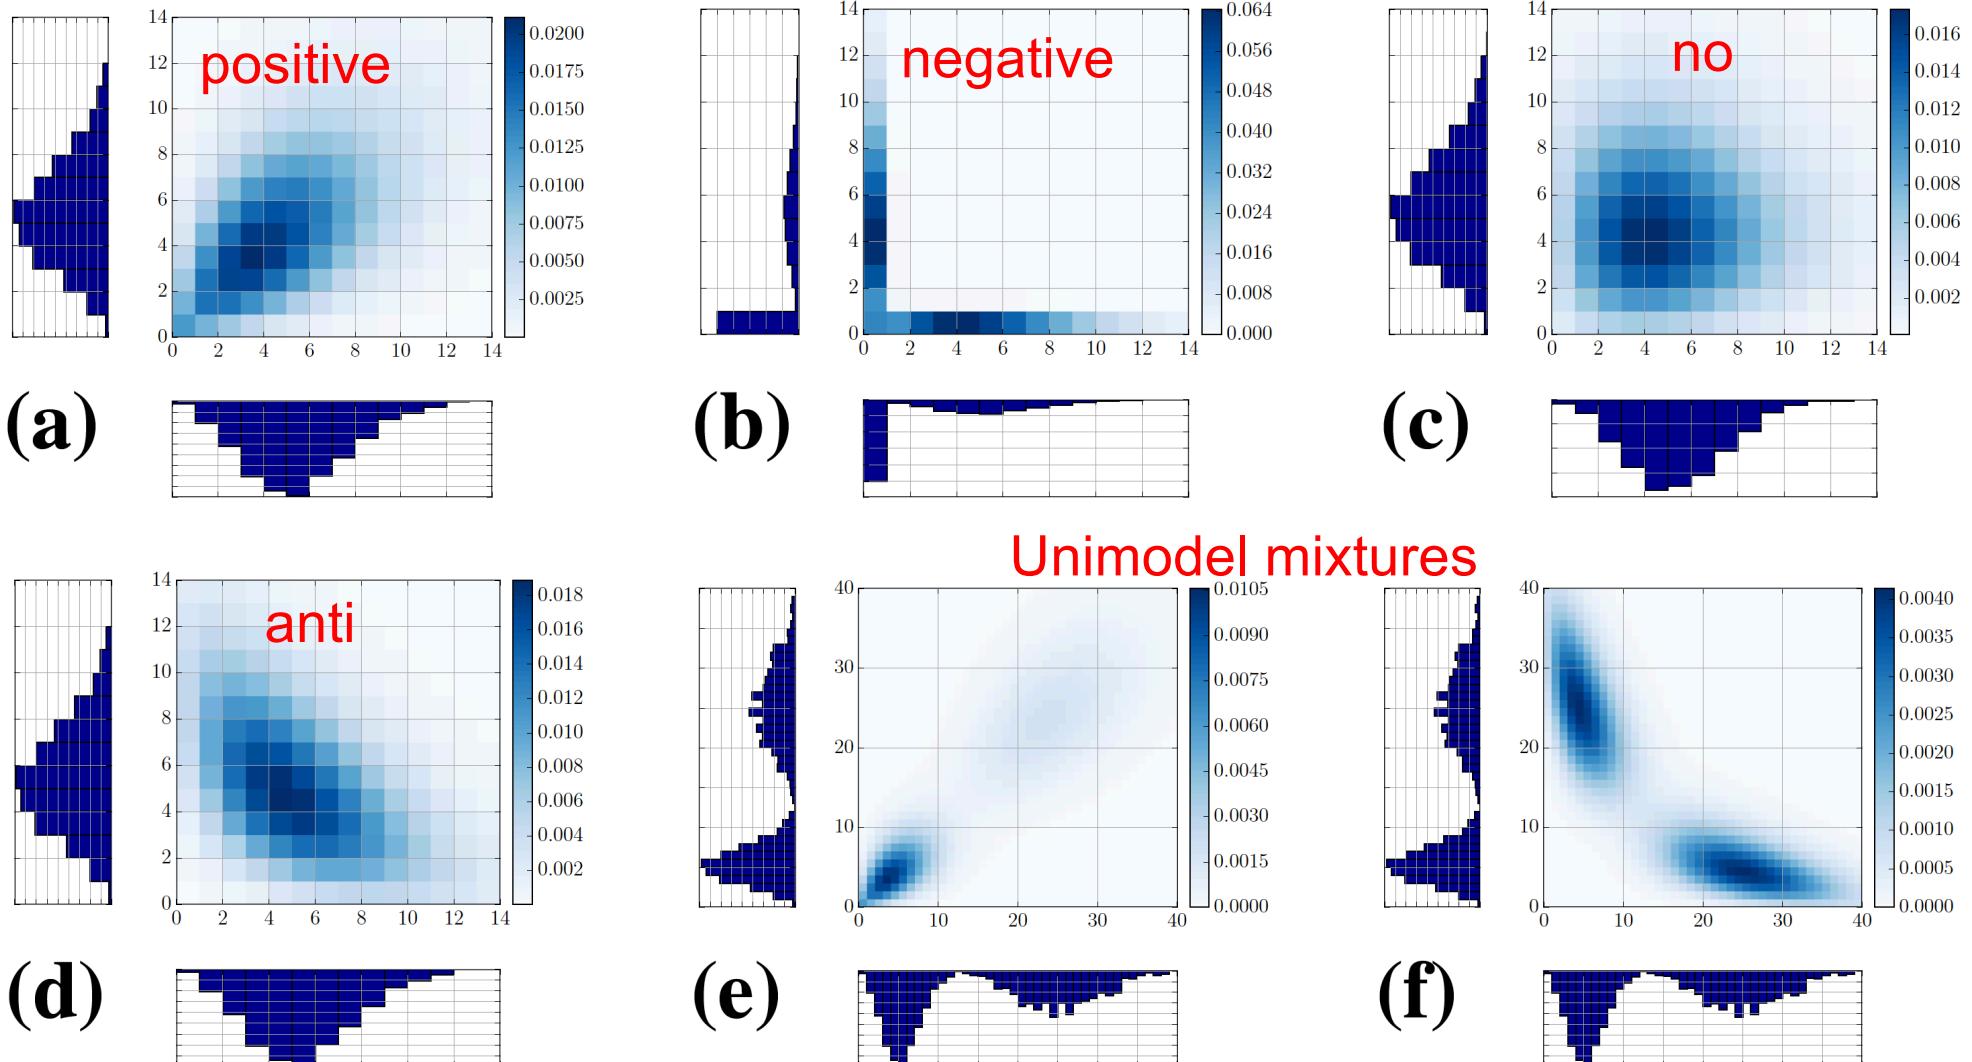
# Inference on Devices



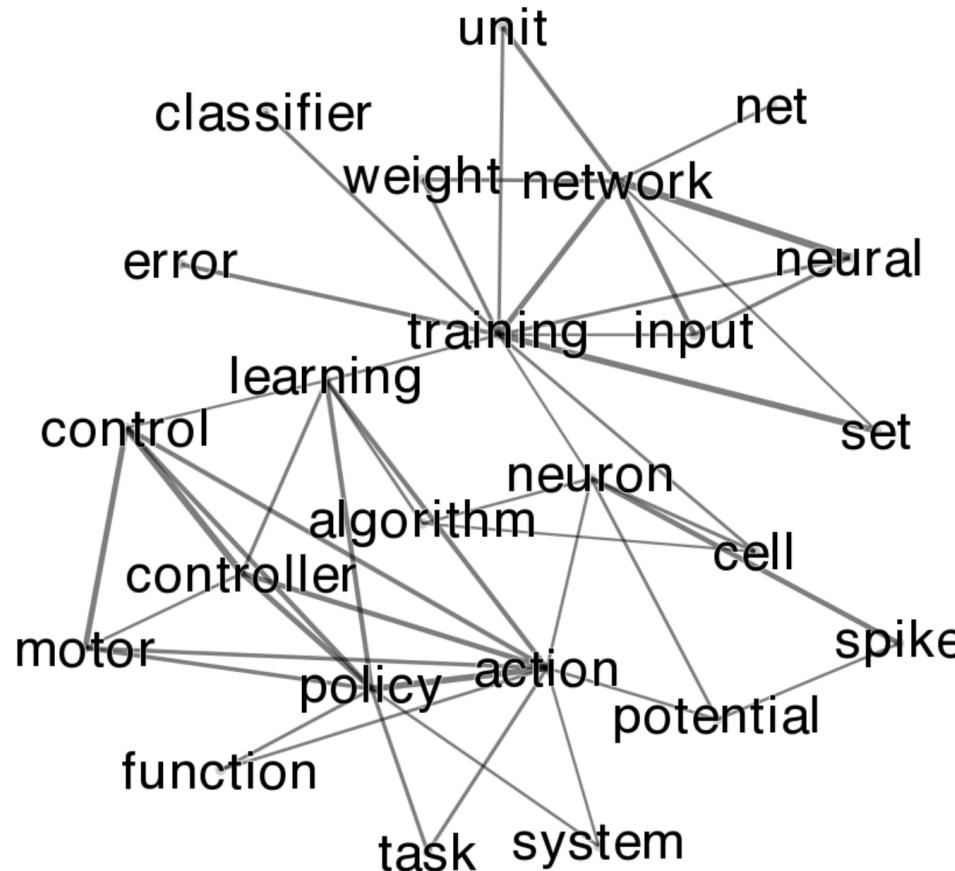
Ultimately, they may even be compiled into simple, flat, library-free code suitable for embedding in real-time applications and devices

SPNs encode polynomials over univariate distributions to build multivariate distributions. They can be fed into symbolic evaluation methods for inference

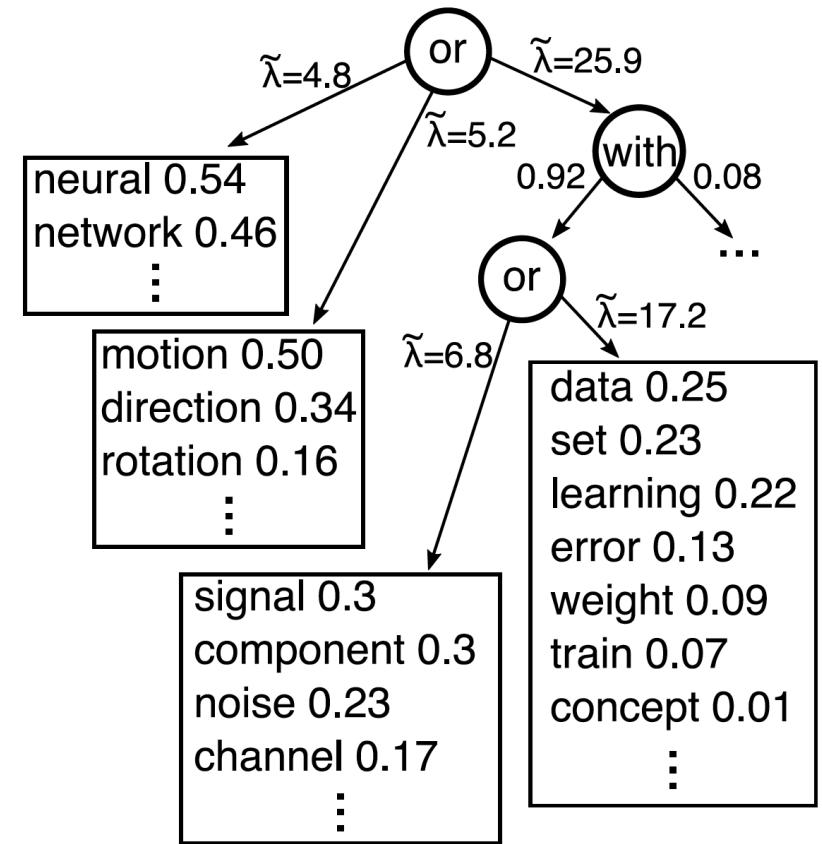
# Flexible Dependencies even for challenging distributions, here multivariate Poisson distribution



# Multiple views on your data due to efficient inference



Mutual Information  
(NIPS corpus)



Poisson Multinomial SPN  
= hierarchical topic model



# Poisson SPNs provide topics: From Topic Models to Poisson Mean Fields

[Bishop, Fienberg, Holland. Discrete Multivariate Analysis: Theory and Practice. Springer, 2007]

Topic-word distributions of LDA can be viewed as a product of independent Poissons

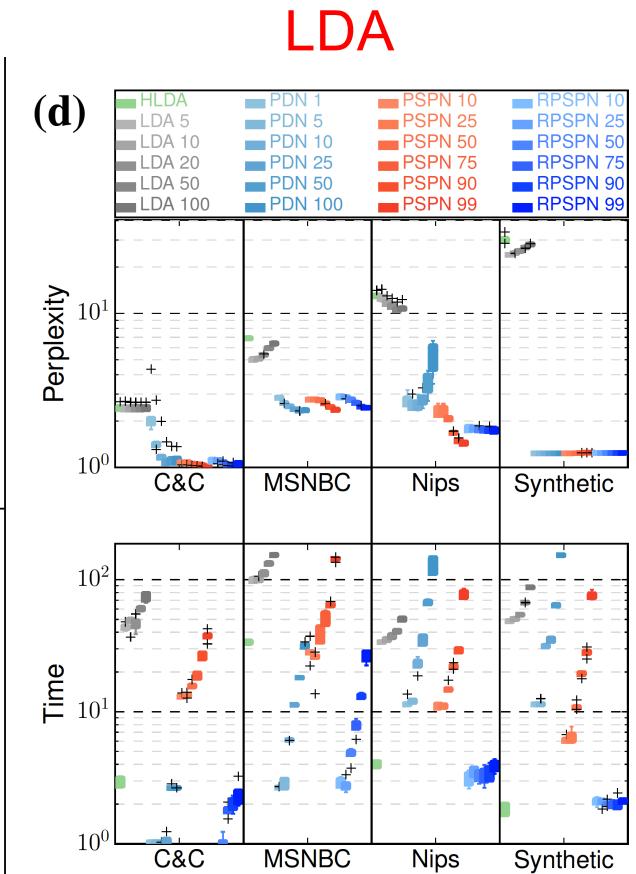
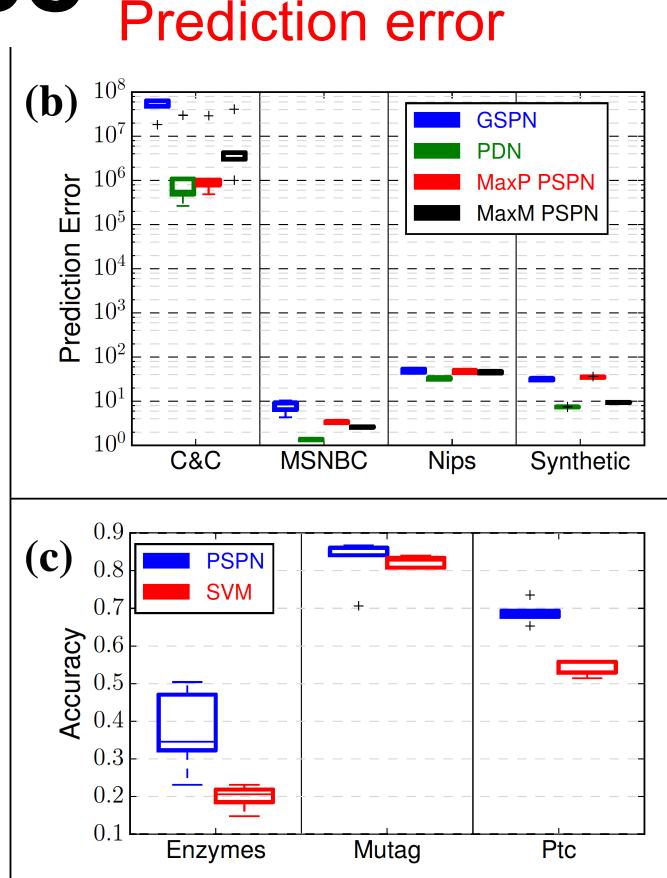
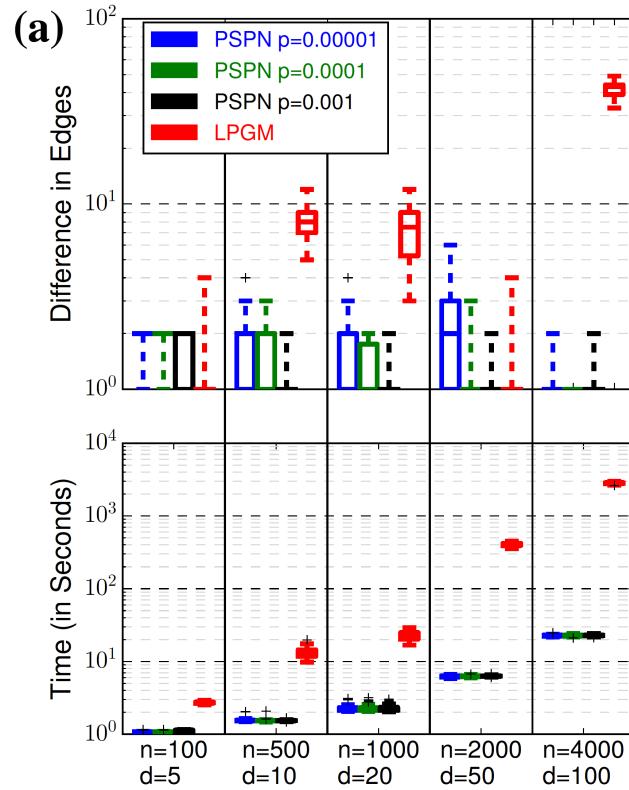
$$\begin{aligned}
 & \Pr_{\text{Poiss}} \left( \tilde{x} \mid \tilde{\lambda} \right) \Pr_{\text{Mult}} \left( \mathbf{x} \mid \theta = (\lambda_1, \dots, \lambda_p) / \tilde{\lambda}, N = \tilde{x} \right) \\
 &= \frac{e^{-\tilde{\lambda}}}{\tilde{x}!} \tilde{\lambda}^{\tilde{x}} \frac{\tilde{x}!}{\prod_{s=1}^p x_s!} \prod_{s=1}^p \left( \frac{\lambda_s}{\tilde{\lambda}} \right)^{x_s} \\
 &= \frac{\tilde{x}!}{\tilde{x}!} \frac{e^{-\tilde{\lambda}}}{\prod_{s=1}^p x_s!} \prod_{s=1}^p \left( \frac{\tilde{\lambda} \lambda_s}{\tilde{\lambda}} \right)^{x_s} \\
 &= \Pr_{\text{Ind. Poiss}} \left( \mathbf{x} \mid \lambda_1, \dots, \lambda_p \right) = \prod_{s=1}^p \frac{e^{-\lambda_s}}{x_s!} \lambda_s^{x_s}
 \end{aligned}$$



# Competitive Predictive Performance



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



Network discovery

Deep Graph Kernels  
(Weisfeiler Lehman)



# Random sum-product networks



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



UNIVERSITY OF  
CAMBRIDGE

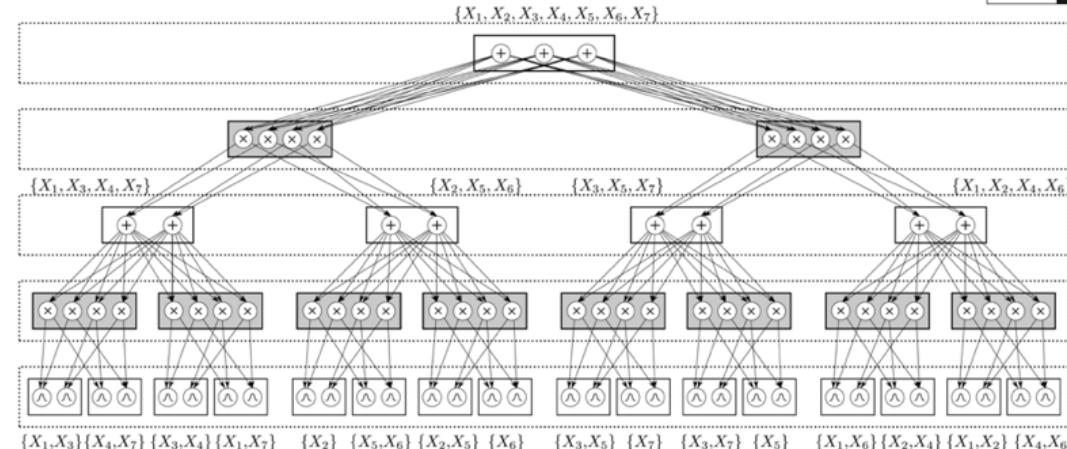


Max Planck Institute for  
Intelligent Systems

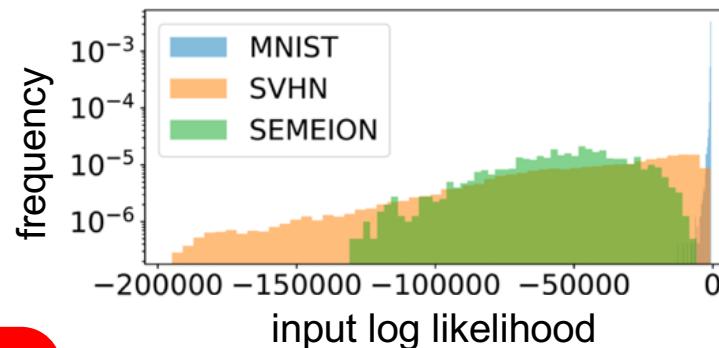


TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

UBER AI Labs



	RAT-SPN	MLP	vMLP
Accuracy	MNIST 98.19 (8.5M)	98.32 (2.64M)	98.09 (5.28M)
	F-MNIST 89.52 (0.65M)	90.81 (9.28M)	89.81 (1.07M)
	20-NG 47.8 (0.37M)	49.05 (0.31M)	48.81 (0.16M)
Cross-Entropy	MNIST 0.0852 (17M)	0.0874 (0.82M)	0.0974 (0.22M)
	F-MNIST 0.3525 (0.65M)	0.2965 (0.82M)	0.325 (0.29M)
	20-NG 1.6954 (1.63M)	1.6180 (0.22M)	1.6263 (0.22M)



SPNs can have  
similar predictive  
performances as  
(simple) DNNs

SPNs can distinguish the  
datasets

Conference on Uncertainty in Artificial Intelligence  
Tel Aviv, Israel  
July 22 - 25, 2019

uai2019

Similar to Random  
Forests, build a random  
SPN structure. This can  
be done in an informed  
way or completely at  
random



SPNs know when they do  
not know by design

# Einsum Networks

[Peharz, Lang, Vergari, Stelzner, Molina, Trapp, Van den Broeck, Kersting, Ghahramani ICML 2020]



TU/e



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

UBER AI Labs

UCLA



ICML | 2020

Thirty-seventh International  
Conference on Machine Learning



(a) Real SVHN images.



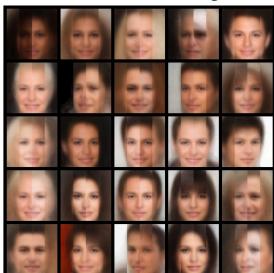
(b) EiNet SVHN samples.



(c) Real images (top), covered images, and EiNet reconstructions



(d) Real Celeba samples.

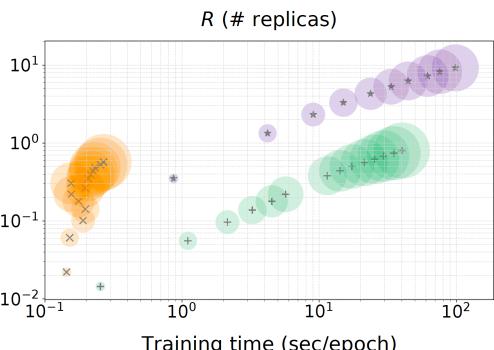
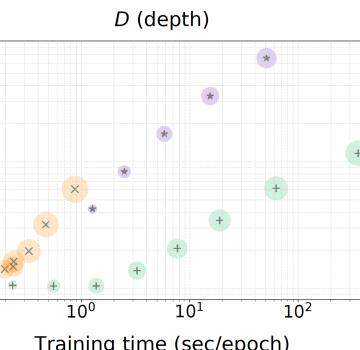
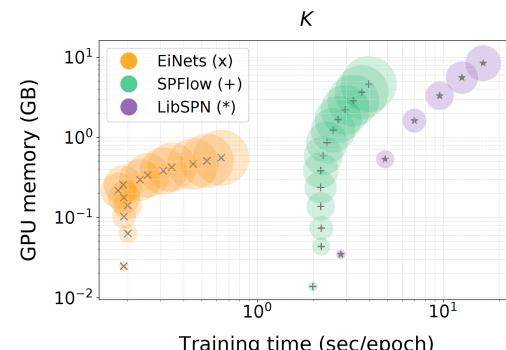
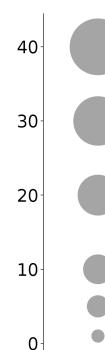
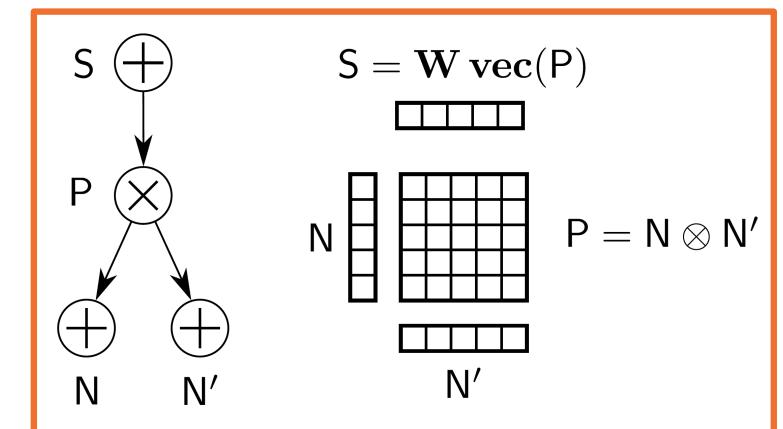


(e) EiNet Celeba samples.



(f) Real images (top), covered images, and EiNet reconstructions

Compilation of PCs into  
efficient to evaluate  
computational graphs



# Unsupervised scene understanding

[Stelzner, Peharz, Kersting ICML 2019, Best Paper Award at TPM@ICML2019] <https://github.com/stelzner/supair>

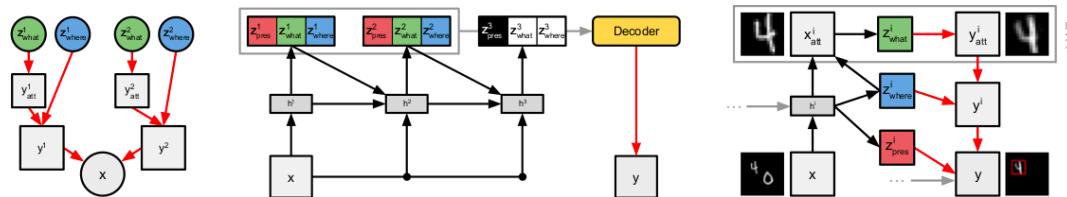


TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

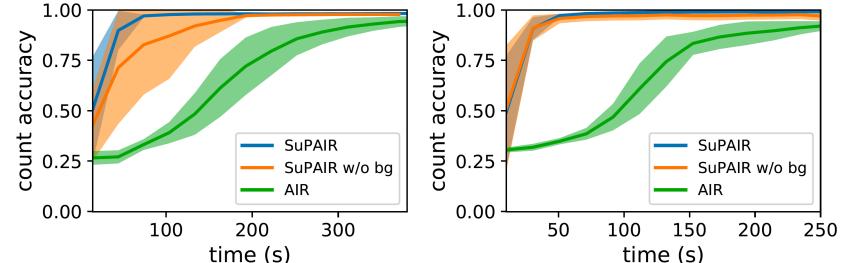
**ICML | 2019**

Thirty-sixth International Conference on  
Machine Learning

Consider e.g. unsupervised scene  
understanding using a generative model  
implemented in a neural fashion

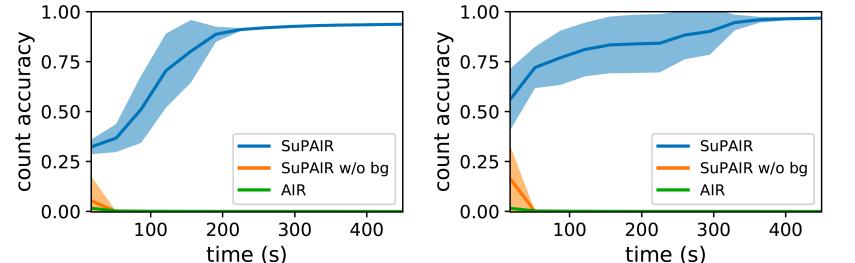


[Attend-Infer-Repeat (AIR) model, Hinton et al. NIPS 2016]



(a) MNIST

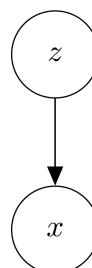
(b) Sprites



(c) Noisy MNIST

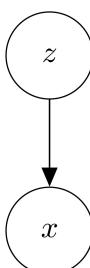
(d) Grid MNIST

VAE

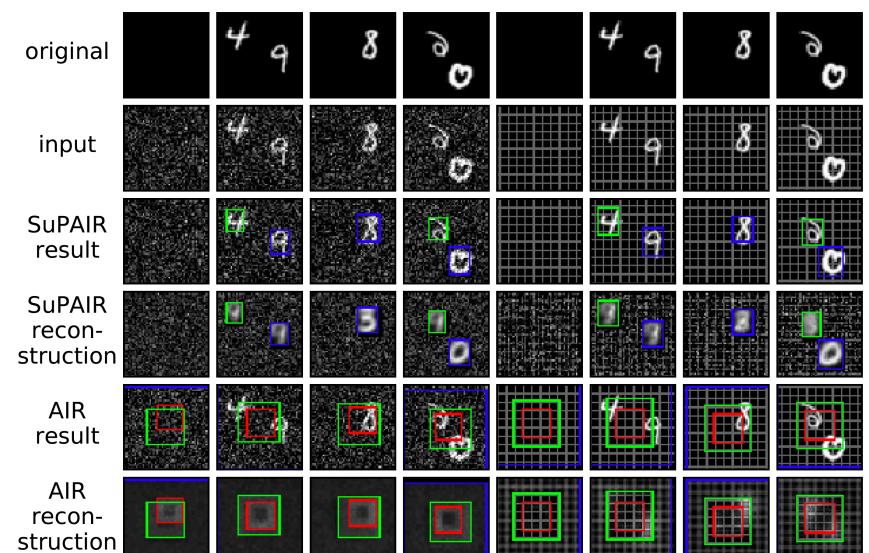


**Replace VAE  
by SPN as  
object model**

SPN



- infinite mixture model
- intractable density
- intractable posterior
- “large” but finite mixture model
- tractable density
- tractable marginals [Peharz et al., 2015]
- tractable posterior [Vergari et al., 2017]



# Unsupervised physics learning

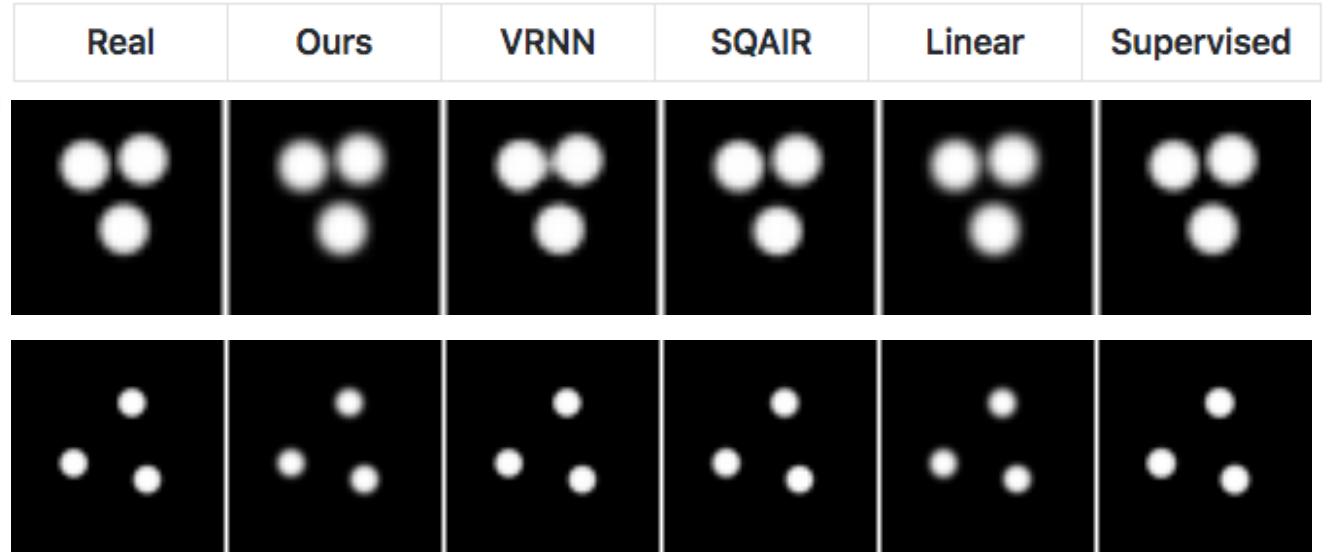
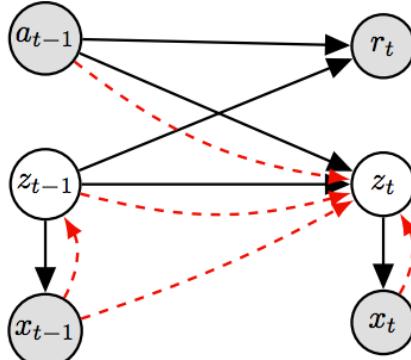
[Kossa, Stelzner, Hussing, Voelcker, Kersting ICLR 2020]

ICLR | 2020

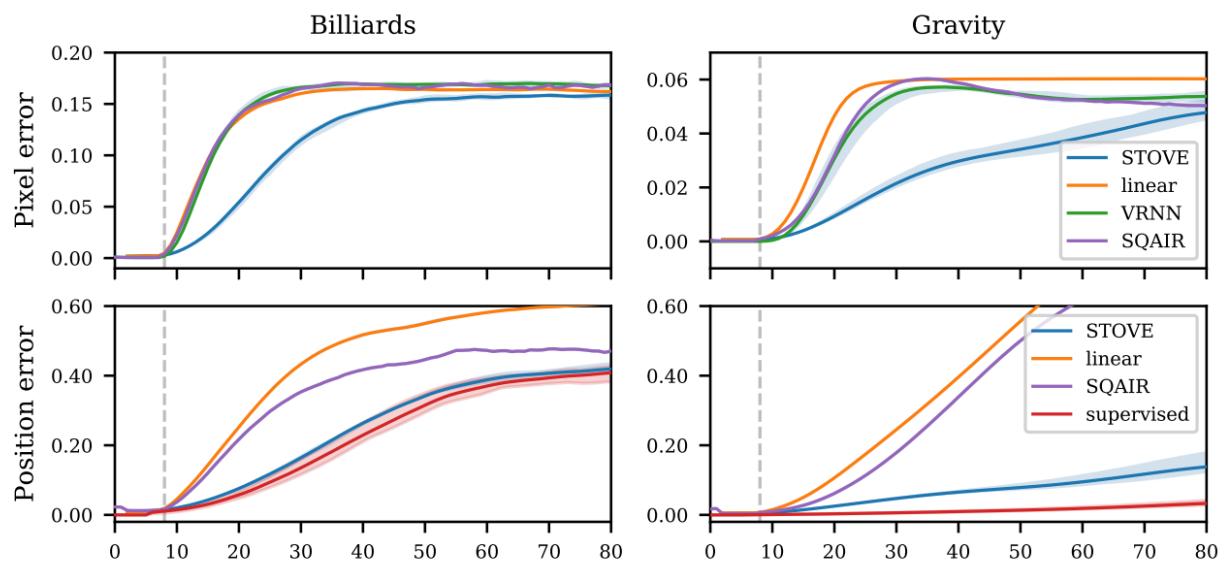
Eighth International Conference on  
Learning Representations



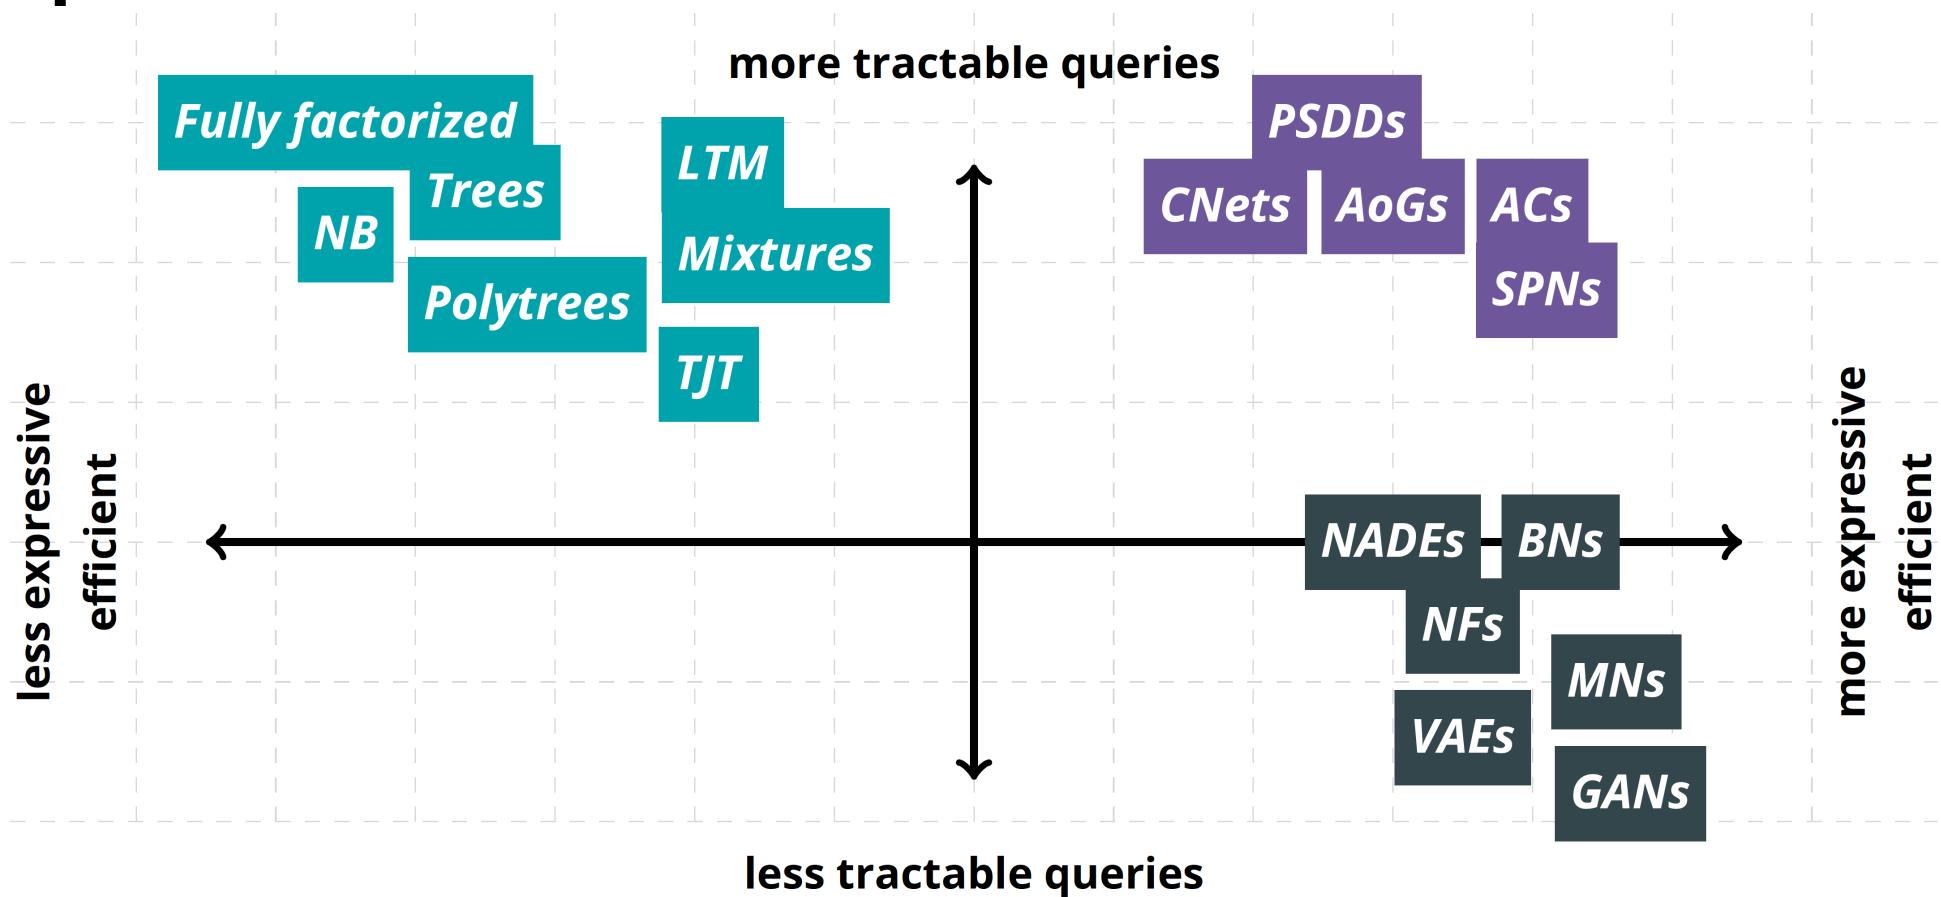
TECHNISCHE  
UNIVERSITÄT  
DARMSTADT



putting  
structure and  
tractable  
inference into  
deep models



**Sum-Product Networks are an example of tractable probabilistic models. There are many others, and one typically speaks of probabilistic circuits**



# What have we learnt about SPNs?

## Sum-product networks (SPNs)

- DAG of sums and products
- They are instances of Arithmetic Circuits (ACs)
- Compactly represent partition function
- Learn many layers of hidden variables

## Efficient marginal inference

Easy learning

**Can outperform well-known alternatives**

