

Statistical Machine Learning

Lecture 09: Classification

Kristian Kersting

TU Darmstadt

Summer Term 2020

Today's Objectives

- Make you understand how to do build a discriminative classifier!
- Covered Topics:
 - Discriminant Functions
 - Multi-Class Classification
 - Fisher Discriminate Analysis
 - Perceptrons
 - Logistic Regression

Outline

- 1. Discriminant Functions**
- 2. Fisher Discriminant Analysis**
- 3. Perceptron Algorithm**
- 4. Logistic Regression**
- 5. Wrap-Up**

Outline

1. Discriminant Functions

2. Fisher Discriminant Analysis

3. Perceptron Algorithm

4. Logistic Regression

5. Wrap-Up

Reminder of Bayesian Decision Theory

- We want to find the **a-posteriori probability** (posterior) of the class C_k given the observation (feature) x

$$p(C_k | x) = \frac{p(x | C_k)p(C_k)}{p(x)} = \frac{p(x | C_k)p(C_k)}{\sum_j p(x | C_j)p(C_j)}$$

- $p(C_k | x)$ - class posterior
- $p(x | C_k)$ - class-conditional probability (likelihood)
- $p(C_k)$ - class prior
- $p(x)$ - normalization term

Reminder of Bayesian Decision Theory

■ Decision rule

- Decide C_1 if $p(C_1 | x) > p(C_2 | x)$

- Equivalent to

$$p(x | C_1) p(C_1) > p(x | C_2) p(C_2) \equiv \frac{p(x | C_1)}{p(x | C_2)} > \frac{p(C_2)}{p(C_1)}$$

- A classifier obeying this rule is called a **Bayes optimal classifier**

Reminder of Bayesian Decision Theory

■ Current approach

- $p(C_k | x) = p(x | C_k) p(C_k) / p(x)$
- Model and estimate the class-conditional density $p(x | C_k)$ and the class prior $p(C_k)$
- Compute posterior $p(C_k | x)$
- Minimize the error probability by maximizing $p(C_k | x)$

■ New approach

- Directly encode the *decision boundary*
- Without modeling the densities directly
- Still minimize the error probability

Discriminant Functions

- Formulate classification using comparisons

- Discriminant functions

$$y_1(x), \dots, y_K(x)$$

- Classify x as class C_k iff

$$y_k(x) > y_j(x) \quad \forall j \neq k$$

- More formally, a **discriminant** maps a vector \mathbf{x} to one of the K available classes

Discriminant Functions

- Example of discriminant functions from the Bayes classifier

$$y_k(x) = p(C_k | x)$$

$$y_k(x) = p(x | C_k) p(C_k)$$

$$y_k(x) = \log p(x | C_k) + \log p(C_k)$$

Discriminant Functions

- Base case with 2 classes

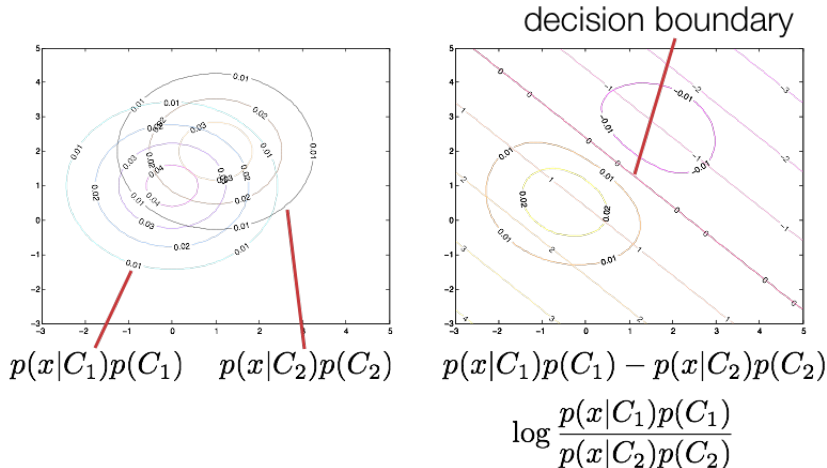
$$\begin{aligned}y_1(x) &> y_2(x) \\ y_1(x) - y_2(x) &> 0 \\ y(x) &> 0\end{aligned}$$

- Example from the Bayes classifier

$$\begin{aligned}y(x) &= p(C_1 | x) - p(C_2 | x) \\ y(x) &= \log \frac{p(x | C_1)}{p(x | C_2)} + \log \frac{p(C_1)}{p(C_2)}\end{aligned}$$

Example - Bayes Classifier

- Base case with 2 classes and Gaussian class-conditionals



Linear Discriminant Functions

- Base case with 2 classes

$y(\mathbf{x}) > 0$ decide class 1, otherwise class 2

- Simplest case: linear decision boundary
 - In **linear discriminants**, the decision surfaces are (hyper)planes
 - **Linear Discriminant Function**


$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Where \mathbf{w} is the normal vector and w_0 the offset


Linear Discriminant Functions

■ Illustration of the 2D case

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0, \quad \mathbf{x} = \begin{bmatrix} x_1 & x_2 \end{bmatrix}^T$$



Linear Discriminant Functions



Linearly separable



Not linearly separable

Discriminant Functions

- Why might we want to use discriminant functions?



- We could easily fit the class-conditionals using Gaussians and use a Bayes classifier

Discriminant Functions

- How about now? Do these points matter for making the decision between the two classes?



Distribution-free Classifiers



- We do not necessarily need to model all the details of the class-conditional distributions to come up with a good decision boundary. (The class-conditionals may have many intricacies that do not matter at the end of the day)



- If we can learn where to place the decision boundary directly, we can avoid some of the complexity
- It would be unwise to believe that such classifiers are inherently superior to probabilistic ones. We shall see why later...

Multi-Class Case

- What if we constructed a multi-class classifier from several 2-class classifiers?




- If we base our decision rule on binary decisions, this may lead to ambiguities

Multi-Class Case - Better Solution

- Use a discriminant function to encode how strongly we believe in each class

$$y_1(x), \dots, y_K(x)$$

- **Decision rule:** Decide k if $y_k(x) > y_j(x) \quad \forall j \neq k$



- If the discriminant functions are linear, the decision regions are **connected** and **convex**

Outline

1. Discriminant Functions

2. Fisher Discriminant Analysis

3. Perceptron Algorithm


4. Logistic Regression

5. Wrap-Up

Linear Discriminant Functions

■ Illustration of the 2D case

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0, \quad \mathbf{x} = \begin{bmatrix} x_1 & x_2 \end{bmatrix}^T$$



First Attempt: Least Squares

- Try to achieve a certain value of the discriminative function

$$y(\mathbf{x}) = +1 \Leftrightarrow \mathbf{x} \in C_1$$

$$y(\mathbf{x}) = -1 \Leftrightarrow \mathbf{x} \in C_2$$

- Training data inputs: $X = \{\mathbf{x}_1 \in \mathbb{R}^d, \dots, \mathbf{x}_n\}$
- Training data labels: $Y = \{y_1 \in \{-1, +1\}, \dots, y_n\}$
- Linear Discriminant Function
 - Try to enforce $\mathbf{x}_i^T \mathbf{w} + w_0 = y_i, \quad \forall i = 1, \dots, n$
 - There is one linear equation for each training data point/label pair

First Attempt: Least Squares

■ Linear system of equations

$$\mathbf{x}_i^T \mathbf{w} + w_0 = y_i, \quad \forall i = 1, \dots, n$$

■ Define $\hat{\mathbf{x}}_i = (\mathbf{x}_i \ 1)^T \in \mathbb{R}^{d \times 1}$, $\hat{\mathbf{w}} = (\mathbf{w} \ w_0)^T \in \mathbb{R}^{d \times 1}$

■ Rewrite the equation system

$$\hat{\mathbf{x}}_i^T \hat{\mathbf{w}} = y_i, \quad \forall i = 1, \dots, n$$

■ In **matrix-vector notation** we have

$$\hat{\mathbf{X}}^T \hat{\mathbf{w}} = \mathbf{y}$$

■ With $\hat{\mathbf{X}} = [\hat{\mathbf{x}}_1, \dots, \hat{\mathbf{x}}_n] \in \mathbb{R}^{d \times n}$ and $\mathbf{y} = [y_1, \dots, y_n]^T$

First Attempt: Least Squares

$$\hat{\mathbf{X}}^T \hat{\mathbf{w}} = \mathbf{y}$$

- An overdetermined system of equations
- There are n equations and $d + 1$ unknowns

First Attempt: Least Squares

- Look for the least squares solution


$$\begin{aligned}\hat{\mathbf{w}}^* &= \arg \min_{\hat{\mathbf{w}}} \left\| \hat{\mathbf{X}}^T \hat{\mathbf{w}} - \mathbf{y} \right\|^2 \\ &= \arg \min_{\hat{\mathbf{w}}} \left(\hat{\mathbf{X}}^T \hat{\mathbf{w}} - \mathbf{y} \right)^T \left(\hat{\mathbf{X}}^T \hat{\mathbf{w}} - \mathbf{y} \right) \\ &= \arg \min_{\hat{\mathbf{w}}} \hat{\mathbf{w}}^T \hat{\mathbf{X}} \hat{\mathbf{X}}^T \hat{\mathbf{w}} - 2 \mathbf{y}^T \hat{\mathbf{X}}^T \hat{\mathbf{w}} + \mathbf{y}^T \mathbf{y}\end{aligned}$$

$$\nabla_{\hat{\mathbf{w}}} \left(\hat{\mathbf{w}}^T \hat{\mathbf{X}} \hat{\mathbf{X}}^T \hat{\mathbf{w}} - 2 \mathbf{y}^T \hat{\mathbf{X}}^T \hat{\mathbf{w}} + \mathbf{y}^T \mathbf{y} \right) = 0$$


$$\hat{\mathbf{w}} = \underbrace{\left(\hat{\mathbf{X}} \hat{\mathbf{X}}^T \right)^{-1}}_{\text{pseudo-inverse}} \hat{\mathbf{X}} \mathbf{y}$$

First Attempt: Least Squares

- Problem: Least-squares is very sensitive to outliers




Without outliers
least-squares discriminant
works



With outliers least-squares
discriminant breaks down

Fisher's Linear Discriminant

- Take a different view on linear classification
- Find a linear projection of our data and classify the projected values



- The same thing as a linear discriminant function
 - Projection: $y = \mathbf{w}^T \mathbf{x}$
 - Checking against a threshold: $\mathbf{w}^T \mathbf{x} \geq -w_0$ or $\mathbf{w}^T \mathbf{x} + w_0 \geq 0$

Fisher's Linear Discriminant

- What is a **good projection \mathbf{w}** ?
 - Idea: **Maximize the “distance” between the two classes to allow for a good separation**
- First attempt: **Maximize the distance between the class means**

$$\mathbf{m}_1 = \frac{1}{|C_1|} \sum_{i \in C_1} \mathbf{x}_i \quad \mathbf{m}_2 = \frac{1}{|C_2|} \sum_{i \in C_2} \mathbf{x}_i$$

- Projection of the means

$$m_1 = \mathbf{w}^T \mathbf{m}_1 \quad m_2 = \mathbf{w}^T \mathbf{m}_2$$

- Maximize squared distance between means

$$\max (m_1 - m_2)^2$$

Fisher's Linear Discriminant

- Maximize squared distance between means

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} (\mathbf{w}^T \mathbf{m}_1 - \mathbf{w}^T \mathbf{m}_2)^2$$

- Obvious problem: Grows unboundedly with the norm of \mathbf{w}
- Obvious solution: Fix the norm of \mathbf{w}

$$\begin{aligned} \max_{\mathbf{w}} \quad & (\mathbf{w}^T \mathbf{m}_1 - \mathbf{w}^T \mathbf{m}_2)^2 \\ \text{s.t.} \quad & \|\mathbf{w}\|^2 = 1 \end{aligned}$$

- Constrained optimization problem!

Fisher's Linear Discriminant

$$\begin{aligned} \max_{\mathbf{w}} \quad & (\mathbf{w}^\top \mathbf{m}_1 - \mathbf{w}^\top \mathbf{m}_2)^2 \\ \text{s.t.} \quad & \|\mathbf{w}\|^2 = 1 \end{aligned}$$

- Necessary conditions


$$\begin{aligned} \nabla_{\mathbf{x}} f(\mathbf{x}) + \lambda \nabla_{\mathbf{x}} g(\mathbf{x}) &= 0 \\ 2(\mathbf{w}^\top \mathbf{m}_1 - \mathbf{w}^\top \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2) + 2\lambda \mathbf{w} &= 0 \end{aligned}$$

- It follows that

$$\mathbf{w} = \frac{\mathbf{m}_1 - \mathbf{m}_2}{\|\mathbf{m}_1 - \mathbf{m}_2\|}$$

Fisher's Linear Discriminant


■ Here's what we get



■ Obvious problem: **large class overlap**

Fisher's Linear Discriminant

- Here's what we **could** get



- Much better separation between classes
- How do we get this?
 - Idea: **Separate the means as far as possible while minimizing the variance of each class**

Fisher's Linear Discriminant

- Second (and final) attempt:
 - Define **within-class variances**:

$$s_1^2 = \sum_{n \in C_1} (\mathbf{w}^\top \mathbf{x}_n - m_1)^2 \quad s_2^2 = \sum_{n \in C_2} (\mathbf{w}^\top \mathbf{x}_n - m_2)^2$$

where $m_1 = \mathbf{w}^\top \mathbf{m}_1$ and $m_2 = \mathbf{w}^\top \mathbf{m}_2$

- **Fisher criterion**

$$\max_{\mathbf{w}} J(\mathbf{w}) = \frac{(m_1 - m_2)^2}{s_1^2 + s_2^2}$$

Fisher's Linear Discriminant

■ Fisher criterion

$$\max_{\mathbf{w}} J(\mathbf{w}) = \frac{(m_1 - m_2)^2}{s_1^2 + s_2^2}$$

■ Rewrite the numerator

$$\begin{aligned}(m_1 - m_2)^2 &= (\mathbf{w}^\top \mathbf{m}_1 - \mathbf{w}^\top \mathbf{m}_2)^2 \\&= (\mathbf{w}^\top (\mathbf{m}_1 - \mathbf{m}_2))^2 \\&= \mathbf{w}^\top \underbrace{(\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^\top}_{\substack{=: \mathbf{S}_B \\ \text{between-class covariance}}} \mathbf{w}\end{aligned}$$

Fisher's Linear Discriminant

■ Fisher criterion

$$\max_{\mathbf{w}} J(\mathbf{w}) = \frac{(m_1 - m_2)^2}{s_1^2 + s_2^2}$$

■ Rewrite the denominator

$$\begin{aligned}
 s_1^2 + s_2^2 &= \sum_{n \in C_1} (\mathbf{w}^\top \mathbf{x}_n - m_1)^2 + \sum_{n \in C_2} (\mathbf{w}^\top \mathbf{x}_n - m_2)^2 \\
 &= \sum_{n \in C_1} (\mathbf{w}^\top (\mathbf{x}_n - \mathbf{m}_1))^2 + \sum_{n \in C_2} (\mathbf{w}^\top (\mathbf{x}_n - \mathbf{m}_2))^2 \\
 &= \sum_{n \in C_1} \mathbf{w}^\top (\mathbf{x}_n - \mathbf{m}_1) (\mathbf{x}_n - \mathbf{m}_1)^\top \mathbf{w} + \sum_{n \in C_2} \mathbf{w}^\top (\mathbf{x}_n - \mathbf{m}_2) (\mathbf{x}_n - \mathbf{m}_2)^\top \mathbf{w} \\
 &= \mathbf{w}^\top \underbrace{\left[\sum_{n \in C_1} (\mathbf{x}_n - \mathbf{m}_1) (\mathbf{x}_n - \mathbf{m}_1)^\top + \sum_{n \in C_2} (\mathbf{x}_n - \mathbf{m}_2) (\mathbf{x}_n - \mathbf{m}_2)^\top \right]}_{\substack{=: \mathbf{S}_W \\ \text{within-class covariance}}} \mathbf{w}
 \end{aligned}$$

Fisher's Linear Discriminant

- Fisher criterion

$$\max_{\mathbf{w}} J(\mathbf{w}) = \frac{(m_1 - m_2)^2}{s_1^2 + s_2^2} = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

- Differentiating w.r.t. \mathbf{w} and setting to 0 we have

$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

- Since $(\mathbf{w}^T \mathbf{S}_B \mathbf{w})$ and $(\mathbf{w}^T \mathbf{S}_W \mathbf{w})$ are scalars, we have that

$$\mathbf{S}_W \mathbf{w} \parallel \mathbf{S}_B \mathbf{w}$$

where \parallel means collinearity

Fisher's Linear Discriminant

- Also, we know that

$$\mathbf{S}_B \mathbf{w} = (\mathbf{m}_1 - \mathbf{m}_2)(\mathbf{m}_1 - \mathbf{m}_2)^\top \mathbf{w} \implies \mathbf{S}_B \mathbf{w} \parallel (\mathbf{m}_1 - \mathbf{m}_2)$$

- Hence, we have

$$\begin{aligned} \mathbf{S}_W \mathbf{w} &\parallel (\mathbf{m}_1 - \mathbf{m}_2) \\ \mathbf{w} &\parallel \mathbf{S}_W^{-1} (\mathbf{m}_1 - \mathbf{m}_2) \end{aligned}$$

- Fisher's Linear Discriminant

$$\mathbf{w} \propto \mathbf{S}_W^{-1} (\mathbf{m}_1 - \mathbf{m}_2)$$

Fisher's Linear Discriminant

$$\mathbf{w} \propto \mathbf{S}_W^{-1} (\mathbf{m}_1 - \mathbf{m}_2)$$

- The Fisher linear discriminant only gives us a **projection**
 - We still need to find the threshold
 - E.g., use Bayes classifier with Gaussian class-conditionals
- **Bayes optimality**
 - Fisher's linear discriminant is Bayes optimal if the class-conditional distributions are equal, with diagonal covariance
- Essentially equivalent to **Linear Discriminant Analysis** (LDA)

Fisher's Linear Discriminant

- We won't go through this here, but Fisher's linear discriminant can be shown to be equivalent to a certain case of a least-squares linear classifier (see Bishop 4.1.5)
- **Problem with this method:** it is still very sensitive to noise!
- By The Way: This method is a **true classic** (it dates back to 1936)
 - Fisher, R.A., *The Use of Multiple Measurements in Taxonomic Problems*. Annals of Eugenics, 7: 179-188 (1936)

Outline

1. Discriminant Functions
2. Fisher Discriminant Analysis
- 3. Perceptron Algorithm**
4. Logistic Regression
5. Wrap-Up

New Strategy

- If our classes are **linearly separable**, we want to make sure that we find a **separating (hyper)plane**



- First such algorithm we will see
 - The perceptron algorithm [Rosenblatt, 1962]




Rosenblatt [1928-1971]

Perceptron Algorithm

■ Perceptron discriminant function

$$y(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$$

■ where $\text{sign}(x) = \{+1, x > 0; 0, x = 0; -1, x < 0\}$




Perceptron Algorithm


■ Perceptron Algorithm

- Initialize the weight vector \mathbf{w} and bias b
- For all pairs of data points (\mathbf{x}_i, y_i) , where $y_i \in \{-1, +1\}$, do
 - If \mathbf{x}_i is correctly classified, i.e., $y(\mathbf{x}_i) = y_i$, do nothing
 - Else if $y_i = 1$ update the parameters with
$$\mathbf{w} \leftarrow \mathbf{w} + \mathbf{x}_i, \quad b \leftarrow b + 1$$
 - Else if $y_i = -1$ update the parameters with
$$\mathbf{w} \leftarrow \mathbf{w} - \mathbf{x}_i, \quad b \leftarrow b - 1$$
- Repeat until convergence


Perceptron Algorithm - Intuition




Perceptron Algorithm - Intuition



Perceptron Algorithm - Intuition



Perceptron Algorithm - Intuition



Perceptron Algorithm

- Why does this algorithm work?
- We have an optimization problem


$$\begin{aligned}\max_{\mathbf{w}} J(\mathbf{w}) &= |\{x \in X : \langle \mathbf{w}, x \rangle < 0\}| \\ &= \sum_{x \in X : \langle \mathbf{w}, x \rangle < 0} \langle \mathbf{w}, x \rangle\end{aligned}$$

- And also a gradient method

$$\frac{\partial J}{\partial \mathbf{w}} = \sum_{x \in X : \langle \mathbf{w}, x \rangle < 0} x$$

But is the Perceptron Algorithm useful?


- How often is data linearly separable?
- A **simple failure** example is the XOR function



- History: Minsky & Papert [1969] criticized the perceptron for not being able to handle this case, which halted research on this and related techniques for decades

Other Feature Spaces

- It took a long time until people had realized that there is a simple way out
- Key idea: Transform the input data nonlinearly so that the problem becomes linearly separable!



- There is an important message to get out from this
 - Create features instead of learning from raw data
 - Neural networks do it *automagically* for you

Outline

1. Discriminant Functions
2. Fisher Discriminant Analysis
3. Perceptron Algorithm
- 4. Logistic Regression**
5. Wrap-Up

Generative vs. Discriminative

- There are two different views to solve the classification problem
- **Generative modelling**
 - We model the class-conditional distributions $p(x | C_2)$ and $p(x | C_1)$
 - We classify by computing the class posterior using Bayes' rule
 - E.g.: Naive Bayes
- **Discriminative modelling**
 - We model the class-posterior directly, e.g. $p(C_1 | x)$
 - Consequence: We only care about getting the classification right, and not whether we fit the class-conditional well
 - E.g.: Logistic Regression

Probabilistic Discriminative Models

- For now, we will write the class posterior using Bayes' rule


$$\begin{aligned} p(C_1 | \mathbf{x}) &= \frac{p(\mathbf{x} | C_1) p(C_1)}{p(\mathbf{x})} = \frac{p(\mathbf{x} | C_1) p(C_1)}{\sum_i p(\mathbf{x}, C_i)} \\ &= \frac{p(\mathbf{x} | C_1) p(C_1)}{\sum_i p(\mathbf{x} | C_i) p(C_i)} \\ &= \frac{p(\mathbf{x} | C_1) p(C_1)}{p(\mathbf{x} | C_1) p(C_1) + p(\mathbf{x} | C_2) p(C_2)} \\ &= \frac{1}{1 + p(\mathbf{x} | C_2) p(C_2) / (p(\mathbf{x} | C_1) p(C_1))} \\ &= \frac{1}{1 + \exp(-a)} = \sigma(a) \rightarrow \text{logistic sigmoid function} \end{aligned}$$

$$\text{with } a = \log \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_2)p(C_2)}$$

Sigmoid

■ Logistic / Sigmoid function

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$



[Wikipedia]

■ Sigmoid: 'S-shaped'

■ Squashes real numbers into the $[0, 1]$ interval

Probabilistic Discriminative Models

- Class posterior

$$p(C_1 | \mathbf{x}) = \sigma(a) \quad \text{with} \quad a = \log \frac{p(\mathbf{x} | C_1) p(C_1)}{p(\mathbf{x} | C_2) p(C_2)}$$

- Logistic regression

- Assume that a is given by a linear discriminant function

$$p(C_1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

- Find \mathbf{w} and w_0 so that the class-posterior is modeled best
- When is this an **appropriate assumption**?
 - When the class conditionals are Gaussians with equal covariance
 - But also for a number of other distributions
 - Some independence of the form of the class-conditionals

Logistic Regression

- Model the class posterior as

$$p(C_1 | \mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

- Maximize the likelihood

- Data (as always) is i.i.d. and define $y_i = \begin{cases} 0 & \mathbf{x}_i \text{ belongs to } C_1 \\ 1 & \mathbf{x}_i \text{ belongs to } C_2 \end{cases}$

$$\begin{aligned} p(Y | X; \mathbf{w}, w_0) &= \prod_{i=1}^N p(y_i | \mathbf{x}_i; \mathbf{w}, w_0) \\ &= \prod_{i=1}^N p(C_1 | \mathbf{x}_i; \mathbf{w}, w_0)^{1-y_i} p(C_2 | \mathbf{x}_i; \mathbf{w}, w_0)^{y_i} \\ &= \prod_{i=1}^N \sigma(\mathbf{w}^T \mathbf{x}_i + w_0)^{1-y_i} (1 - \sigma(\mathbf{w}^T \mathbf{x}_i + w_0))^{y_i} \end{aligned}$$

Logistic Regression

- We won't do the derivation here (see Bishop 4.3), but basically you can apply the logarithm to $p(Y | X; \mathbf{w}, w_0)$ and do gradient descent
- Similar to what we have seen in regression, we can get **more robust classifiers by incorporating priors and taking a Bayesian approach**
- Later, we will turn to a very different interpretation of this:
 - **Logistic regression as a neural network**

Outline

1. Discriminant Functions
2. Fisher Discriminant Analysis
3. Perceptron Algorithm
4. Logistic Regression
- 5. Wrap-Up**

5. Wrap-Up

You know now:

- What a Bayesian Optimal Classifier is
- What a discriminant function is
- How to formalize (with intuition and mathematically) the classification problem as linearly-separable
- How to compute the least squares solution for classification and why it fails
- What Fisher's Linear Discriminant is and how it differs from least-squares
- What the perceptron is, why it fails in the XOR problem and how to overcome it with feature spaces
- The difference between Generative and Discriminative modelling
- What logistic regression is

Self-Test Questions

- How do we get from Bayesian optimal decisions to discriminant functions?
- How to derive a discriminant function from a probability distribution?
- How to deal with more than two classes?
- What does linearly-separable mean?
- What is Fisher discriminant analysis? How does it relate to regression?
- Is Fisher's linear discriminant Bayes optimal?
- What are perceptrons? How can we train them?
- What is logistic regression? How to derive the parameter update rule?

Homework

- Reading Assignment for next week
 - Bishop 7.1.5 and 12.1
 - Murphy 6.5 and 12.2