

Lifted Belief Propagation [Jaimovich-UAI07, Singla-AAAI08, Kersting-UAI09]
Lifted Bisimulation/Mini-buckets [Sen-VLDB08, Sen-UAI09]
Lifted Importance Sampling [Gogate-UAI11, Gogate-AAAI12]
Lifted Relax, Compensate & Recover (Generalized BP) [VdB-UAI12]
Lifted MCMC [Niepert-UAI12, Niepert-AAAI13, Venugopal-NIPS12]
Lifted Variational Inference [Choi-UAI12, Bui-StarAI12]
Lifted MAP-LP [Mladenov-AISTATS14, Apsel-AAAI14] and many more ...

Statistical Relational AI

Probabilistic Relational Models

- Lifted Approximate Inference



**Kristian
Kersting**

Thanks to Rina Dechter, Luc De Raedt, Pedro Domingos, Peter Flach, Vibhav Gogate, Carlos Guestrin, Daphen Koller, Nir Friedman, Ray Mooney, Sriraam Natarajan, David Poole, Fabrizio Riguzzi, Dan Suciu, Guy van den Broeck, and many others for making their slides publically available

Relational Probabilistic Models also provide opportunities to speed up inference and learning



A simple example



Guy van den Broeck
UCLA

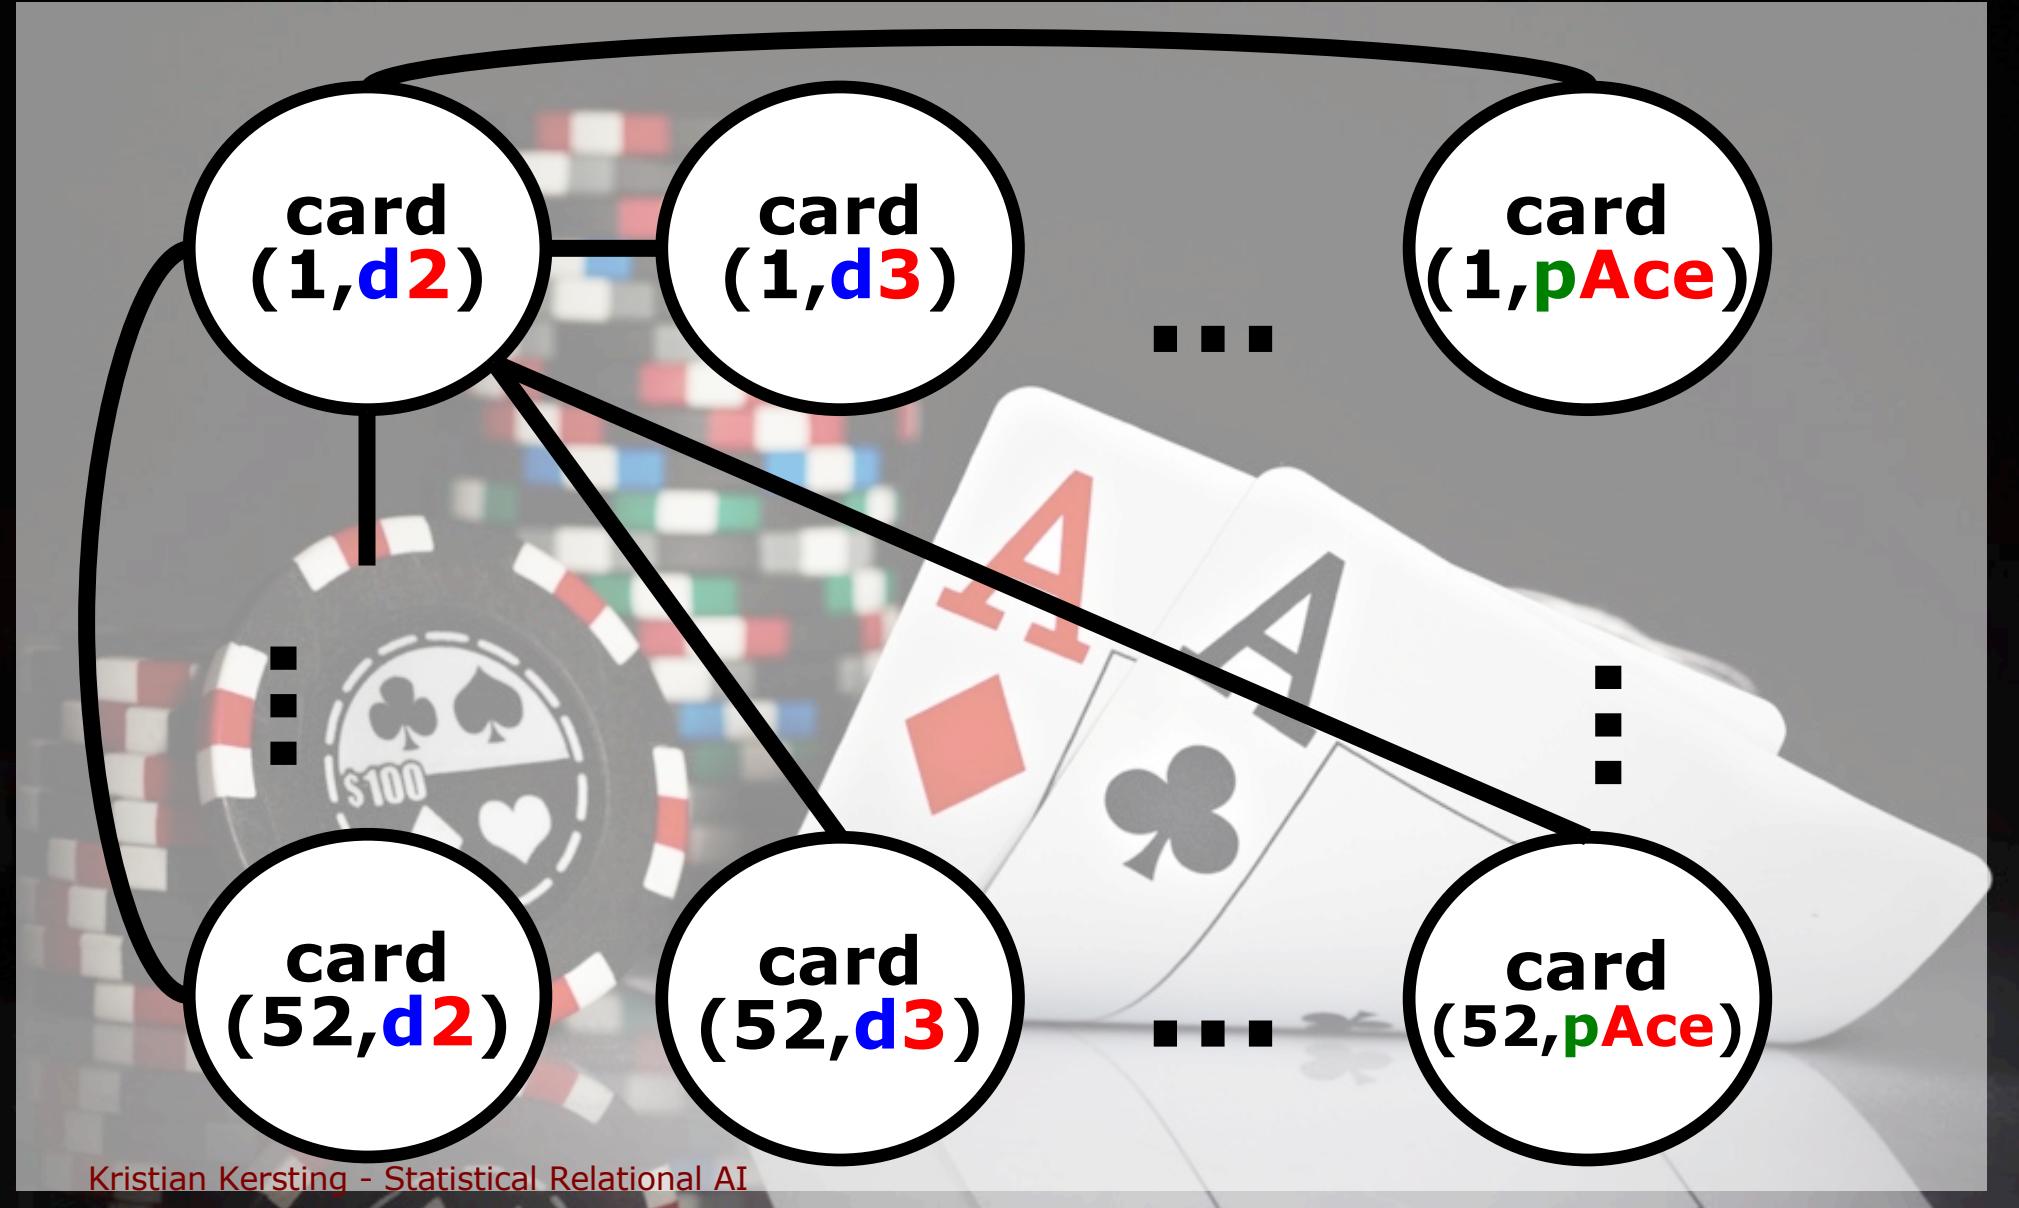
What is the probability that the first card of a randomly shuffled deck with 52 cards is an Ace?



A simple example



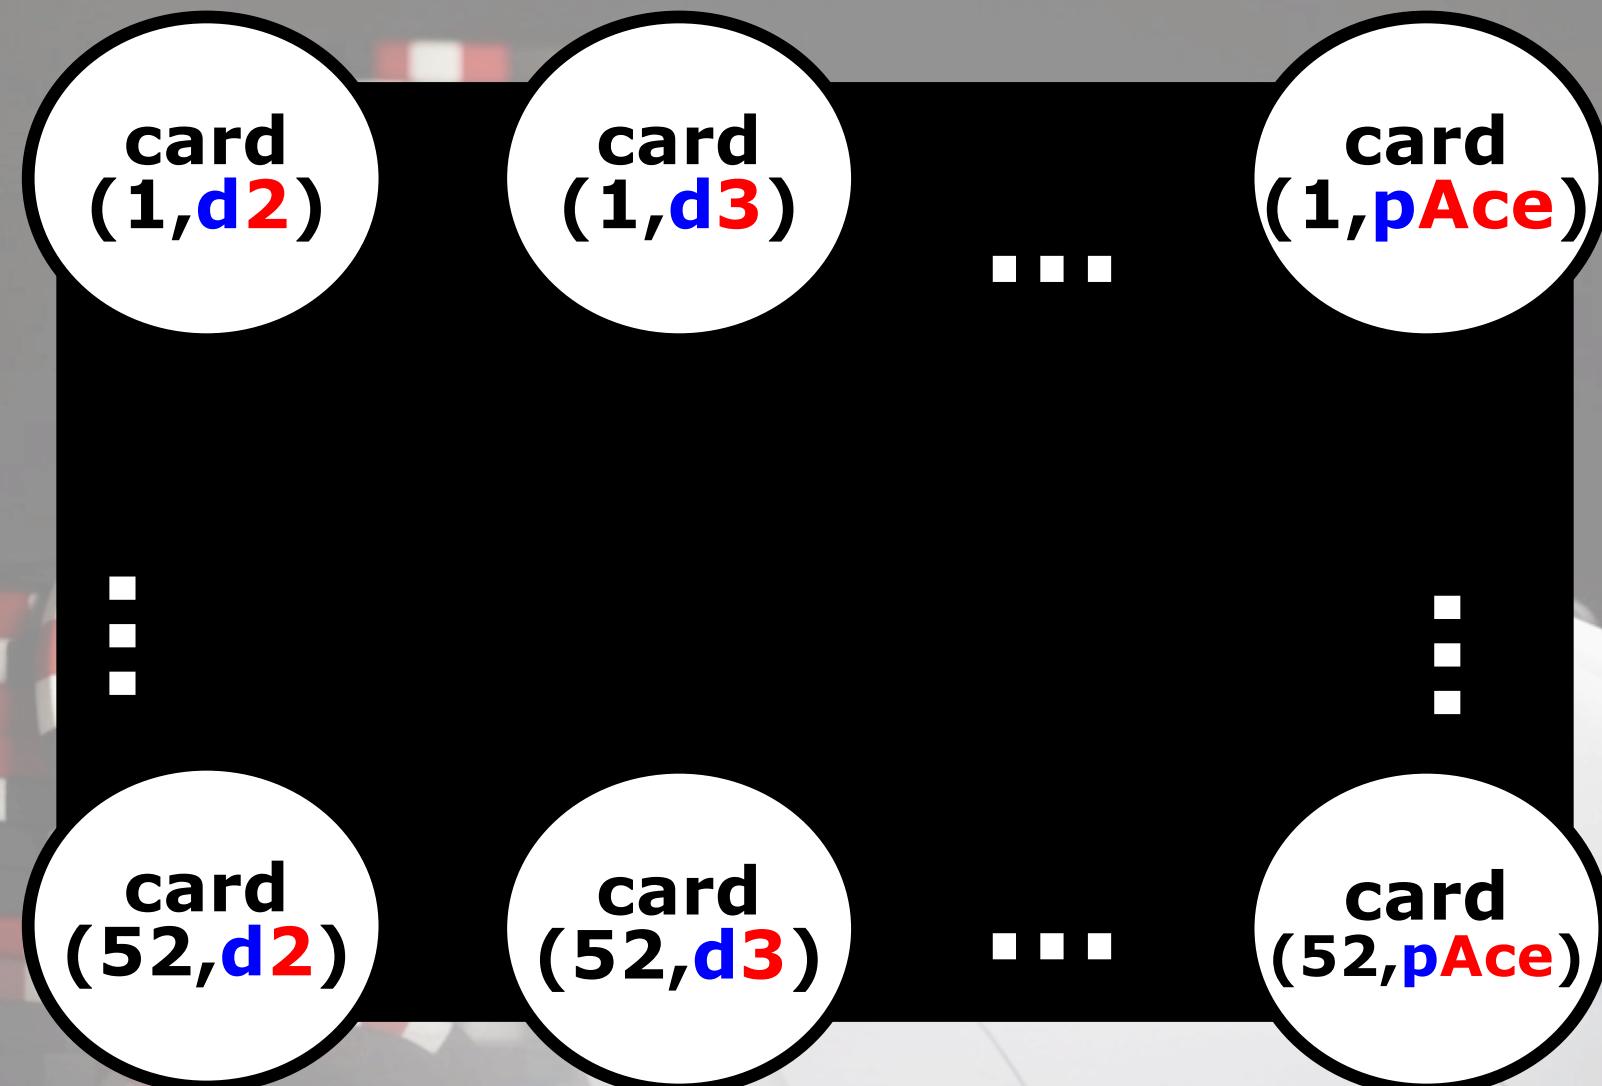
Guy van den Broeck
UCLA



A simple example



Guy van den Broeck
UCLA



A simple example



Guy van den Broeck
UCLA

card
(1,d2)

card
(1,d3)

...

card
(1,pAce)

No independencies.

Fully connected.

2^{2704} states

⋮
card
(52,d2)

card
(52,d3)

...

⋮
card
(52,pAce)

A simple example



Guy van den Broeck
UCLA

card
(1,d2)

card
(1,d3)

...

card
(1,pAce)

**A machine will not solve
the problem**

⋮
⋮
card
(52,d2)

card
(52,d3)

...

card
(52,pAce)

What are we missing ?

Faster modelling

Let's use a high-level language e.g.

Markov Logic Networks (MLNs)

w1: $\forall p, x, y: \text{card}(P, X), \text{card}(P, Y) \Rightarrow x = y$

w2: $\forall c, x, y: \text{card}(X, C), \text{card}(Y, C) \Rightarrow x = y$

and symmetry- and language-aware inference

Faster inference and learning

Positions and cards are exchangable but the machine is not aware of these symmetries



Lifted Approximate Inference

- There are many other, in particular, lifted inference approaches, in particular for exact inference such as
 - Knowledge compilation
 - Lifted recursive conditiong
- Here, we want to take a “group-theoretic” view on approximate lifted inference, also because it connects to Graph Neural Networks
- To do so, we start by lifting (loopy) belief propagation



[Singla, Domingos AAAI'08; Kersting, Ahmadi, Natarajan UAI'09; Ahmadi, Kersting, Mladenov, Natarajan MLJ'13]

Lifted Loopy Belief Propagation

Exploiting computational symmetries

**Big
Model**

Run
Loopy Belief Propagation

automatically
compressed



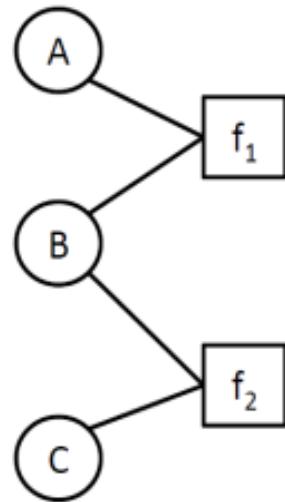
**Small
Model**

Run a modified
Loopy Belief Propagation

What are symmetries in (loop) belief propagation?



Compression: Coloring the graph

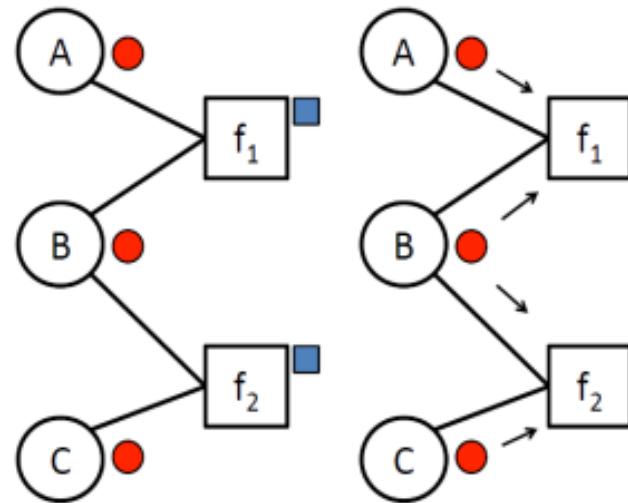


- **Color nodes according to the evidence you have**
 - No evidence, say **red**
 - State „one“, say **brown**
 - State „two“, say **orange**
 - ...
- **Color factors distinctively according to their equivalences**

For instance, assuming f_1 and f_2 to be identical and B appears at the second position within both, say **blue**



Compression: Pass the colors around

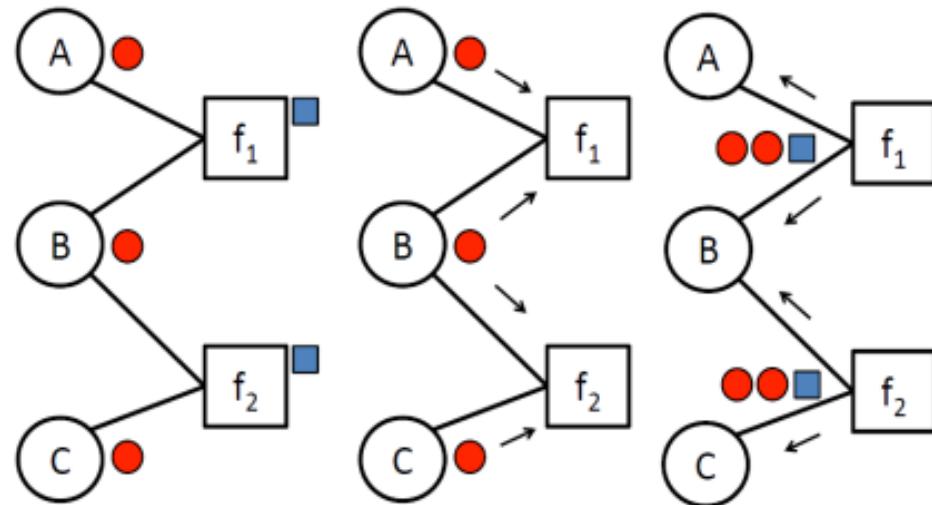


1. Each factor collects the colors of its neighboring nodes



Compression: Pass the colors around*

*can also be done at the „lifted“, i.e., relational level

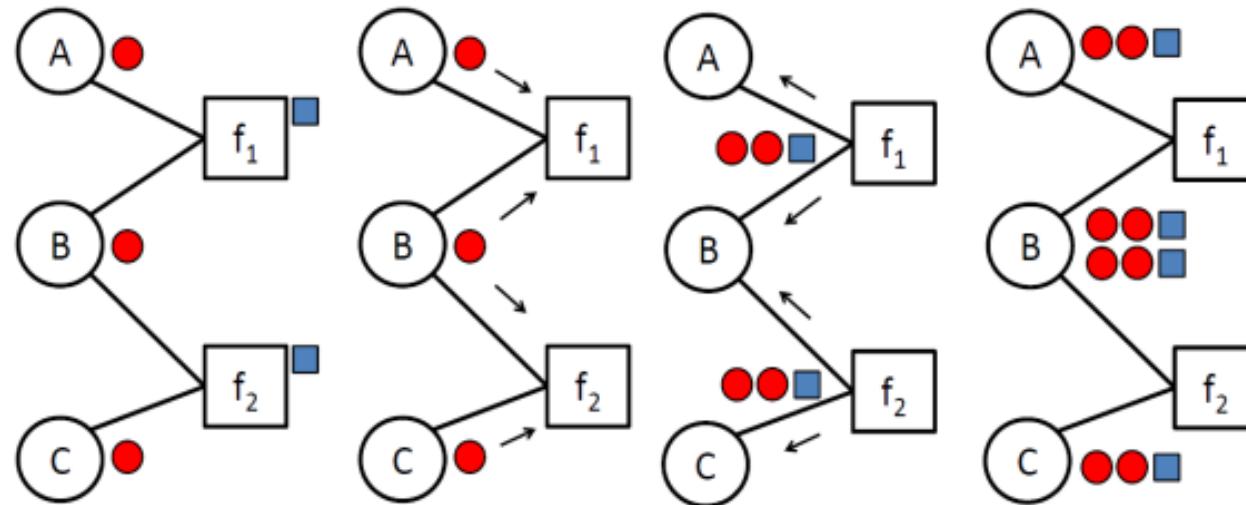


1. Each factor collects the colors of its neighboring nodes
2. Each factor „signs“ ist color signature with its own color



Compression: Pass the colors around*

*can also be done at the „lifted“, i.e., relational level

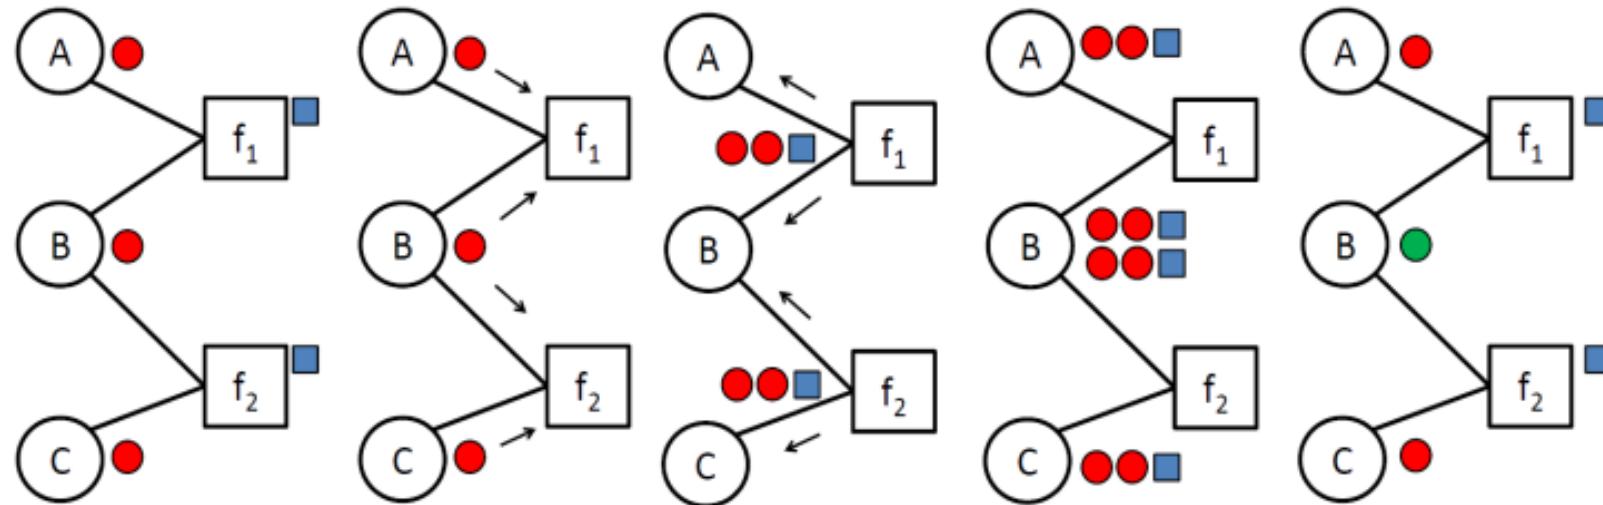


1. Each factor collects the colors of its neighboring nodes
2. Each factor „signs“ its color signature with its own color
3. Each node collects the signatures of its neighboring factors



Compression: Pass the colors around*

*can also be done at the „lifted“, i.e., relational level

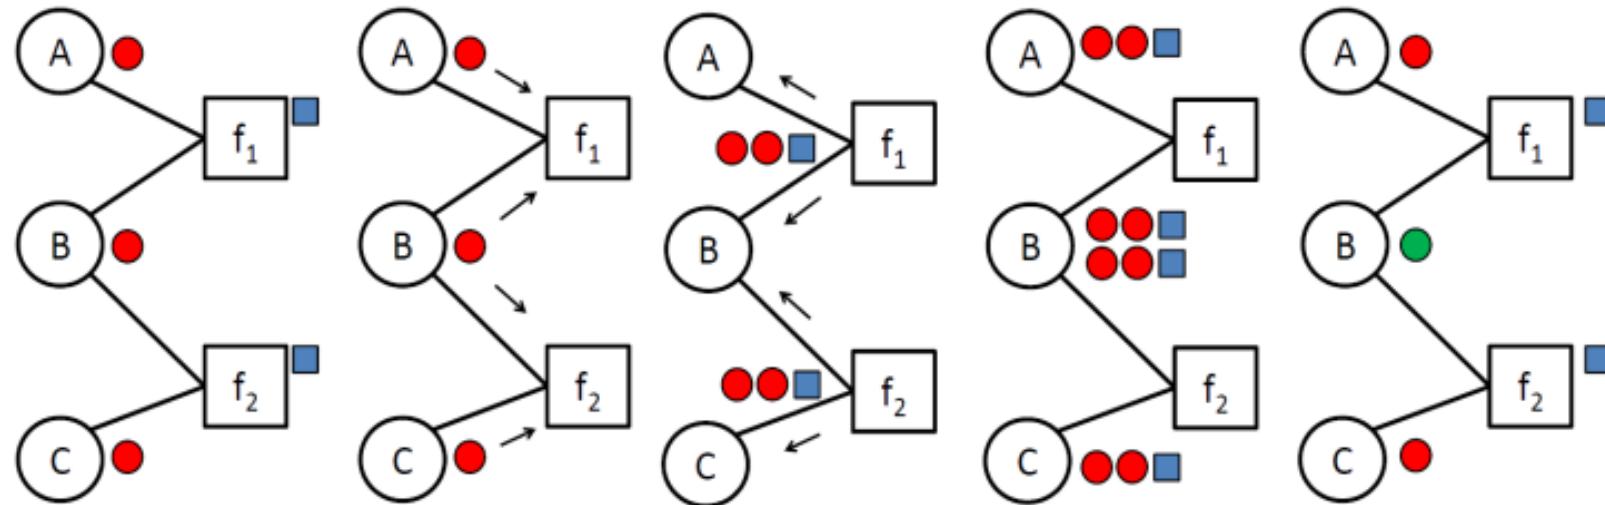


1. Each factor collects the colors of its neighboring nodes
2. Each factor „signs“ its color signature with its own color
3. Each node collects the signatures of its neighboring factors
4. Nodes are recolored according to the collected signatures



Compression: Pass the colors around*

*can also be done at the „lifted“, i.e., relational level

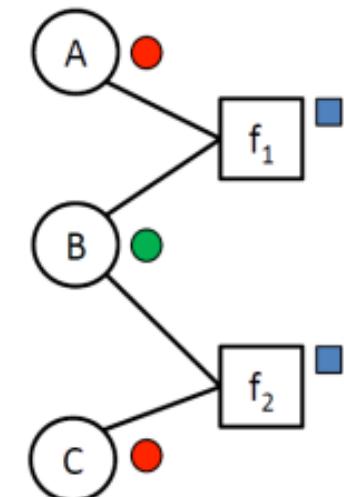
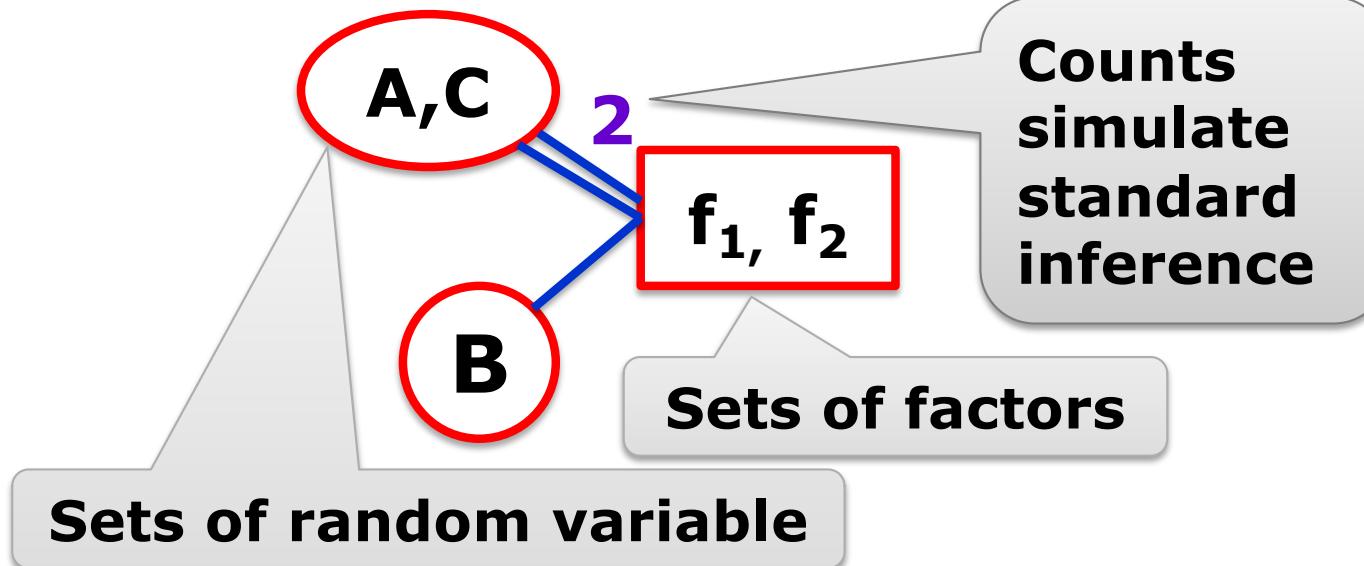


1. Each factor collects the colors of its neighboring nodes
2. Each factor „signs“ its color signature with its own color
3. Each node collects the signatures of its neighboring factors
4. Nodes are recolored according to the collected signatures
5. If no new color is created stop, otherwise go back to 1

Finally, we compute the quotient factor graph and run a modified loopy belief propagation

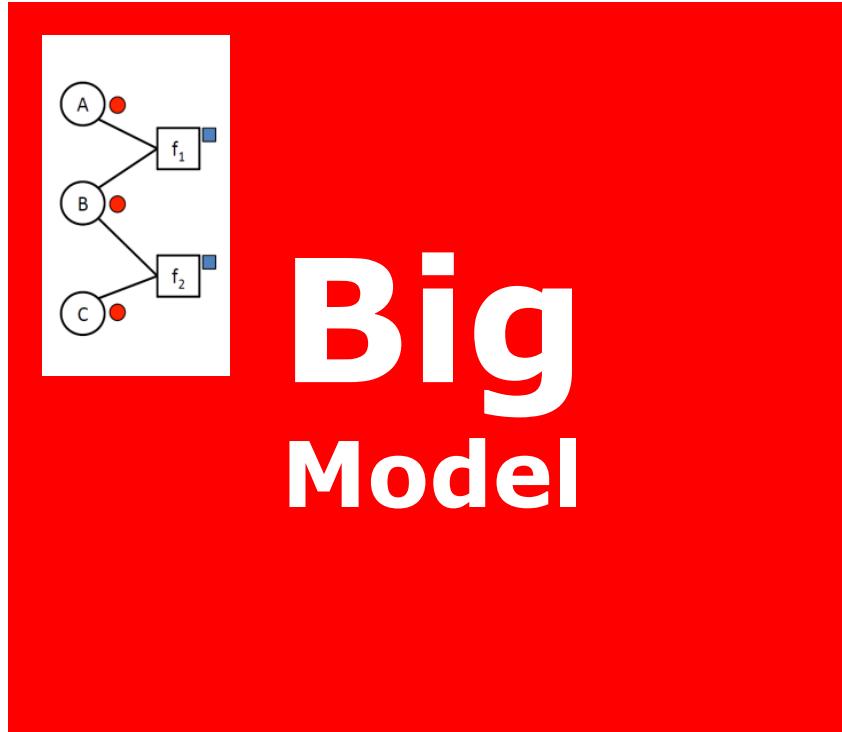
$$b_i(x_i) = \prod_{\mathfrak{f} \in \text{nb}(\mathfrak{X}_i)} \mu_{\mathfrak{f} \rightarrow \mathfrak{X}_i}(x_i)^{c(\mathfrak{f}, \mathfrak{X})}$$

$$\mu_{\mathfrak{X} \rightarrow \mathfrak{f}}(x) = \mu_{\mathfrak{f} \rightarrow \mathfrak{X}}(x)^{c(\mathfrak{f}, \mathfrak{X})-1} \cdot \prod_{\mathfrak{h} \in \text{nb}(\mathfrak{X}) \setminus \{\mathfrak{f}\}} \mu_{\mathfrak{h} \rightarrow \mathfrak{X}}(x)^{c(\mathfrak{h}, \mathfrak{X})}$$



Lifted Loopy Belief Propagation

Exploiting computational symmetries

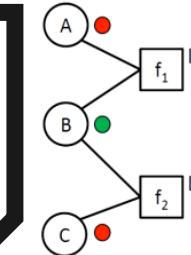
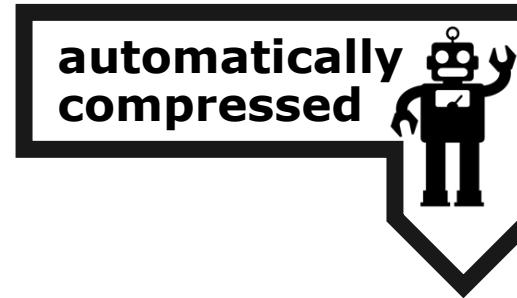


**Big
Model**

Run
Loopy Belief Propagation

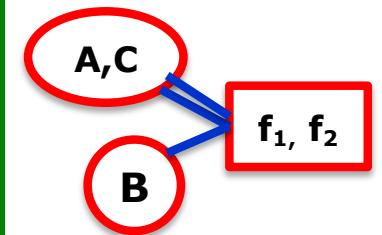
quasi-linear time

automatically
compressed



**Small
Model**

Run a modified
Loopy Belief Propagation



[Singla, Domingos AAAI'08; Kersting, Ahmadi, Natarajan UAI'09; Ahmadi, Kersting, Mladenov, Natarajan MLJ'13; Van Haaren, Van den Broeck, Meert, Davis MLJ'16]

Compression can considerably speed up inference and training

Probabilistic inference using lifted (loopy) belief propagation

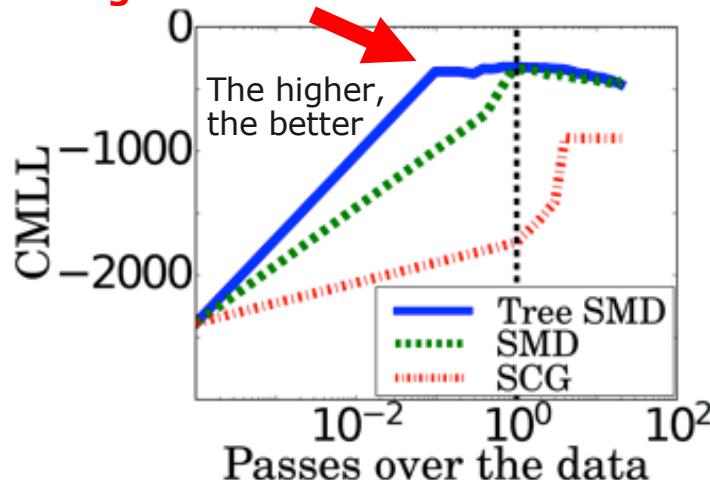
Domain	Time (in seconds)						The lower, the better		No. of (Super) Features	
	Construction		BP		Total					
	Ground	Lifted	Ground	Lifted	Ground	Lifted	Ground	Lifted	Ground	Lifted
Cora	263.1	1173.3	12368.4	3997.7	12631.6	5171.1	2078629	295468		
UW-CSE	6.9	22.1	1015.8	602.5	1022.8	624.7	217665	86459		
Friends & Smokers	38.8	89.7	10702.2	4.4	10741.0	94.2	1900905	58		

114x faster

Parameter training using a lifted stochastic gradient

CORA entity resolution

converges before data has been seen once



What is going on algebraically?

Can we generalize this to other data science approaches?

State-of-the-art

**WL computes a fractional automorphism
of some matrix A**

$$\mathbf{X}_Q \mathbf{A} = \mathbf{A} \mathbf{X}_P$$

where \mathbf{X}_Q and \mathbf{X}_P are doubly-stochastic matrixes (relaxed form of automorphism)

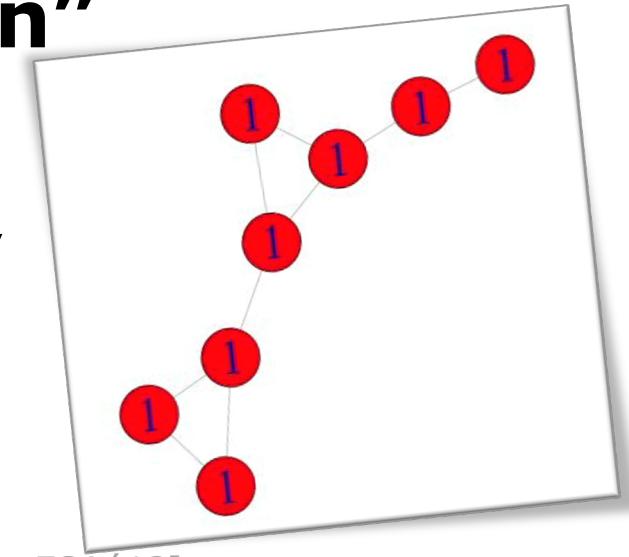
**It turns out that color passing is well
known in graph theory**

The Weisfeiler-Lehman Algorithm

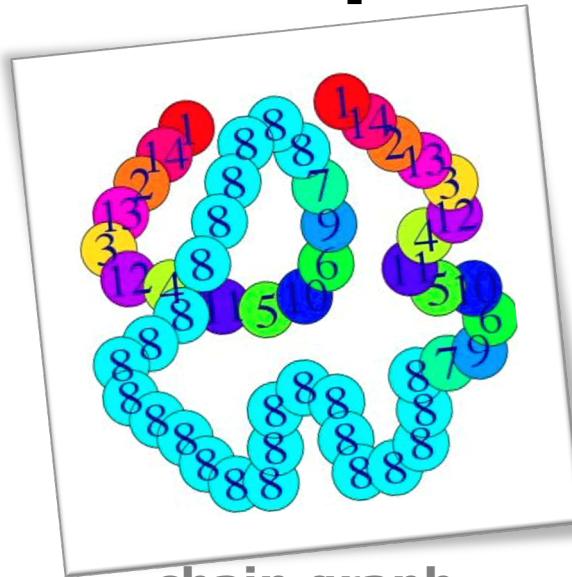


Weisfeiler-Lehman (WL) Algorithmus aka “naive vertex classification”

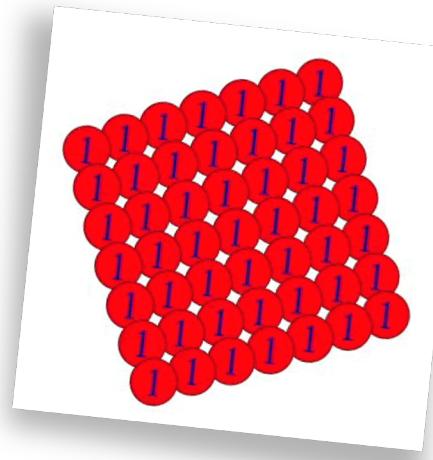
- Basic subroutine for GI testing
- Computes LP-relaxations of GA-ILP,
aka. **fractional automorphisms**
- **Quasi-linear** running time
 $O((n+m)\log(n))$ when using
asynchronous updates [Berkholz, Bonsma, Grohe ESA '13]
- **Part of graph tool SAUCY** [See e.g. Darga, Sakallah, Markov DAC '08]
- Has lead to highly performant graph kernels
[Shervashidze, Schweitzer, van Leeuwen, Mehlhorn, Borgwardt JMLR 12:2539-2561 '11]
- Can be extended to weighted graphs/real-valued
matrices [Grohe, Kersting, Mladenov, Selman ESA '14]
- Actually a Frank-Wolfe optimizer and can be
viewed as recursive spectral clustering
[Kersting, Mladenov, Garnett, Grohe AAAI '14]



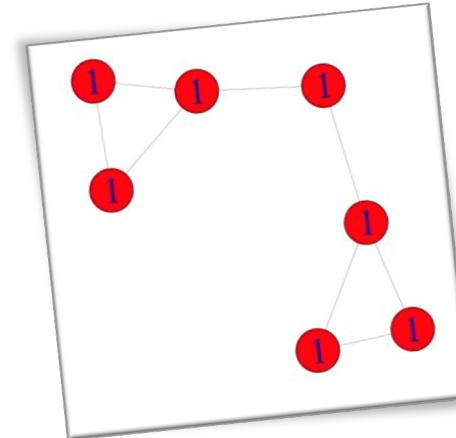
Examples



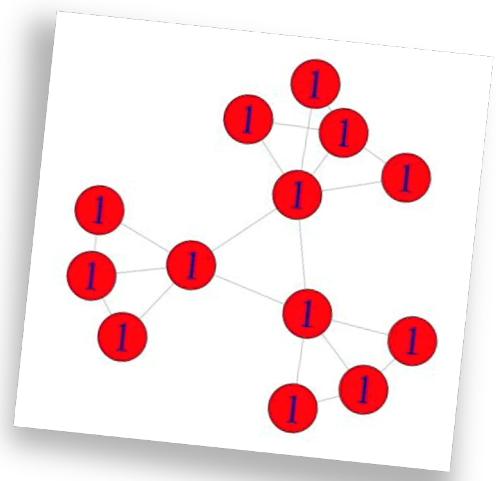
chain graph



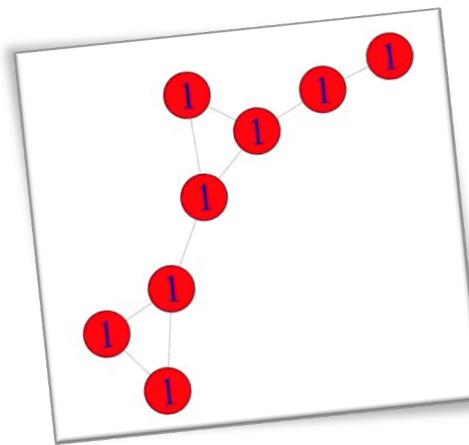
grid graph



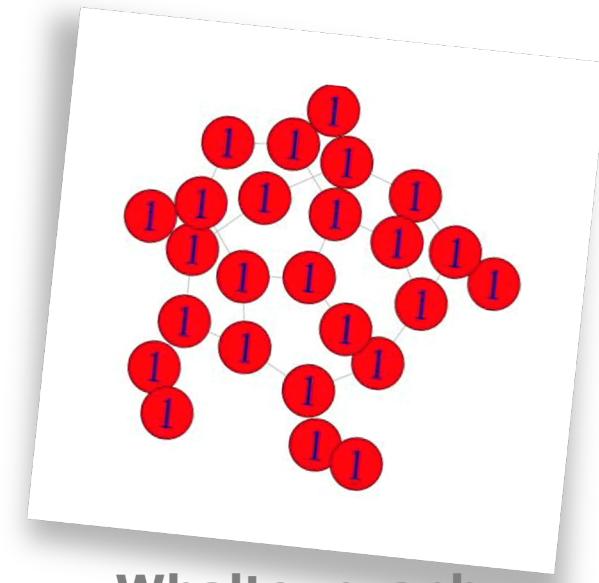
McKay graph



social rules



„walk“ graph



Whalter graph



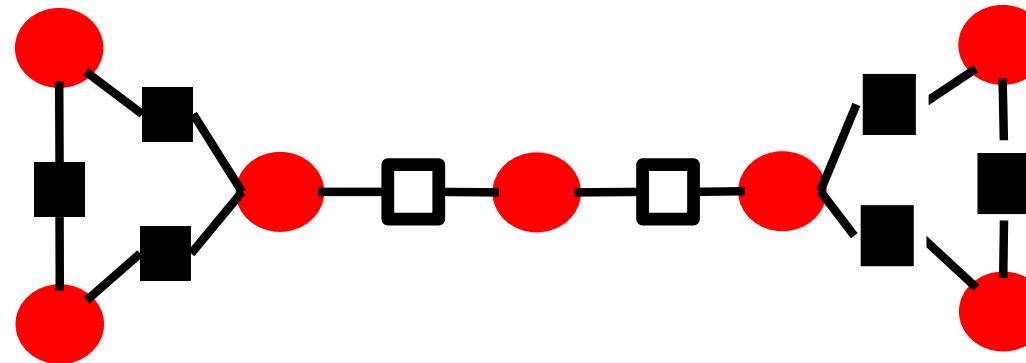
Why quasi-linear running time?



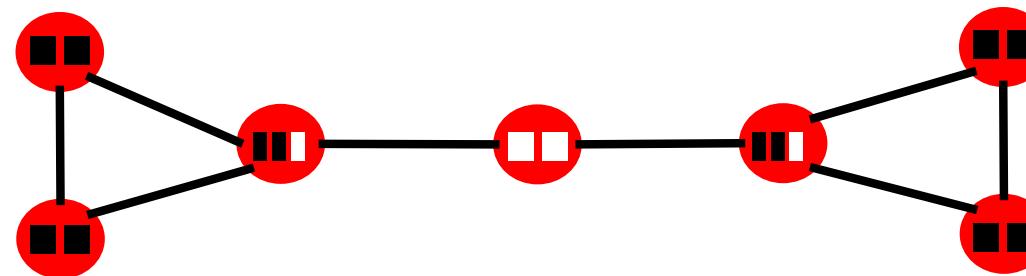
- Keep an active set of color classes and send color message asynchronously
 - Select an active color class but not the largest (halving argument)
 - Send message to the neighboring nodes of the class
 - If a color class refines, make it activ.
- Due to the halving argument, every node can become active only $O(\log(n))$ times.
- Message update is of order m



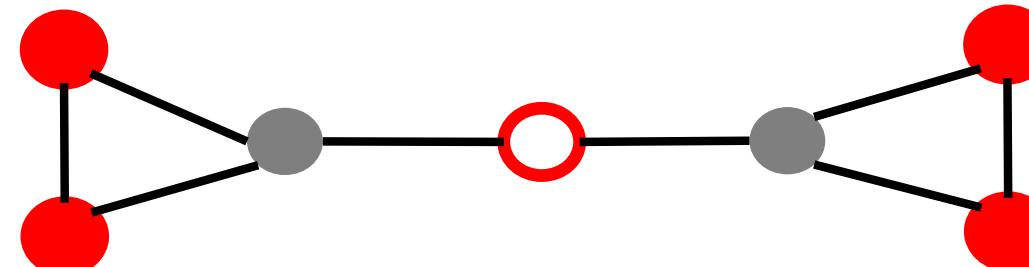
From Factor Graphs to Graphs



- Encode the factor colors into the node colors



- Then run Weisfeiler Lehman / Color Passing just on the graph



... more importantly, it paves the way to an algebraic understanding:

Instead of looking at ML through the glasses of probabilities over possible worlds, let's approach it using optimization



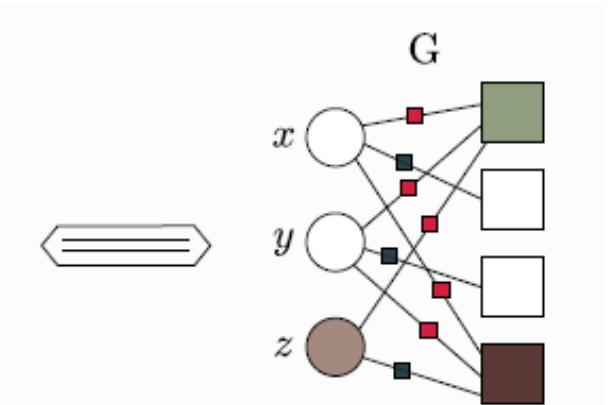
Lifted Linear Programming

$$\max_{[x,y,z]^T \in \mathbb{R}^3} 0x + 0y + 1z$$

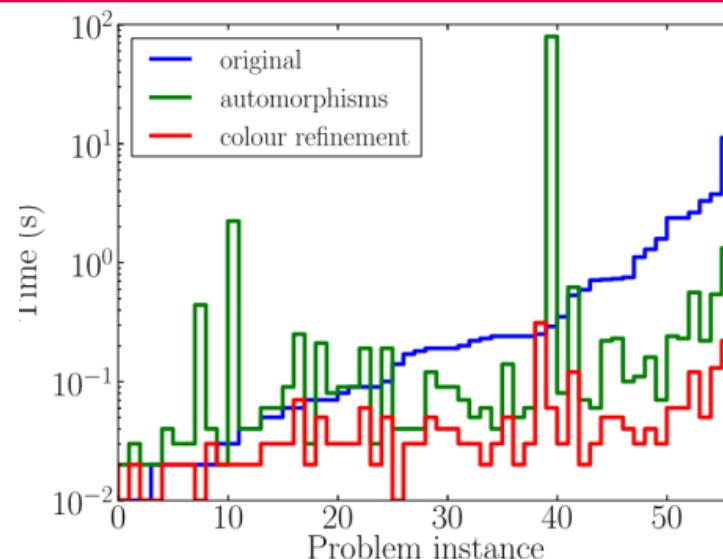
s.t.

$$\begin{bmatrix} 1 & 1 & 1 \\ -1 & 0 & 0 \\ 0 & -1 & 0 \\ 1 & 1 & -1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} \leq \begin{bmatrix} 1 \\ 0 \\ 0 \\ -1 \end{bmatrix}$$

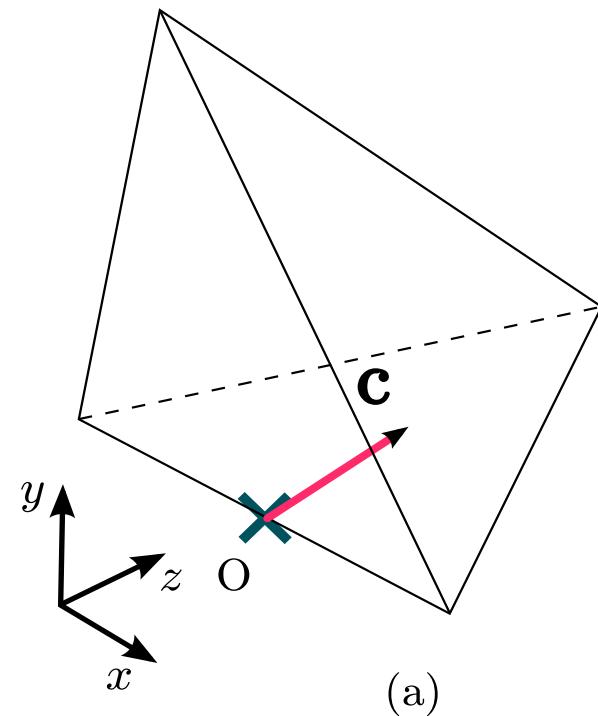
	x	y	z
c			█
A	█	█	█
b	█		



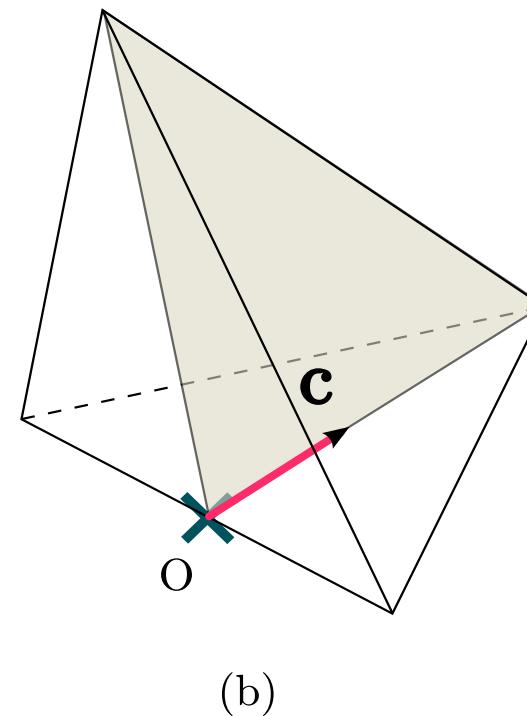
- (1) Reduce the LP by running WL on the LP-Graph
- (2) Run any solver on the (hopefully) smaller LP



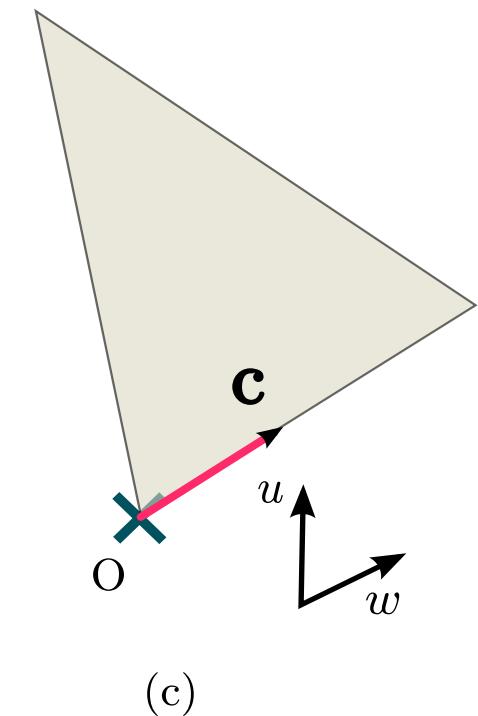
quasi-linear overhead
that may result in
exponential speed up



**Feasible region
of LP and the
objective
vectors**



**Span of the
fractional auto-
morphism of the LP**

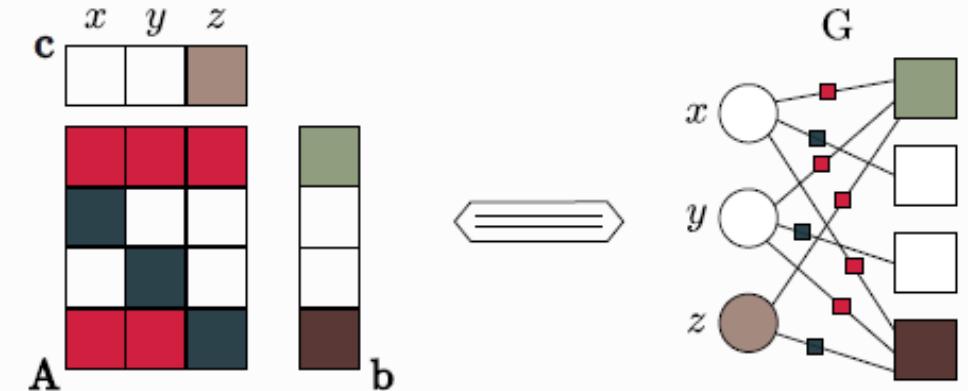


**Projections of the
feasible region
onto the span of
the fractional
auto-morphism**

Why does this work?



Compute Equitable Partition (EP) of the LP using WL



$$\mathcal{P} = \underbrace{\{P_1, \dots, P_p\}}_{\text{Partition of LP variables}; \text{Partition of LP constraints}}, \underbrace{\{Q_1, \dots, Q_q\}}_{\text{Partition of LP constraints}}$$

Intuitively, we group together variables resp. constraints that interact in the very same way in the LP.



Fractional Automorphisms of LPs

The EP induces a fractional automorphism of the coefficient matrix \mathbf{A}

$$\mathbf{X}_Q \mathbf{A} = \mathbf{A} \mathbf{X}_P$$

where \mathbf{X}_Q and \mathbf{X}_P are doubly-stochastic matrixes (relaxed form of automorphism)

$$(\mathbf{X}_P)_{ij} = \begin{cases} 1/|P| & \text{if both vertices } i, j \text{ are in the same } P, \\ 0 & \text{otherwise.} \end{cases}$$

$$(\mathbf{X}_Q)_{ij} = \begin{cases} 1/|Q| & \text{if both vertices } i, j \text{ are in the same } Q, \\ 0 & \text{otherwise} \end{cases}$$

Fractional Automorphisms Preserve Solutions

If \mathbf{x} is feasible, then $\mathbf{X}_p \mathbf{x}$ is feasible, too.

By induction, one can show that left-multiplying with a double-stochastic matrix preserves directions of inequalities. Hence,

$$\mathbf{A}\mathbf{x} \leq \mathbf{b} \Rightarrow \mathbf{X}_Q \mathbf{A}\mathbf{x} \leq \mathbf{X}_Q \mathbf{b} \Leftrightarrow \mathbf{A}\mathbf{X}_P \mathbf{x} \leq \mathbf{b}$$



Fractional Automorphisms Preserve Solutions

If \mathbf{x}^* is optimal, then $\mathbf{X}_p \mathbf{x}^*$ is optimal, too.

Since by construction $\mathbf{c}^T \mathbf{X}_P = \mathbf{c}^T$ and hence

$$\mathbf{c}^T (\mathbf{X}_P \mathbf{x}) = \mathbf{c}^T \mathbf{x}$$



What have we established so far?

Instead of considering the original LP

$$(\mathbf{A}, \mathbf{b}, \mathbf{c})$$

It is sufficient to consider

$$(\mathbf{A}\mathbf{X}_P, \mathbf{b}, \mathbf{X}_P^T \mathbf{c})$$

i.e. we “average” parts of the polytope.

But why is this dimensionality reduction?



Dimensionality Reduction

The doubly-stochastic matrix $\mathbf{X}_{\textcolor{green}{P}}$ can be written as

$$\mathbf{X}_{\textcolor{green}{P}} = \mathbf{B}\mathbf{B}^T$$

$$\mathbf{B}_{i\textcolor{green}{P}} = \begin{cases} \frac{1}{\sqrt{|\textcolor{green}{P}|}} & \text{if vertex } i \text{ belongs to part } \textcolor{green}{P}, \\ 0 & \text{otherwise.} \end{cases}$$

Since the column space of \mathbf{B} is equivalent to the span of $\mathbf{X}_{\textcolor{green}{P}}$, it is actually sufficient to consider only

$$(\mathbf{A}\mathbf{B}_{\textcolor{green}{P}}, \mathbf{b}, \mathbf{B}_{\textcolor{green}{P}}^T \mathbf{c})$$

This is of reduced size, and actually we can also drop any constraint that becomes identical



$$\mathbf{X}_Q \mathbf{A} = \mathbf{A} \mathbf{X}_P$$

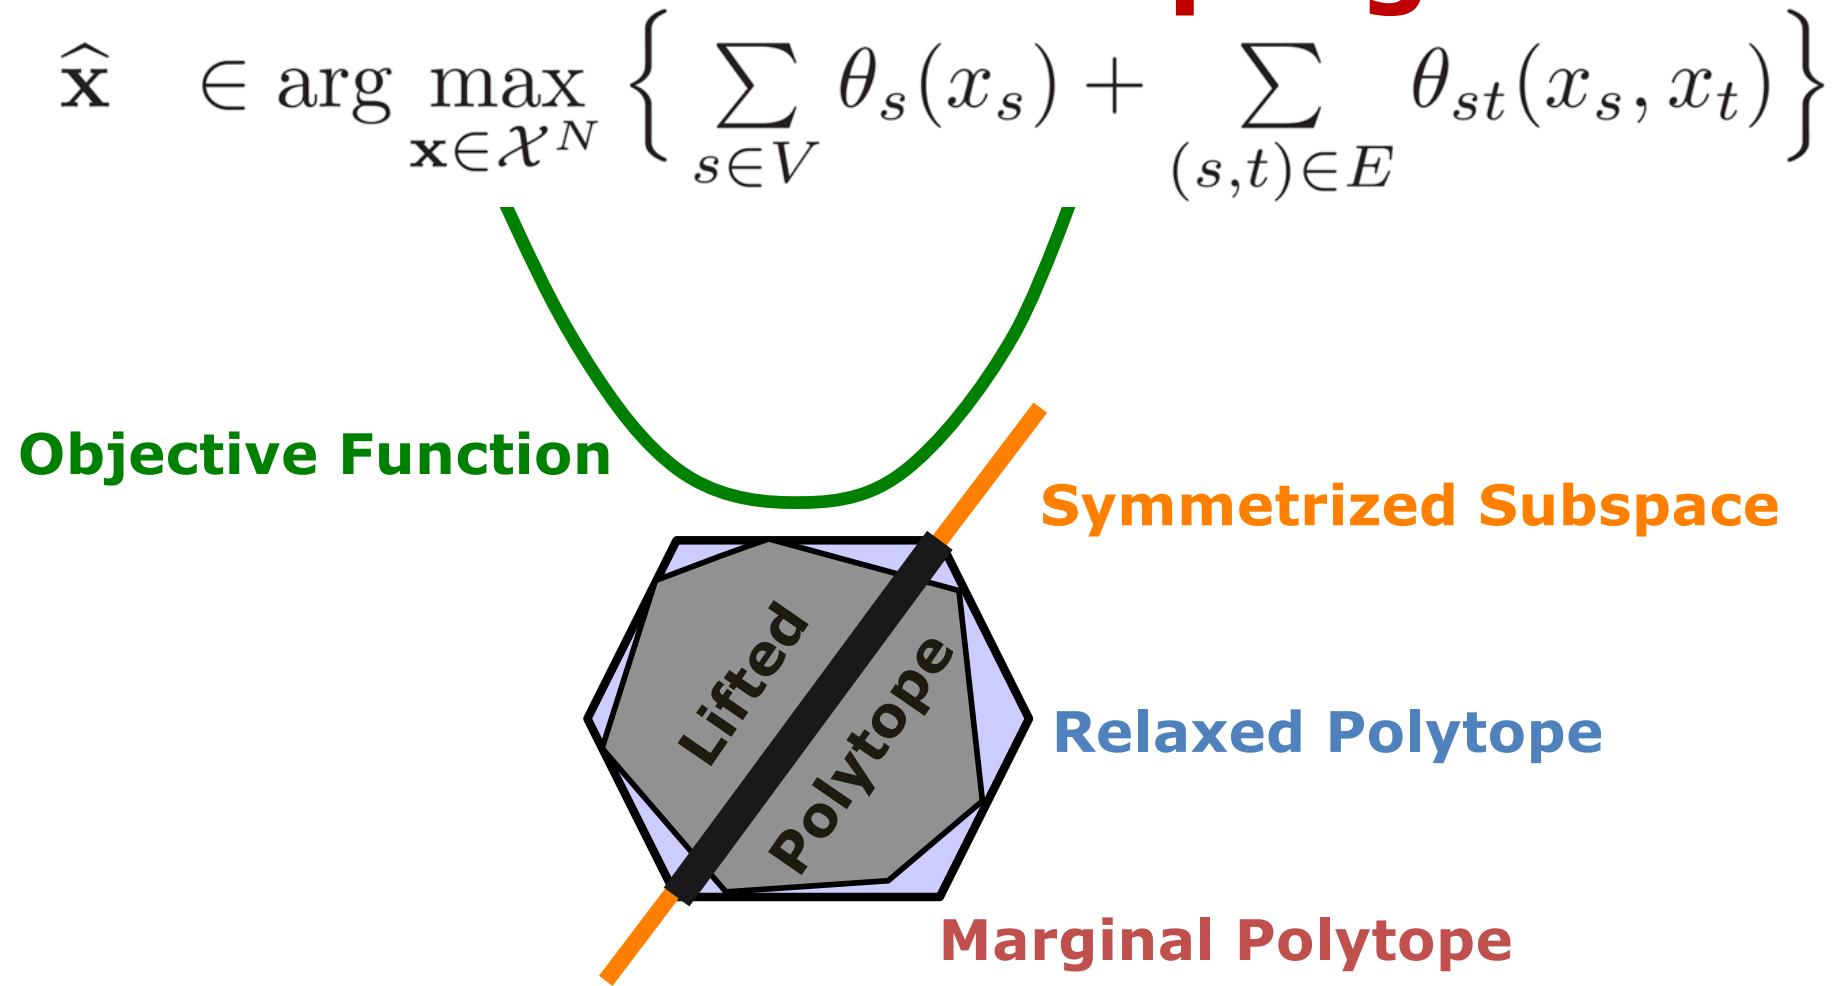
**Fractional automorphisms
provide an algebraic tool to
study lifted inference**

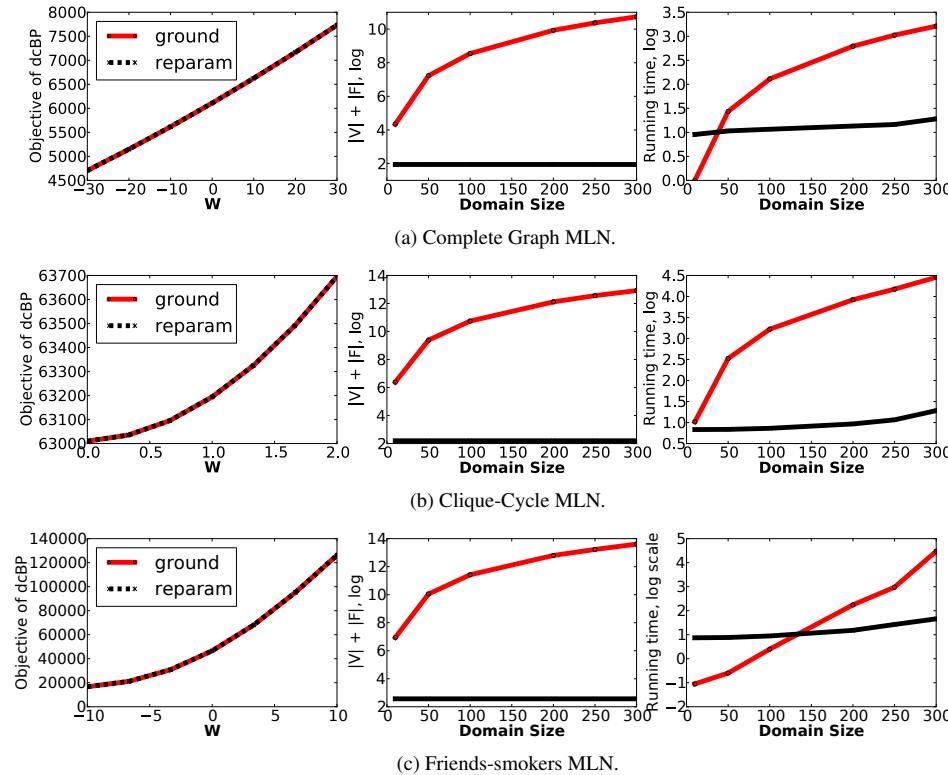
Actually, there is a whole body of work on (fractional) automorphisms for probabilistic inference, see the book, and we have focused here on the arguably simplest view.

**This has resulted in a number
of insights ...**



Approximate inference in graphical models is closely connected to linear programs



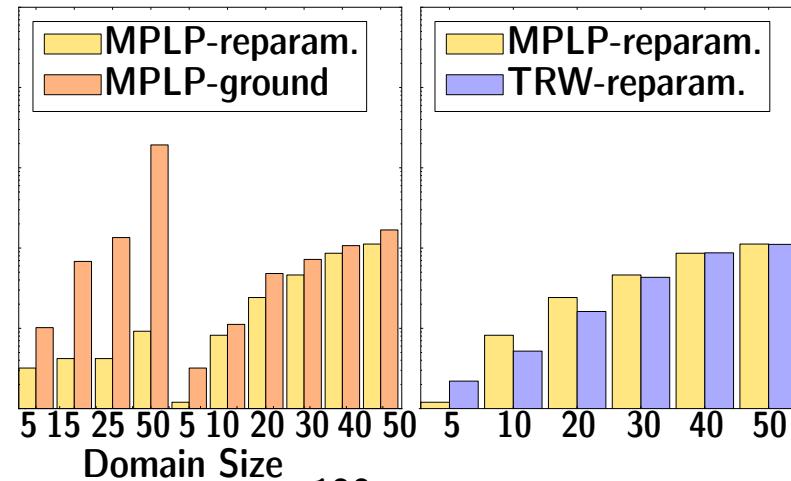


Marginals

Any concave free energy is liftable

MAP

Any MAP-LP message-passing approach is liftable



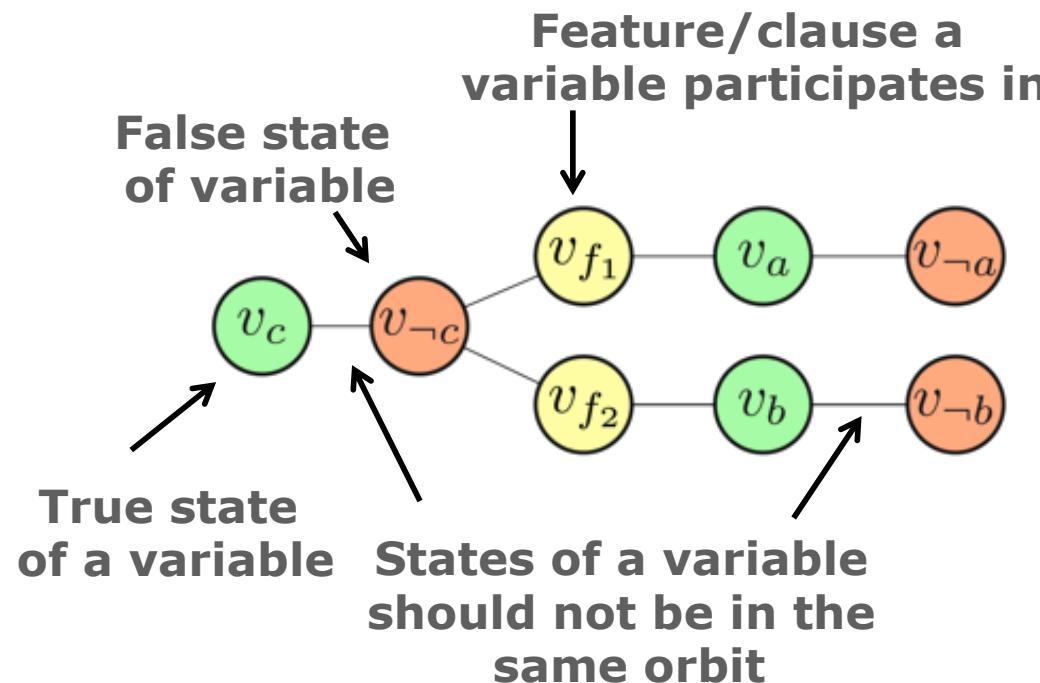
SYMMETRIES CAN ALSO BE EXPLOITED TO SPEED UP SAMPLING



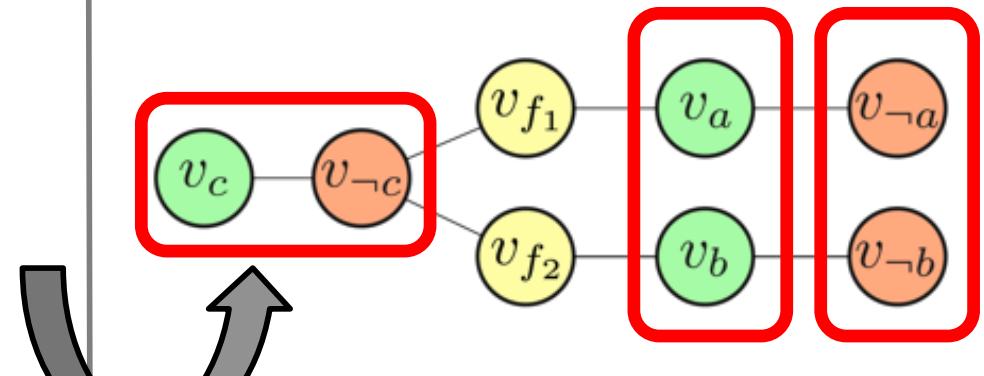
Orbital Markov Chain Monte Carlo

[Niepert UAI 2012, Van den Broeck, Niepert AAAI 2015]

true and false states have the same color, and all clauses/features that have the same weight



Symmetry classes of variables



Jump between symmetric states uniformly



Orbital MCMC Sampling

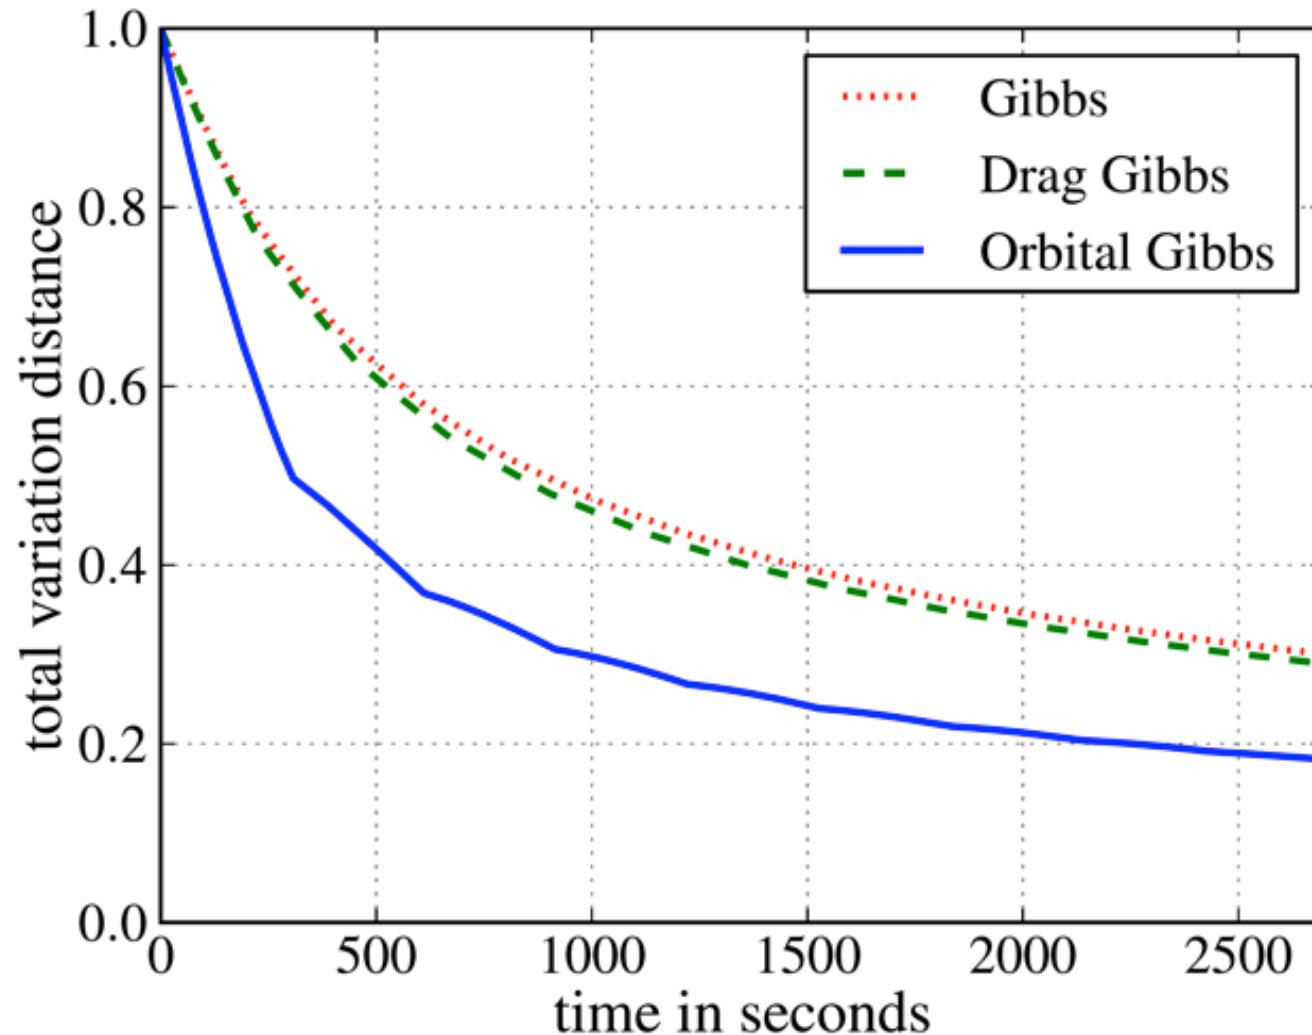
[Niepert UAI 2012, Van den Broeck, Niepert AAAI 2015]

In each sampling iteration:

- 1. run a step of a traditional MCMC chain TM first and then**
- 2. sample the state of M at the next time uniformly at random from the orbit of the state of the original chain TM at time t, i.e., select an equivalent state uniformly at random**



Orbital MCMC on a 6x6 Ising grid



Lifted Metropolis-Hastings

Given an orbital Metropolis chain A:

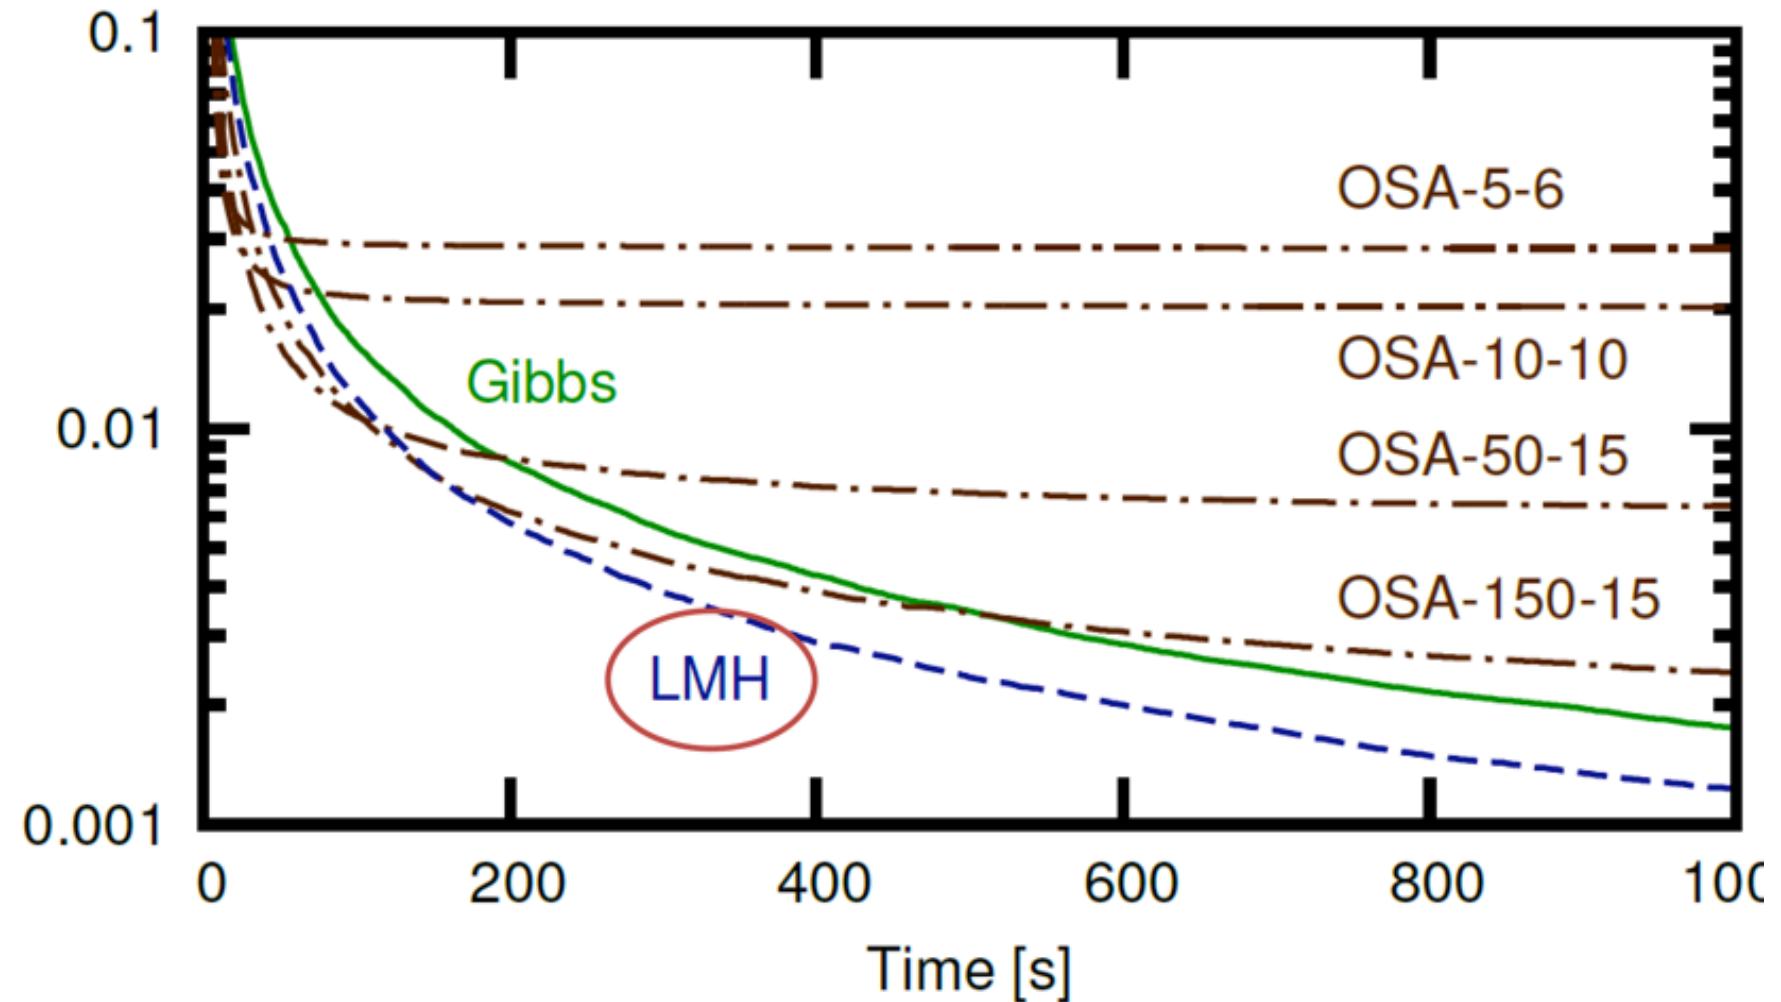
- Given symmetry group G (approx. symmetries)
 - Orbit x^G contains all states approx. symm. to x
 - In state x :
 1. Select y uniformly at random from x^G
 2. Move from x to y with probability $\min\left(\frac{\Pr(y)}{\Pr(x)}, 1\right)$
 3. Otherwise: stay in x (reject)
 4. Repeat
- Color-passing /Saucy**

and an ordinary (base) Markov chain B, with prob. a follow B and with (1-a) follow A

This can also account for evidence that may break symmetries, using e.g. approx. symmetries

Lifted Metropolis-Hastings on WebKB

[Niepert UAI 2012, Van den Broeck, Niepert AAAI 2015]



What have we learnt?

- Lifted inference exploits (fractional) symmetries
- Fractional symmetries can be computed in quasi-linear time
- Symmetries allow one to study lifted inference in an algebraic way, i.e., independent of the underlying algorithm
- Essentially, the whole family of approximate inference methods is liftable

